

**CS-5009: Advanced
Operating Systems**

Serial No:

Midterm Exam
Total Time: 2 Hours
Total Marks: 90

Wednesday, 03rd November, 2021

Course Instructor

Dr. Ehtesham Zahoor

Signature of Invigilator

Student Name

Roll No

Section

Signature

DO NOT OPEN THE QUESTION BOOK OR START UNTIL INSTRUCTED
Instructions:

1. Attempt on question paper. Attempt all of them. Read the question carefully, understand the question, and then attempt it.
2. No additional sheet will be provided for rough work. You can however use the back of last page for any rough work.
3. After asked to commence the exam, please verify that you have Nine (09) different printed pages including this title page. There are 2 questions in total.
4. Use permanent ink pens only. Any part done using soft pencil will not be marked and cannot be claimed for rechecking.
5. You need to be as specific as possible and provide the answers in the space provided. Cutting, overwriting and ambiguous answers would be considered as incorrect.
6. Please make sure to properly comment your code and write clearly, failure to do so can cost you 10 marks. Code containing overwritten text and corrections cannot be claimed for rechecking. You need to write code specifically in the space given.

The paper is open book and you can use any printed helping material. Keep it clean and do not share anything. You cannot use any electronic device in the exam.

	Q-1	Q-2	Total
Marks Obtained			
Total Marks	45	45	90

You need to be as specific as possible and provide the answers in the space provided. Overwriting and ambiguous answers would be considered as incorrect. It is important to read the question carefully and clearly provide the answer so think before you write anything.

Question 1 - Short Questions [15x3 = 45Marks]

- I. We studied about the 50,000-core supercomputer built on AWS for cancer research. How do *AWS staff and administrators* help build the supercomputer?

The answer is no because as we studied that Aws staff and administration has no contribution in building the supercomputers for cancer research.

- II. In reference to the Beowulf cluster that we built, what is the purpose of the command below and what does the `|` in the command do?

```
echo "/mirror *(rw, sync)" | sudo tee -a /etc/exports
```

Here the output of one command will be given as input to other .
The tee command is used for read and write purpose. The `\mirror` folder in master node will be shared with slave nodes in read-write mode.

- III. In reference to the Beowulf cluster that we built, is it compulsory to add the mount information at Slave machines to the ***/etc/fstab***

Fstab is used in order to mount it on every boot. We mount the external storage so that every slave node can get access of it . So thats why use this.

- IV. In reference to the Beowulf cluster that we built, what would be the implications on connectivity if we change the home directory of `mpiuser` to some other directory than the `/mirror`

As we know that only mirror folder is shared so if we change the home directory the `/mirror` will not be accessible by ssh.

- V. Consider that `n` processes are using bakery algorithm for synchronization. If the token received by a process is higher than that of all other processes, in which **state**, as per the Five-state Protocol Model, the process would be? Write one-line justification as well.

IX. We built a docker image containing a Web server. The content of the *dockerfile* are as below. You need to discuss the purpose of both lines clearly stating what is *nginx* and *alpine* and what is being copied from where to where?

```
FROM nginx:alpine
COPY . /usr/share/nginx/html
```

Here we are copying the web page from Local directory to nginx html folder.

X. In the code segment below (assume all includes to be present and no syntax errors) is there a race condition and/or critical section. If so, fix it elegantly using any synchronization primitive we have studied. You need to only write the modified part in the **space provided**.

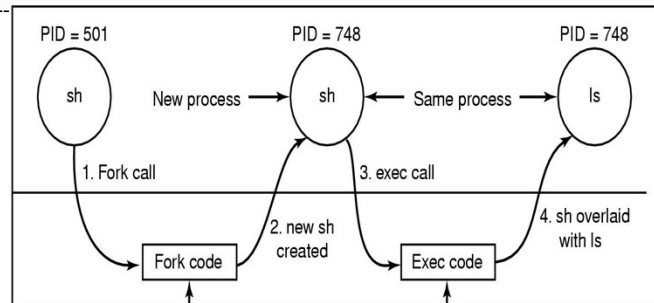
```
#define NUM_THREADS 10
int sharedData = 21;
void* incrementData(void* arg) {
    localData = sharedData;
    pthread_exit(NULL);
}
int main()
{
    pthread_t threadID;
    for (int counter=0; counter<NUM_THREADS;counter++) {
        pthread_create(&threadID, NULL, incrementData, NULL);
    }
    cout << "ThreadCount:" << sharedData << endl;
    pthread_exit(NULL);
}
```

In the above code there is no critical section that exists.

XI. Is it possible that a computer system runs out of file descriptors? Discuss this in reference to *file_struct*, a kernel-resident array data structure containing the details of open files.

Yes it is possible that computer system runs out of file descriptors. It contains information about the open files in it but it may end up with this.

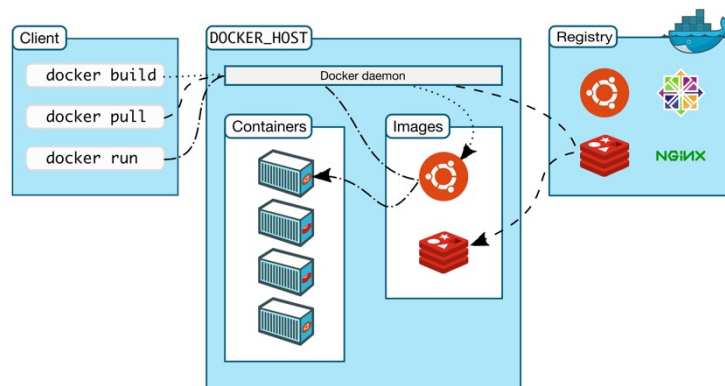
There will be busy wait which means it will be in Ready state.



VI. Consider the steps in executing the command `ls` type to the shell, as shown in the figure above. Is it possible to run the `ls` command, **without first** spawning a new process using `fork`, i.e. by removing the step 1&2? What would be the implications of this choice then?

Yes it can be possible

But there will be limitation that the `ls` command will be executed only once and no other will be allowed to be executed after this.



VII. Consider running the hello-world docker image on your personal system for the very first time. You need to discuss which elements act as client, docker host and the registry, show above.

Docker_hub will be registry

Our own machine will be docker host

Terminal will be client

VIII. We built a docker image containing a Web server. The command needed to run the image is `docker run -p 80:80 ehtesham/staticweb`. What is the purpose of `-p` flag?

It Maps the system port to the web server port. So it is basically used for port Mapping.

XII. Why everything in Unix (and its flavors) is considered as a file (or specifically a file-descriptor)?

In Unix, everything is considered as a file in which two functions are available: read and write.

XIII. What is the logical (not syntactic) error in the code below and how it can be fixed? You need to just discuss the solution without modifying the code provided.

```
int main(int argc, char* argv[]) {
    int sum, counter, inputList[6] = {11,45,3,5,12,-3};
    #pragma omp parallel num_threads(2)
    {
        #pragma omp for schedule(static, 3)
        for (counter=0; counter<6; counter++) {
            sum+=inputList[counter];
        }
    }
    printf("The summed up Value: %d", sum);
}
```

The shared variable `sum` is the logical error and it can be prevented by using Locks.

XIV. Why the Semaphore and pthread mutex locks are called the hardware locks?

When an instruction can do both set and lock, it is a hardware lock and in case of interrupts, prevention hardware locking mechanisms are used.

XV. How blockchain is related to our course? More specifically, which aspect of decentralized blockchain relates to our course?

It uses a support system that runs in the background of a computer system. The synchronization is related to our course.

Question 2 - Fork, Pthreads, OpenMP and MPI [10+15+10+10 = 45 Marks]

You need to be as specific as possible and provide the answers in the space provided. Overwriting and ambiguous answers would be considered as incorrect. It is important to think before you write anything.

Please make sure to properly comment your code and write clearly, failure to do so will result in a loss of 10 marks. Code containing overwritten text and corrections cannot be claimed for marks.

You have been asked to write code for different parts of the problem in parts, please note that the complete program should work as a whole, that is only code for some sub problem is not acceptable.

- I. You need to write a program where the parent process initializes an array of 10 integers. It then creates a child process, which then finds the maximum number amongst the array. The child process reports back to parent process, which displays the results. (3 points)

/ write the code to initialize the array (3 points) */*

```
int arr[10] = {1,2,3,4,5,6,7,8,9,10};
```

```
pid= fork();
```

/ write the code for the rest of the question (1+3+3 points) */*

```
if(pid==0)  \\child
{
for(i=1; i<10; i++)  \\using sorting for array containing 10 elements
{
if(arr[0]<arr[i])
{
arr[0]= arr[i];
}
}
int maximum= arr[0];
else  \\parent
{
wait(&status);
\\returns the exit status of the child
Printf(" The maximum number is : %d\\n", WEXITSTATUS(status));
}
}
```

- II. You need to write a program where the master thread initializes an array of n (you can initialize n to any number) integers. It then creates n/2 threads (using pthread API) and passes an index number to each thread. Each thread then sums up two numbers from the array (index and index+1, where index is passed by the master thread) and writes the sum to a shared variable in a synchronized manner. (3 points)

/ write the code to initialize the array */*

```
int *arr = (int *)malloc((n) * sizeof(int));
printf("The element of arrays are: ");
for (int i = 0; i < n; i++)
{
arr[i] = rand() % 11;
printf("%d ", arr[i]); }
}
```

/ write the code to create n/2 threads (using pthread API) and pass an index number to each thread (6points)*

```
int N,n;
printf("Enter the number of threads: ");
// scanf("%d", &N);
n= N/2;
printf("Number = %d",n);
pthread_t th[n];
int i;
for (i = 0; i < n; i++) {
    int* a = malloc(sizeof(int));
    *a = i * 2;
    if (pthread_create(&th[i], NULL, &routine, a) != 0) {
        perror("Failed to create thread"); }
}
```

/ write the code for each thread to sum given numbers from the array (3points) and write it to the shared variable in a synchronized manner (6points)*

```
int Sum = 0;
for (i = 0; i < n; i++) {
    int* r;
    if (pthread_join(th[i], (void**) &r) != 0) {
        perror("Failed to join thread");
    }
    Sum += *r;
    free(r);
}
printf("Sum: %d\n", Sum);
```

III. You need to redo the question in previous part using OpenMP instead of pthreads. ignore the array initialization part and more specific details are given below.

/ write the code to create n/2 threads (using openMP) and somehow pass an index number to each thread (2+3points)*

```
#pragma omp parallel num_threads(n/2)
```

/ write the code for each thread to sum given numbers from the array (3points) and write it to the shared variable in a synchronized manner (2points) */*

```
for(i=0; i<=n/2; i++)
{
sum= arr[i]+arr[i+1];
i=i+1;
#pragma omp critical
{
total_sum=total_sum+sum;
}
```

IV. You need to write a program, using MPI, where each process (amongst any n processes) initializes a number to be its rank*2, process ranked 4 initializes the number to be its rank*4, process then **broadcasts** its number, and when the code completes each process has the **numbers** of all the processes. For communication, you are allowed to use **MPI_Bcast only**.

/ write the code where each process initializes its number to be its rank*2 (4points) */*

```
MPI_Init(&argc, &argv);
MPI_Comm_size(MPI_COMM_WORLD, &size);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Get_processor_name(processorName, &nameLen);

for(i=0; i<rank_num.size(); i++)
{
rank_num[i]=rank*2;
}
```

/ write the code where each process broadcast its number to all others, that is, when the code executes each process has the localNumbers of all process. **(6points)***

```
for(i=0; i<rank_num.size(); i++)  
{  
    MPI_Bcast(rank_num[i], 13, MPI_INT, i, MPI_COMM_WORLD);  
    printf( " Message from process %d\n", rank_num[i]);  
}
```