

# Operating Systems

---

## **2. Process Concept**

# Process: The Concept

---

- Process = A program in execution
  - Execution progress in a sequential fashion
- Example processes
  - OS kernel
  - OS shell
  - Program executing after compilation
  - www-browser
- Process management by OS
  - Allocate resources
    - Processor, main memory, I/O modules, timers, ...
  - Schedule: Interleave their execution
    - Watch for processor utilization, response time
  - Inter-process communication, synchronization
    - Check potential deadlocks
    - Interleaving and non determinism imply increased difficulties!

# The Process

---

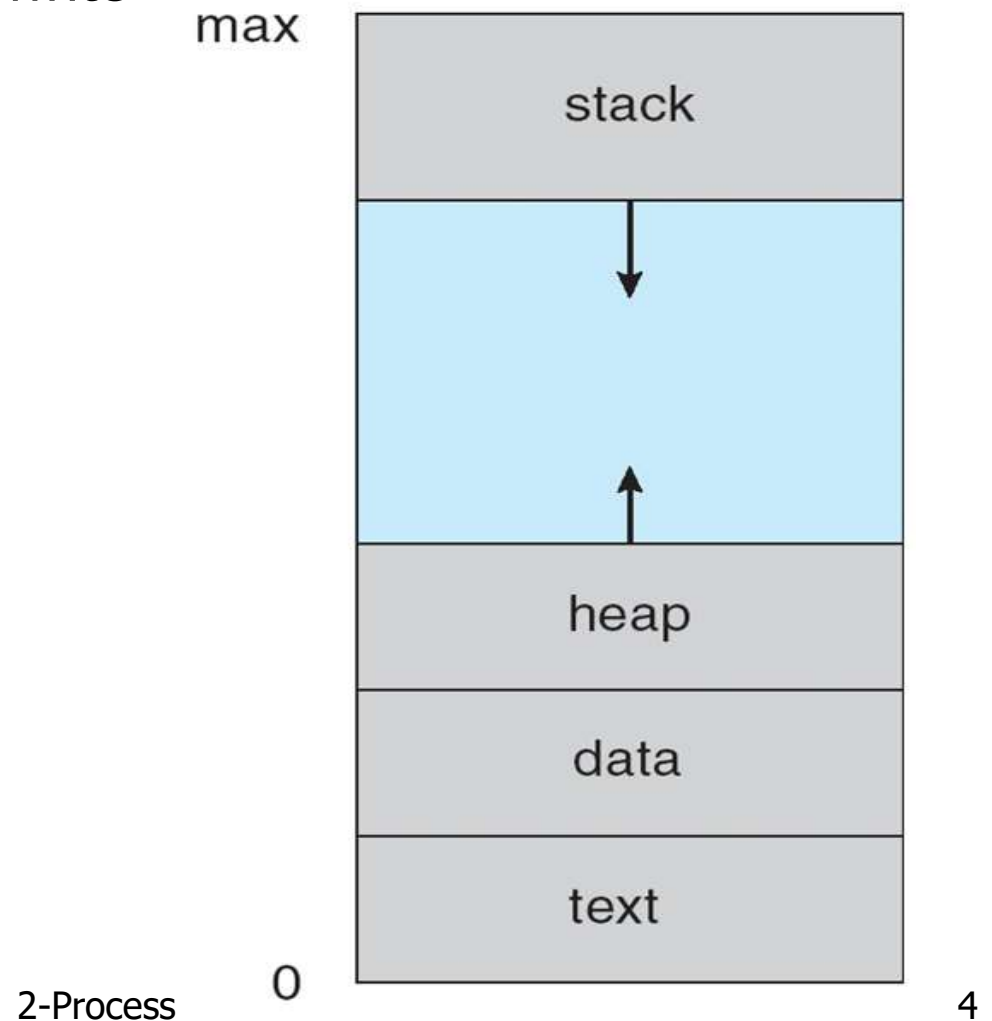
A process consists of **multiple parts**

- Program code, also called text section
- Current activity including program counter, processor registers
- Stack containing temporary data
  - Function parameters, return addresses, local variables
- Data section containing global variables
- Heap containing memory dynamically allocated during run time

# Process Address Space

---

- A list of memory locations from some min (usually 0) to some max that a process can read and write



# Relation Between Program And Process

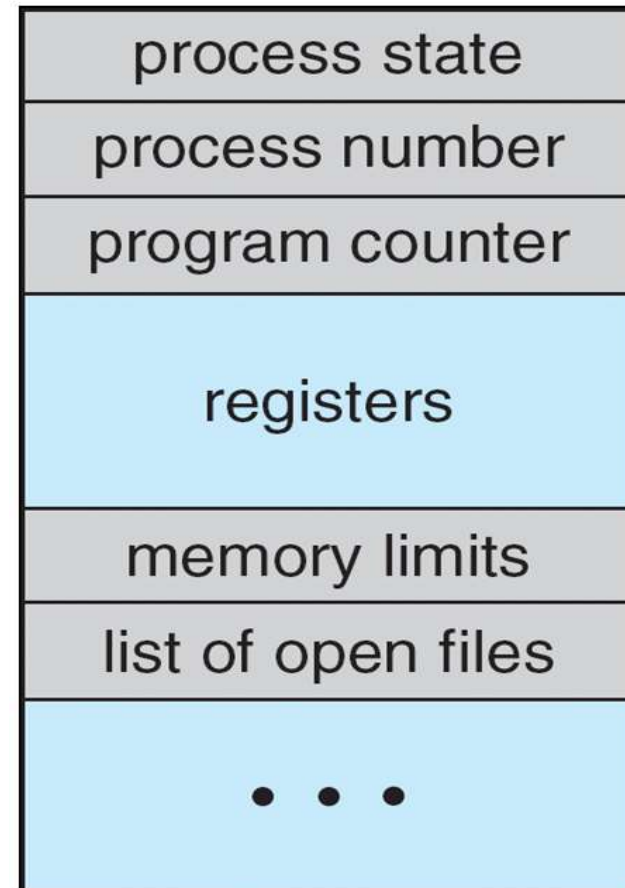
---

- Program is passive entity, process is active
  - Program: Static code + static data
  - Process: Dynamic instantiation of code + data + more
- Program becomes process when executable file loaded into memory
- No 1 to 1 mapping between program and process
  - One program may run many processes
    - Separate execution sequences
  - Multiple users may execute the same program
    - Text section is equivalent

# Process Control Block

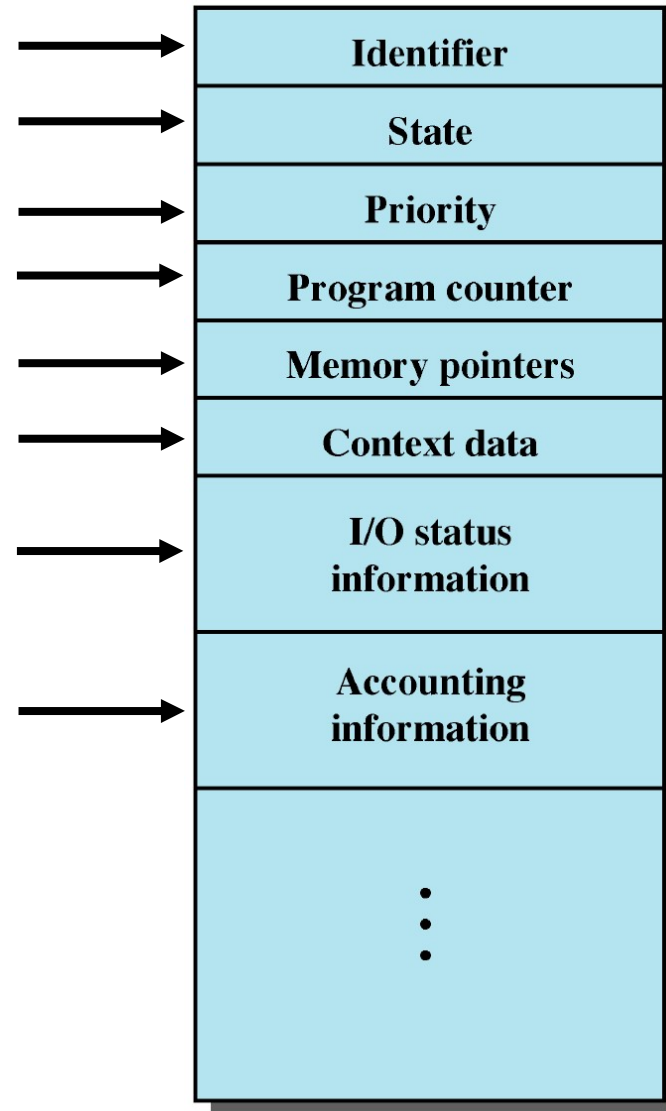
---

- OS keeps all the data it needs about a process in the process control block (PCB)
  - Process state
  - Program counter
  - CPU registers
  - CPU scheduling information
  - Memory-management information
  - Accounting information
  - I/O status information



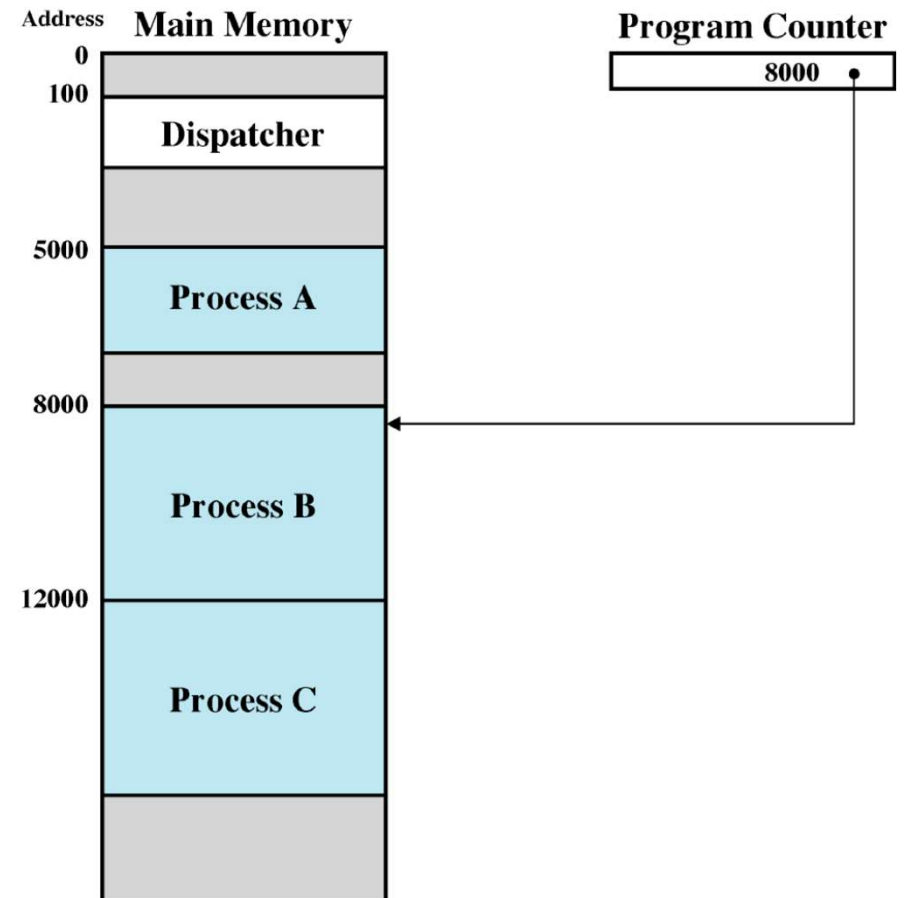
# Process Control Block

- Distinction between processes
- Knowledge on processes ready to execute
- Discriminate between processes
- Know the next instruction to execute
- Find the program and data stored in memory
- Keep information about registers
- Keep track about resources, e.g., I/O devices
- Collect information on utilization of the system



# Dispatcher/Scheduler

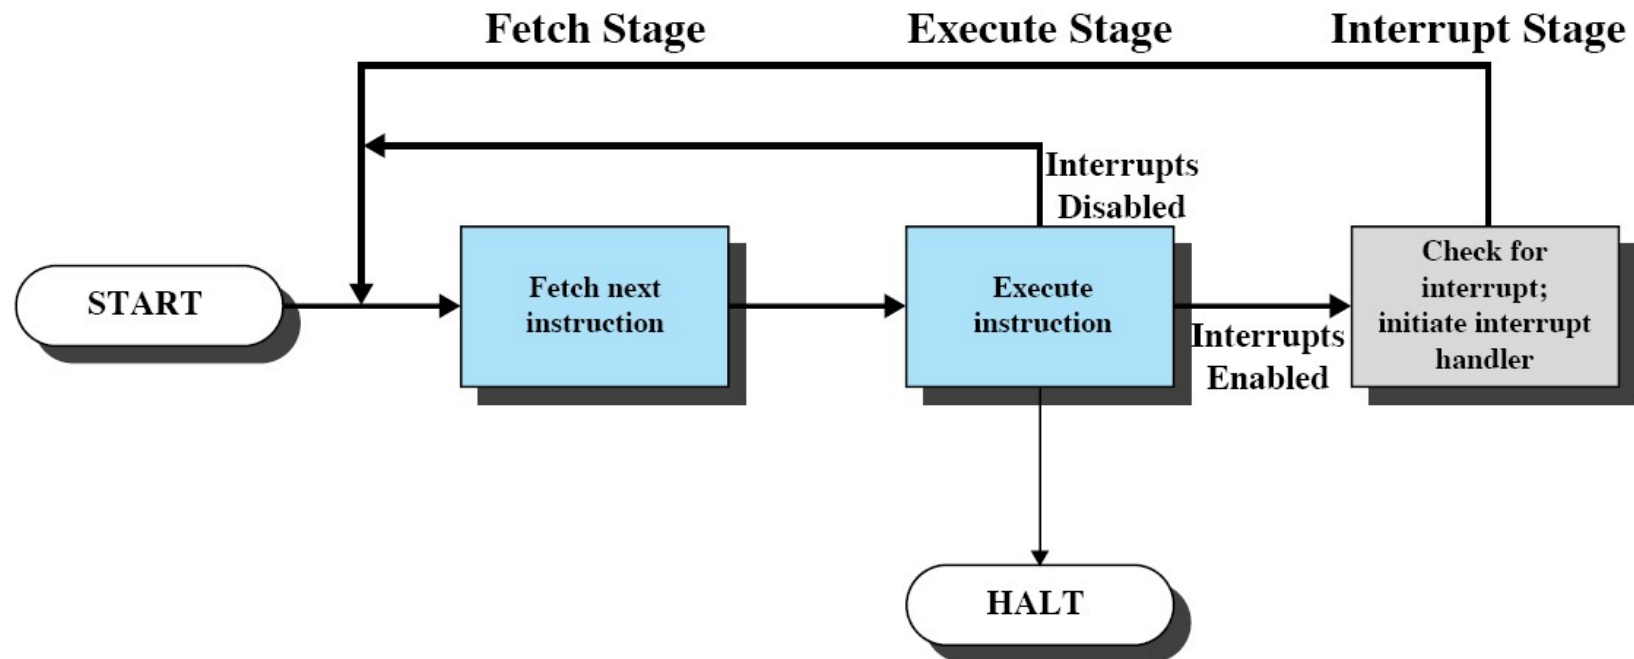
- OS program that decides which process to run next
- Uses Process Block Control (PBC) information
  - Decision what is the next process to execute





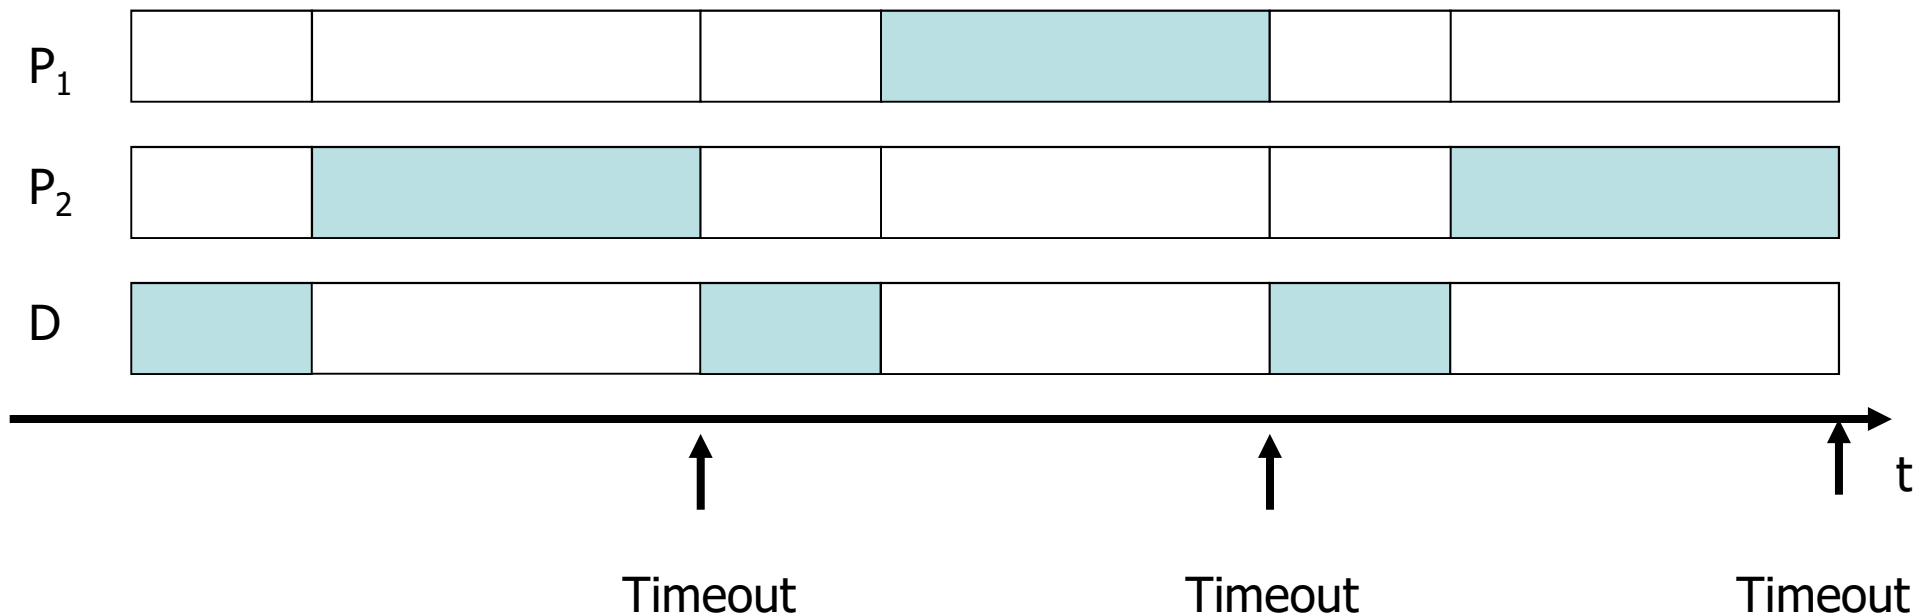
# How Does the Dispatcher Gains Control?

- Register interrupt handlers
  - Timeout
  - I/O request



## Execution with Timeouts


- On timeout dispatcher gains control
- Decides which process to execute next
- Dispatcher should perform as few instructions as possible



# Two-State Process Model

---

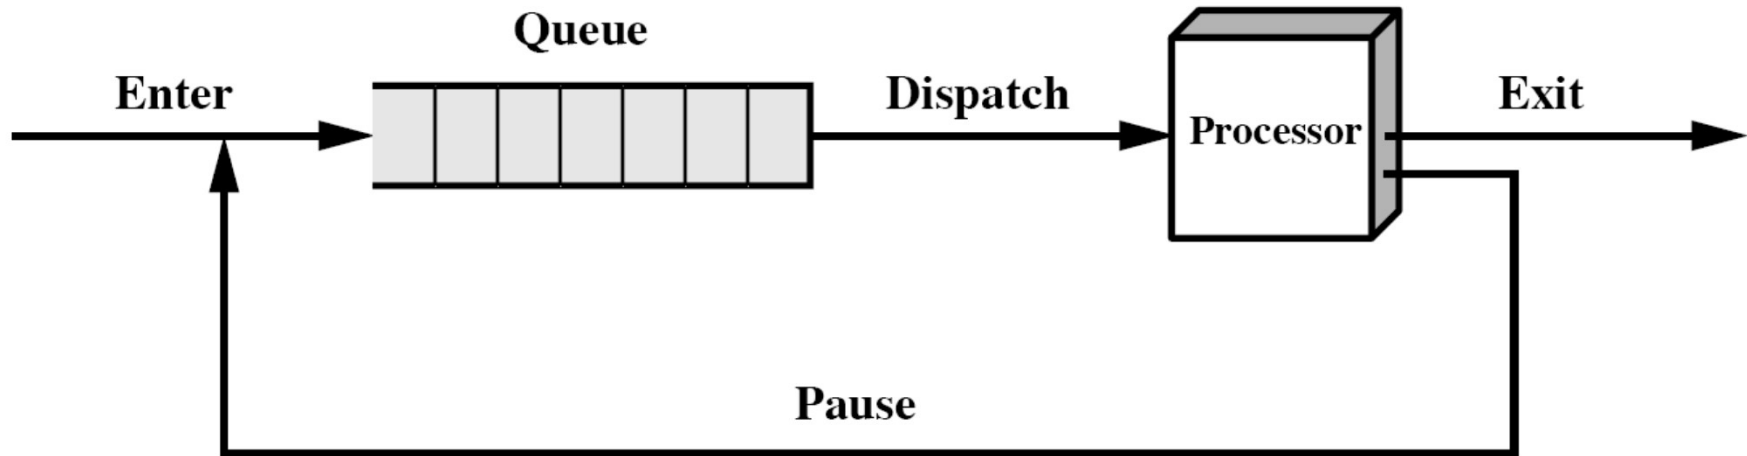
- At any time a process is either being executed by a processor or not



# How to Use States for Scheduling?

---

- Queue of not-running processes
- Queues
  - FIFO, Priority Queue
- Design goal
  - Execute only few operations



# Is Two Process Model Sufficient?

---

Ready


```
→ a := 1
  b := a + 1
  c := b + 1
  read a file
  a := b - c
  c := c * b
  b := 0
```

Ready

```
→ a := 1
  read a file
  b := a + 1
  c := b + 1
  a := b - c
  c := c * b
  b := 0
```

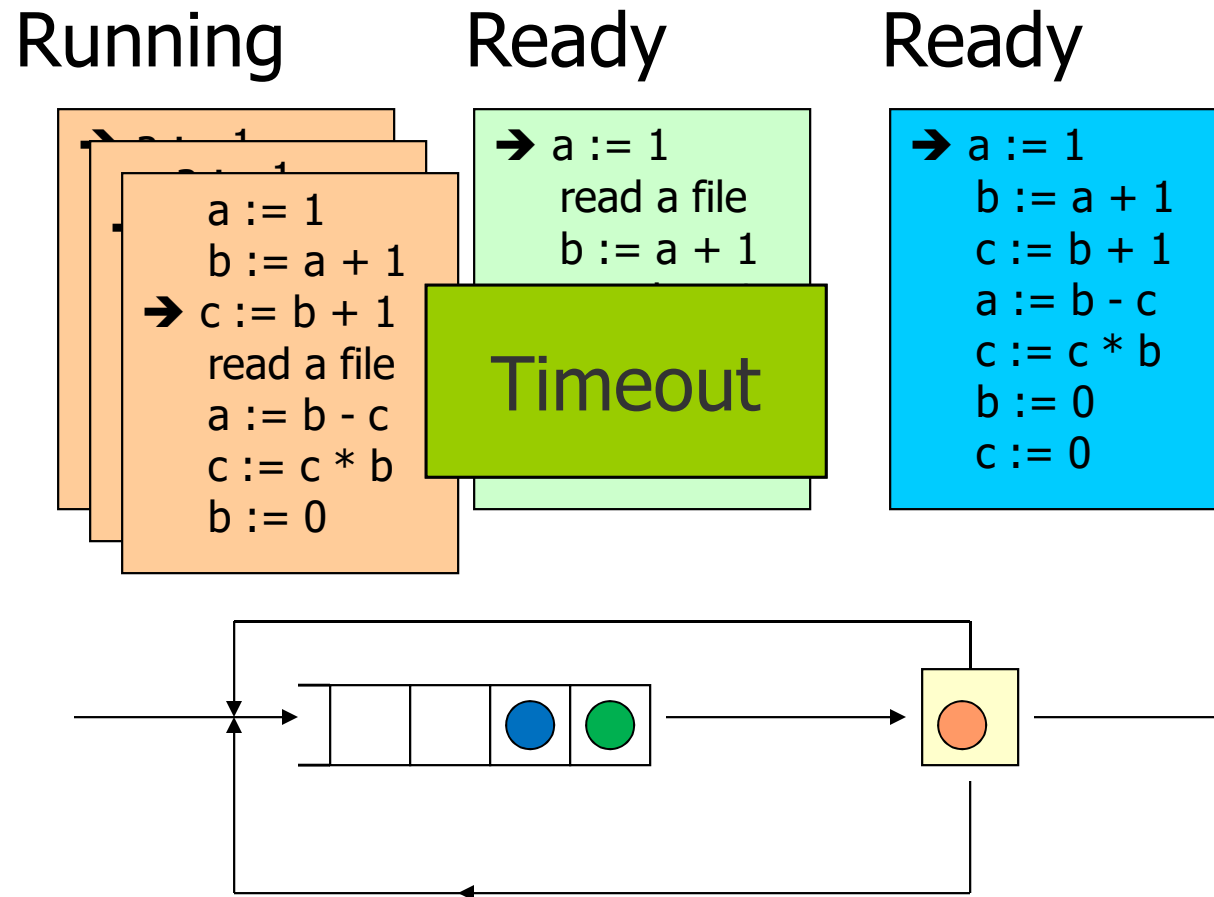
Ready

```
→ a := 1
  b := a + 1
  c := b + 1
  a := b - c
  c := c * b
  b := 0
  c := 0
```

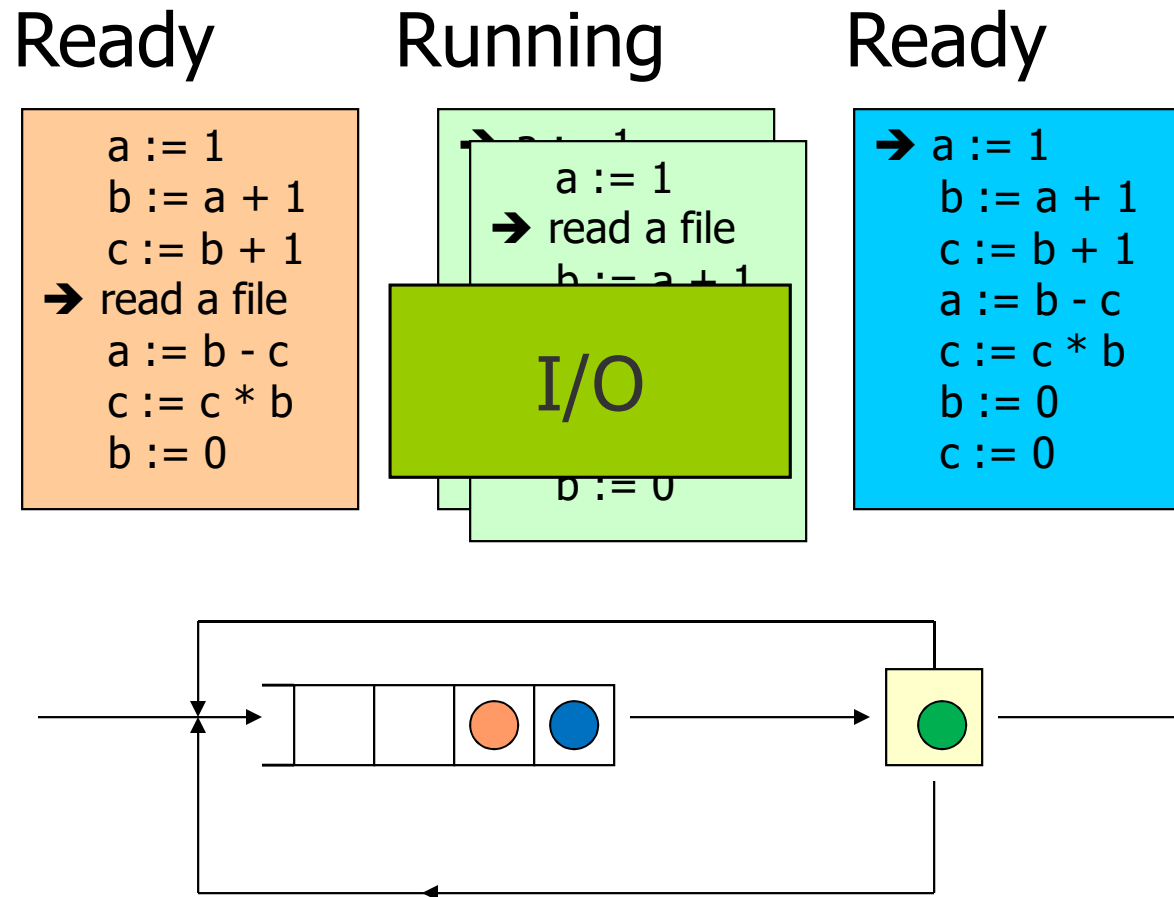


2-Process

# Is Two Process Model Sufficient?

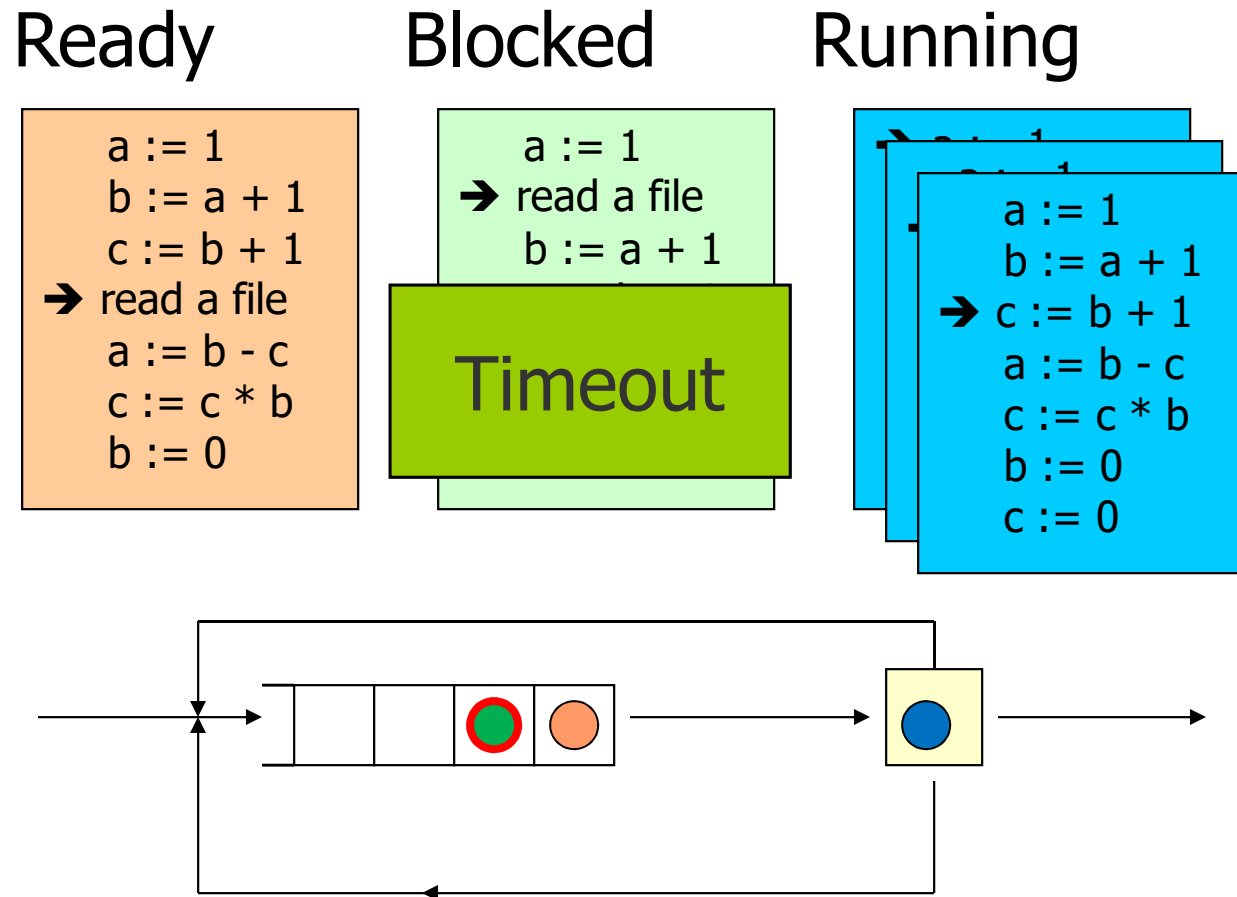


# Is Two Process Model Sufficient?



# Is Two Process Model Sufficient?

---





# Is Two Process Model Sufficient?

Running

```
a := 1
b := a + 1
c := b + 1
→ read a file
a := b - c
c := c * b
b := 0
```


Blocked

```
a := 1
→ read a file
b := a + 1
```

I/O

Ready

```
a := 1
b := a + 1
c := b + 1
→ a := b - c
c := c * b
b := 0
c := 0
```



2-Process

# Is Two Process Model Sufficient?

Blocked


```
a := 1
b := a + 1
c := b + 1
→ read a file
a := b - c
c := c * b
b := 0
```

Blocked

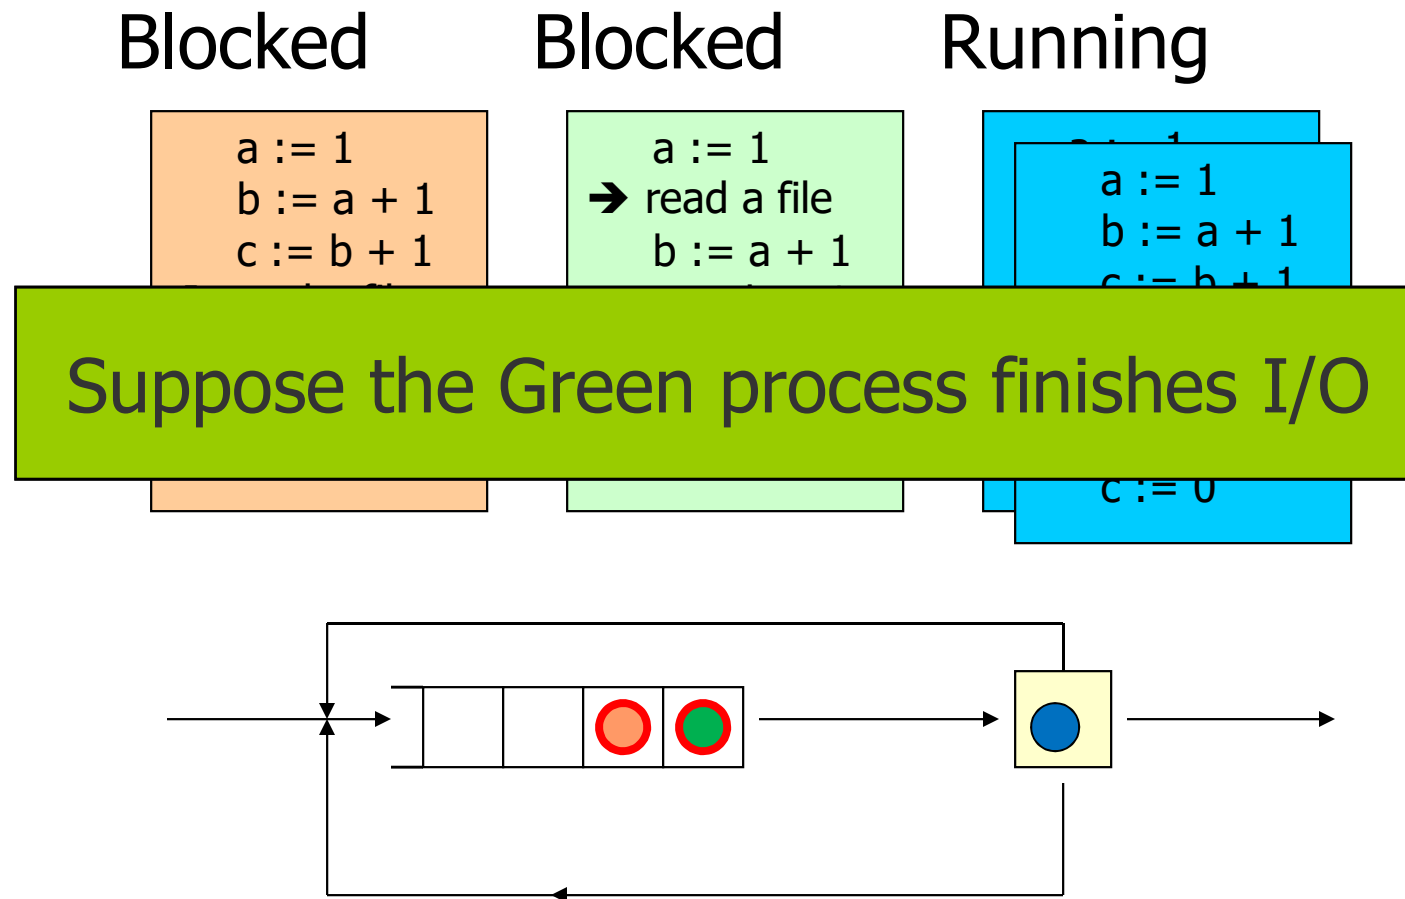
```
a := 1
→ read a file
b := a + 1
c := b + 1
a := b - c
c := c * b
b := 0
```

The Next Process to Run cannot be simply selected from the front

```
c := c + b
b := 0
c := 0
```



# Is Two Process Model Sufficient?



# Is Two Process Model Sufficient?

Blocked

```
a := 1
b := a + 1
c := b + 1
→ read a file
a := b - c
c := c * b
b := 0
```


Ready

```
a := 1
read a file
→ b := a + 1
```

Timeout

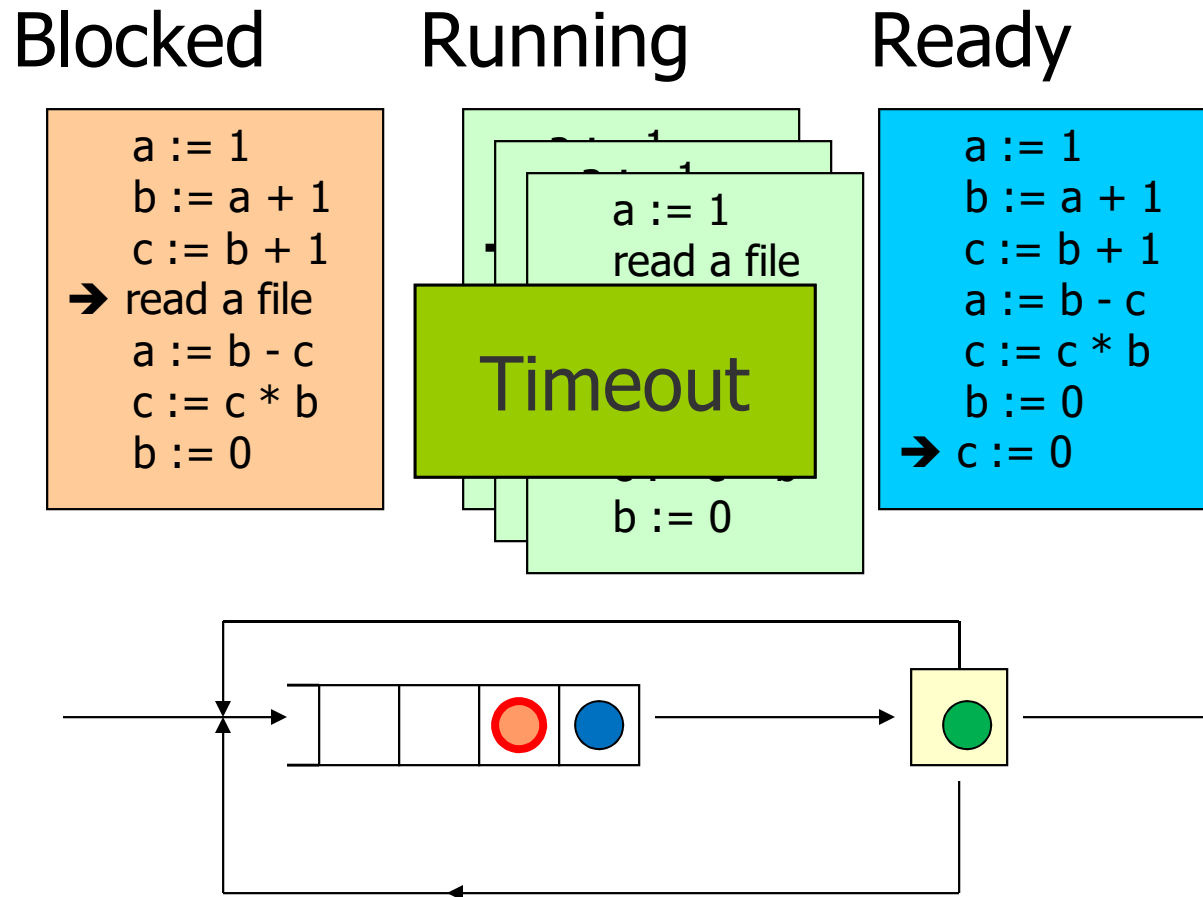
Running

```
a := 1
b := a + 1
c := b + 1
a := b - c
c := c * b
→ b := 0
c := 0
```



2-Process

# Is Two Process Model Sufficient?



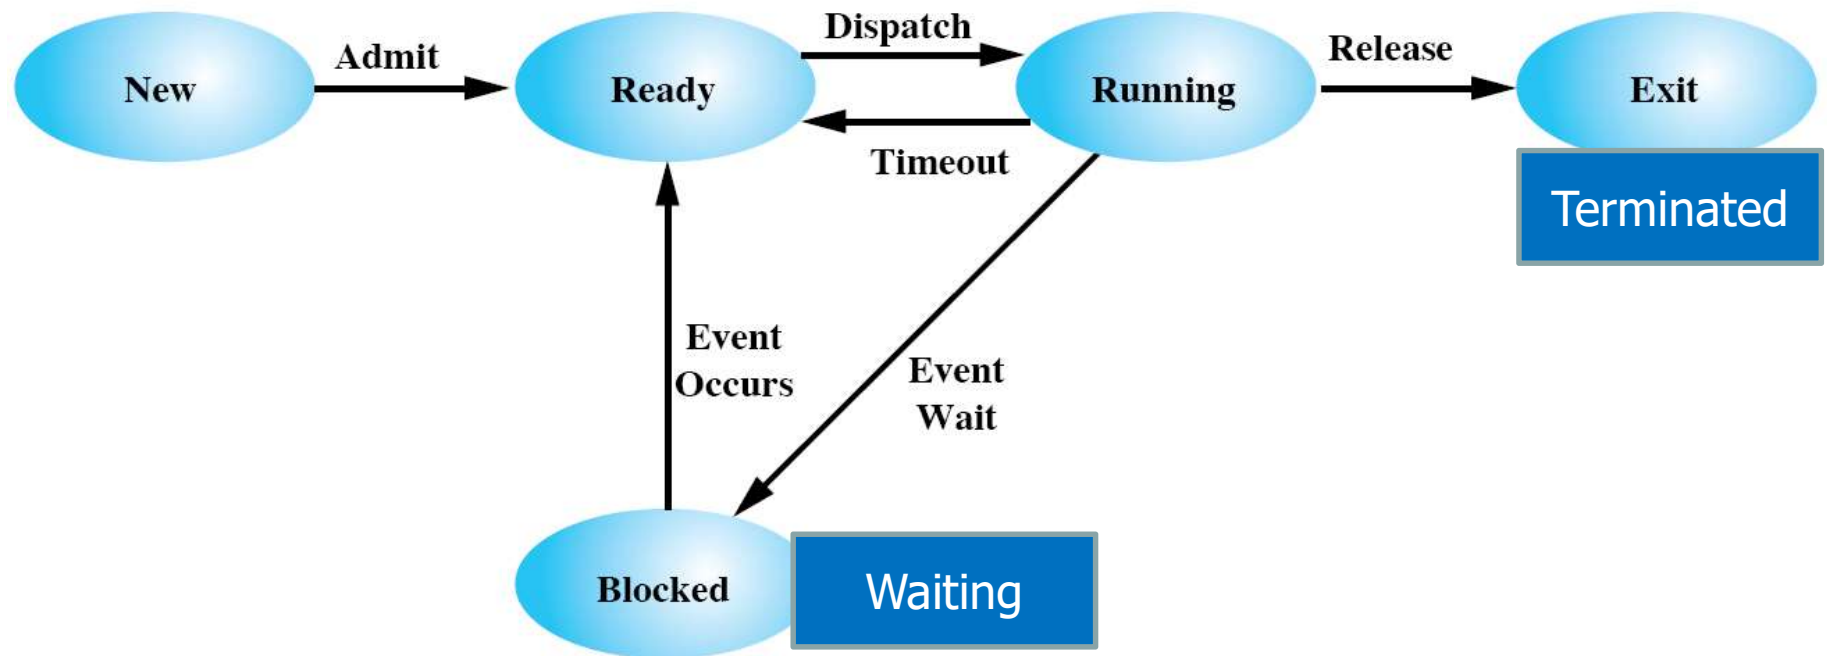
# Limitations of Two Process Model

---

- **Timeouts:** The process with highest priority will be dispatched to the CPU (... regardless of its capabilities!)
- Wasteful in case of
  - Waiting for I/O operations
  - Waiting for access to a resource
  - In this case dispatcher cannot select the process at the front
  - In worst case dispatcher has to scan the whole queue
- Add more states ...
  - Waiting (i.e., blocked), ready

# Five State Model

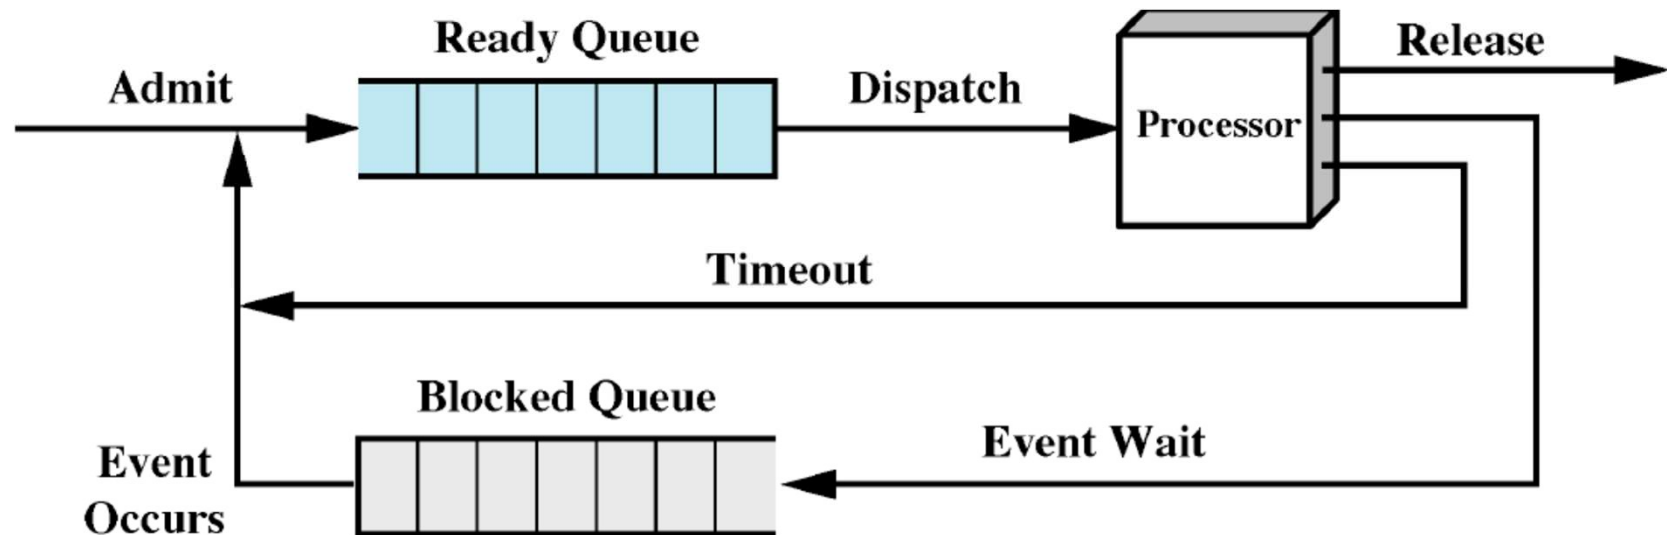
---



- Admittance to control resource usage
- Processes which need to wait for events can be blocked

# How Can the Dispatcher Use the Five States?

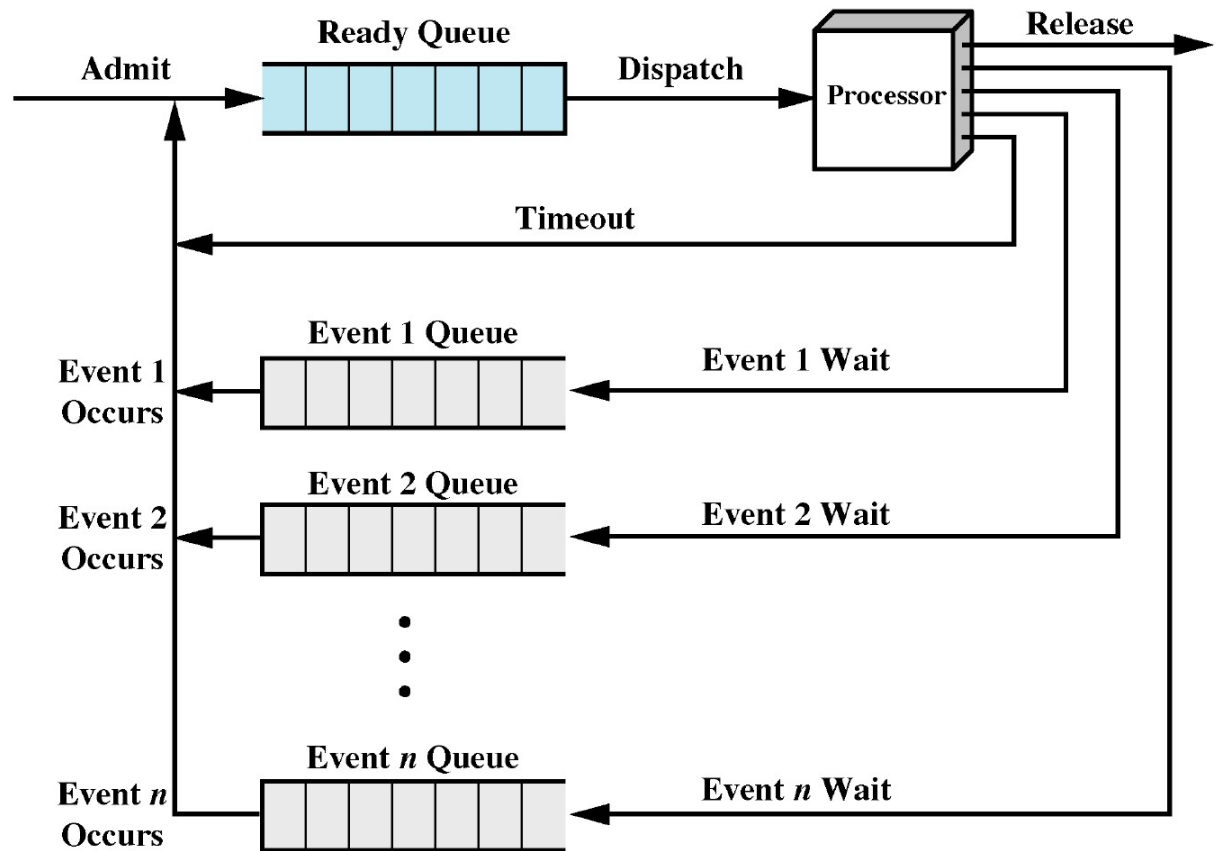
- Maintain different queues
  - One for ready processes
  - One/many for blocked processes
- How expensive is it to transfer a process from blocked to ready state?






# Using Multiple Queues

- Dispatcher can efficiently react to distinct events in the operating system ...
- One queue per event
- $O(1)$  operation to transfer process from **running** to **blocked** state
- $O(1)$  operation to transfer process from **blocked** to **ready** state
- Low dispatch latency



# Using Multiple Queues: Ready And I/O Queues



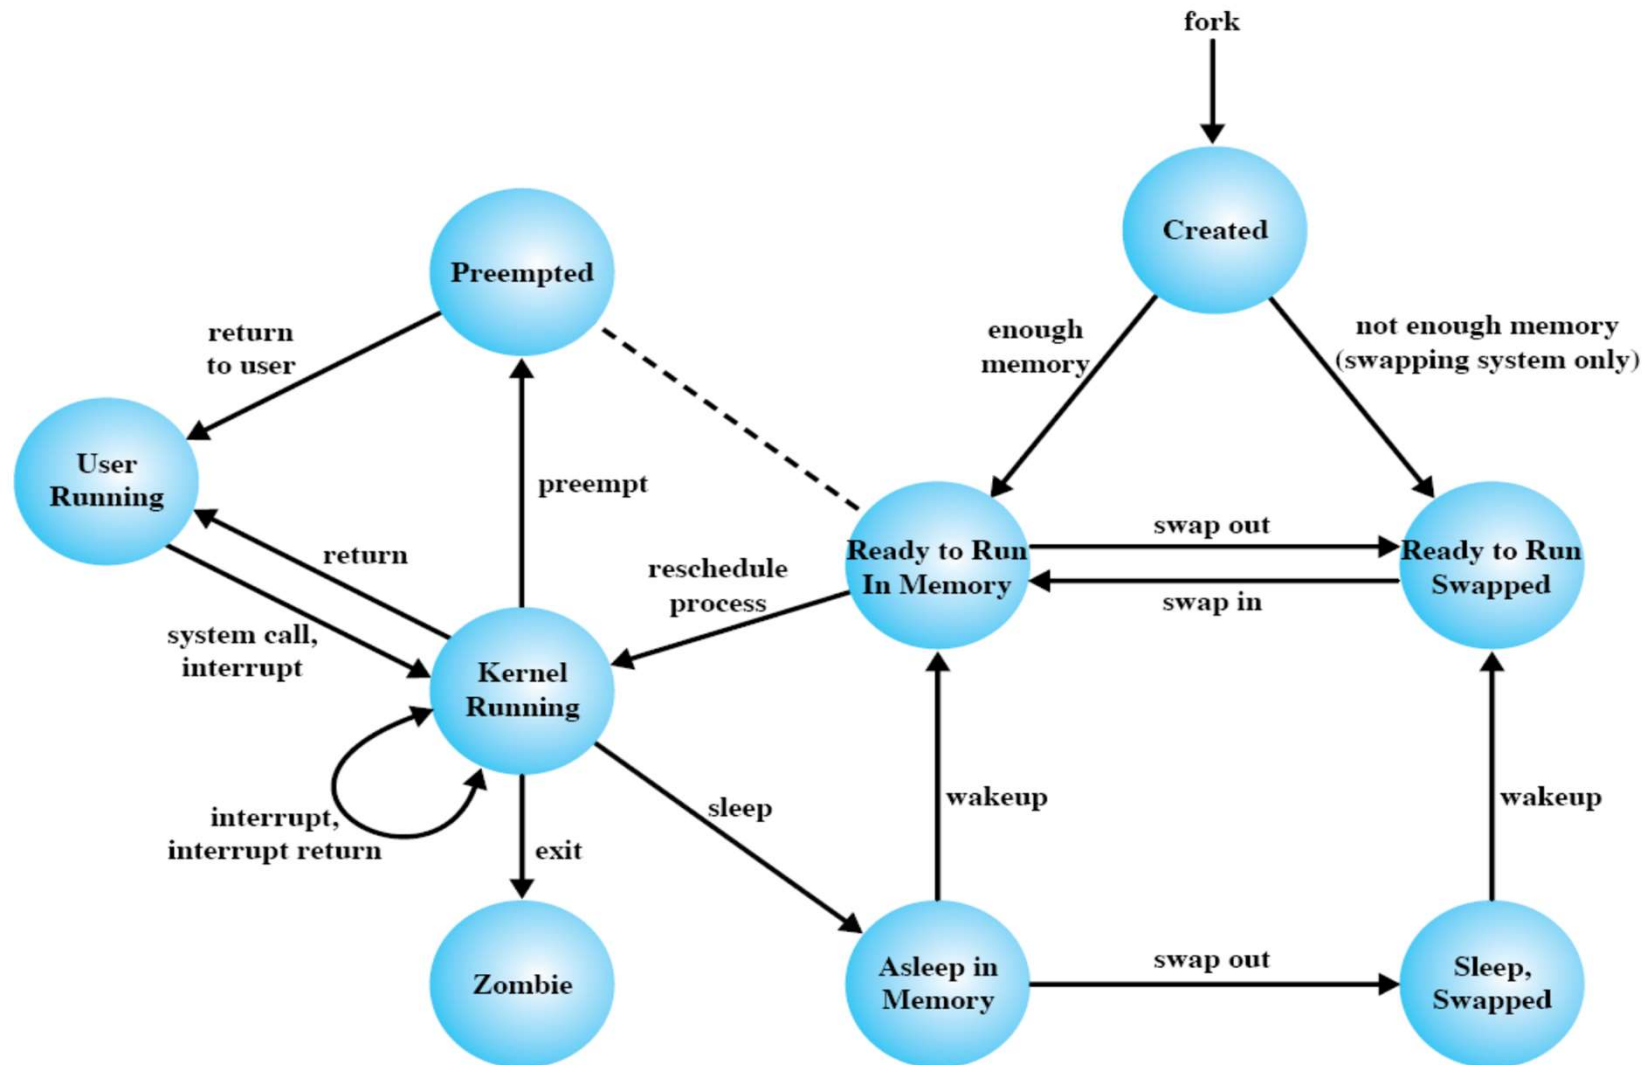
2-Process

# Limitations of the Five State Model

---

- Available resource usage depends on the number of processes deployed
  - For example, the available memory drops with the number of actively running processes
  - Low hit rate in accessing the memory hierarchy and slow performance
- Idea: **Swap** some processes to disk
  - For example, free up more memory and achieve higher hit rate for other processes
- Process is in **suspend state** when swapped to disk
  - Nothing of the process image resides in main memory!

# Unix V Process State Transition

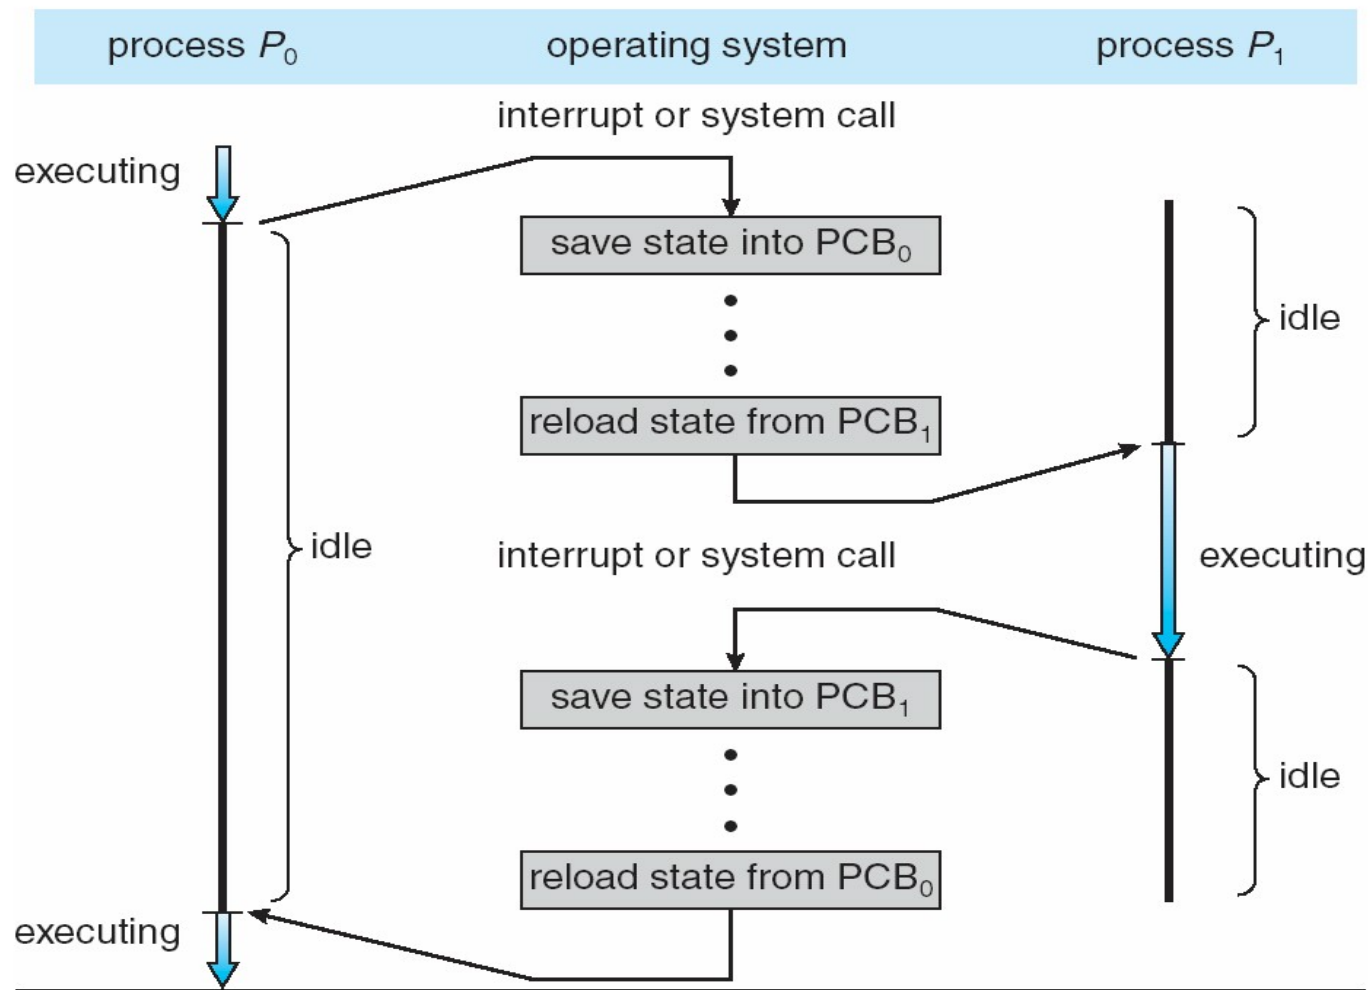


# Context Switching

---

- **Context** of a process represented in the PCB
- When CPU switches to another process, OS must
  - **Save context** of processor including program counter and other registers (in PCB)
  - **Move PCB** to appropriate queue (ready, blocked, ...)
  - Select another process for execution
  - **Update the PCB** of the process selected
  - **Update memory-management** data structures
  - **Restore context** (in processor) of the selected process
- Context switch time is overhead
  - The system does no useful work while switching
  - Time is also dependent on hardware support
    - For example, some hardware provides multiple sets of registers per CPU

# Context Switching

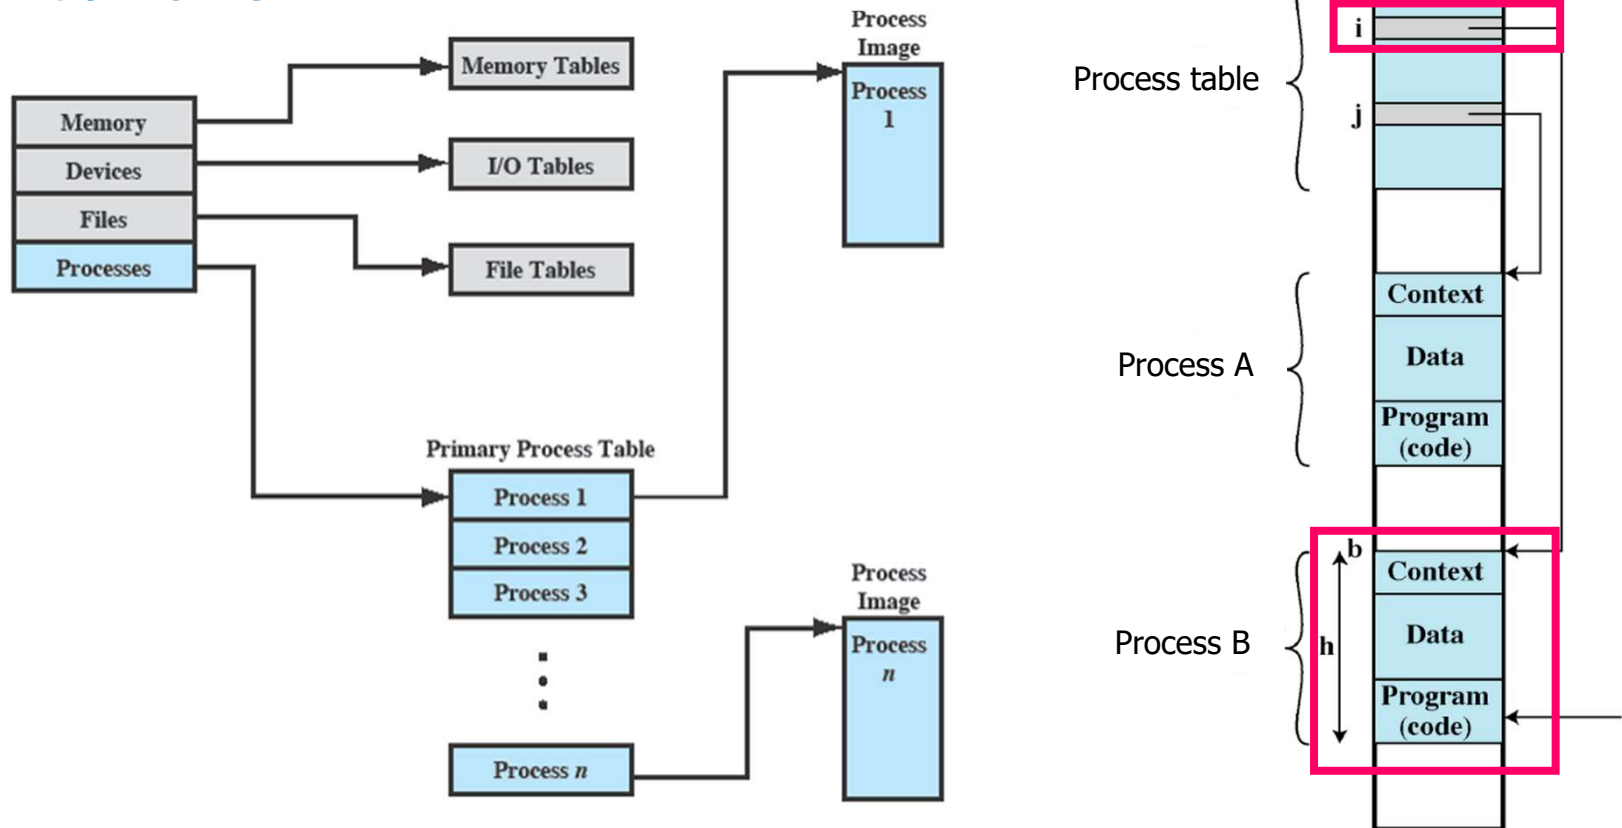


---

# Process Creation

# Process Creation

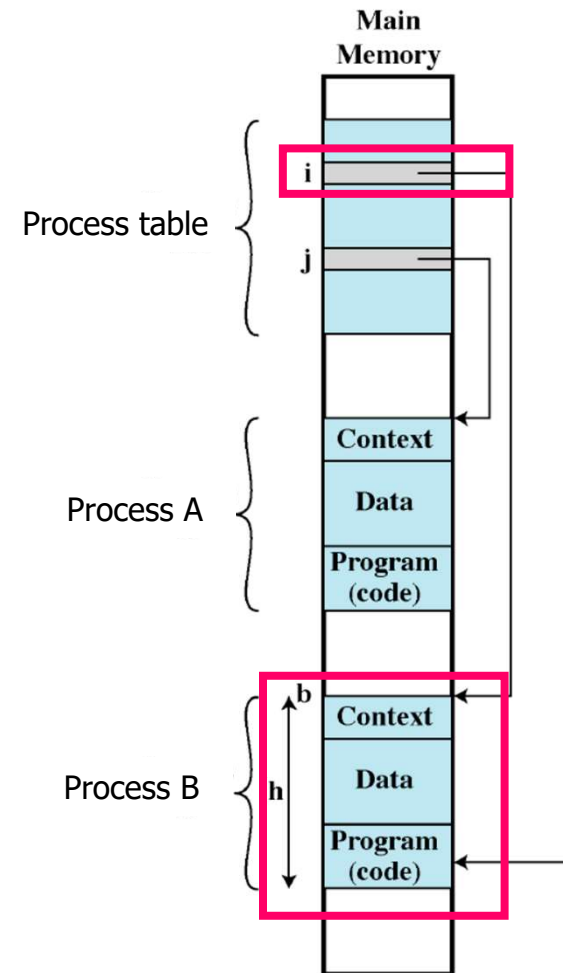
1. Assign a unique identifier (pid) to the new process
2. Allocate space for the process
3. Initialize PCB





# Process Creation

1. Assign a unique identifier (pid) to the new process
2. Allocate space for the process
3. Initialize PCB
4. Setup appropriate linkage
  - E.g., add to ready or ready suspend queue
5. Create or expand other data structures



# Process Creation

---

- Traditionally, the OS created all processes
- But it can be useful to let a running process create another
- This action is called process spawning
  - Parent Process is the original, creating, process
  - Child Process is the new process

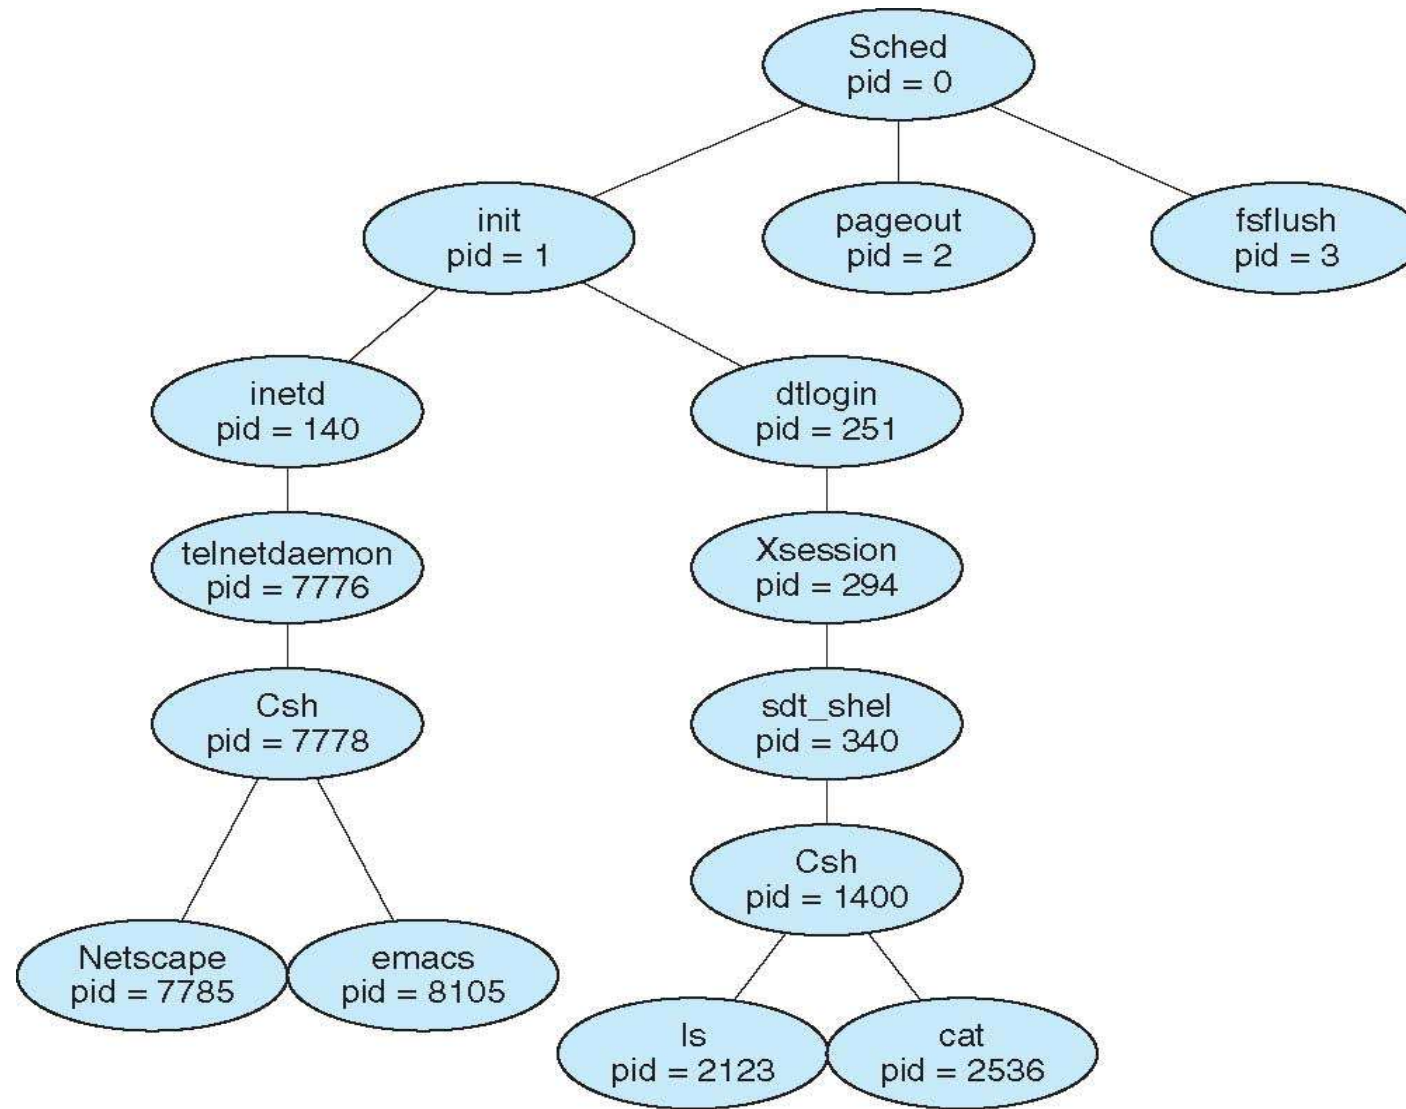
# Process Creation

---

- **Parent** process create **child** processes, which, in turn create other processes, forming a tree of processes
  - Process identified and managed via a **process identifier (pid)**
- **Resource sharing**
  - Parent and children share all resources
  - Children share subset of parent's resources
  - Parent and child share no resources
- **Execution**
  - Parent and children execute concurrently
  - Parent waits until children terminate
- **Address space**
  - Child duplicate of parent
  - Child has a program loaded into it

# Process Tree on Solaris

---



2-Process

# Process Creation in Linux

- **fork** system call creates new process

```
int pid;  
int status = 0;  
  
pid = fork();  
if (pid > 0) {  
    /* parent */  
    .....  
    pid = wait(&status);  
} else {  
    /* child */  
    .....  
    exit(status);  
}
```

**fork** creates an exact copy of the parent process

process ID to the parent

child pid and status.

**wait** variants allow wait on a specific child, or notification

Child process passes status back to parent on **exit**, to report success/failure

# `fork()` System Call

---

- Child process inherits
  - Stack
  - Memory
  - Environment
  - Open file descriptors
  - Current working directory
  - Resource limits
  - Root directory
- Child process does not inherits
  - Process ID and parent process ID
  - Timers and pending signals
  - Resource utilization and CPU times (initialized to zero)
  - Memory and file locks

# Zombie (or Defunct) Process

---

- A process that has **completed execution** but still has an **entry in the process table**
- Process table entry is needed to **allow the parent process to read the child's exit status**
  - Parent can read the status by executing the `wait()` system call
- **No memory allocated to zombie process** except for the process table entry
  - On exit, all of the memory and resources are deallocated
  - However, **process table** can only have **limited number of entries**

# Orphan Process

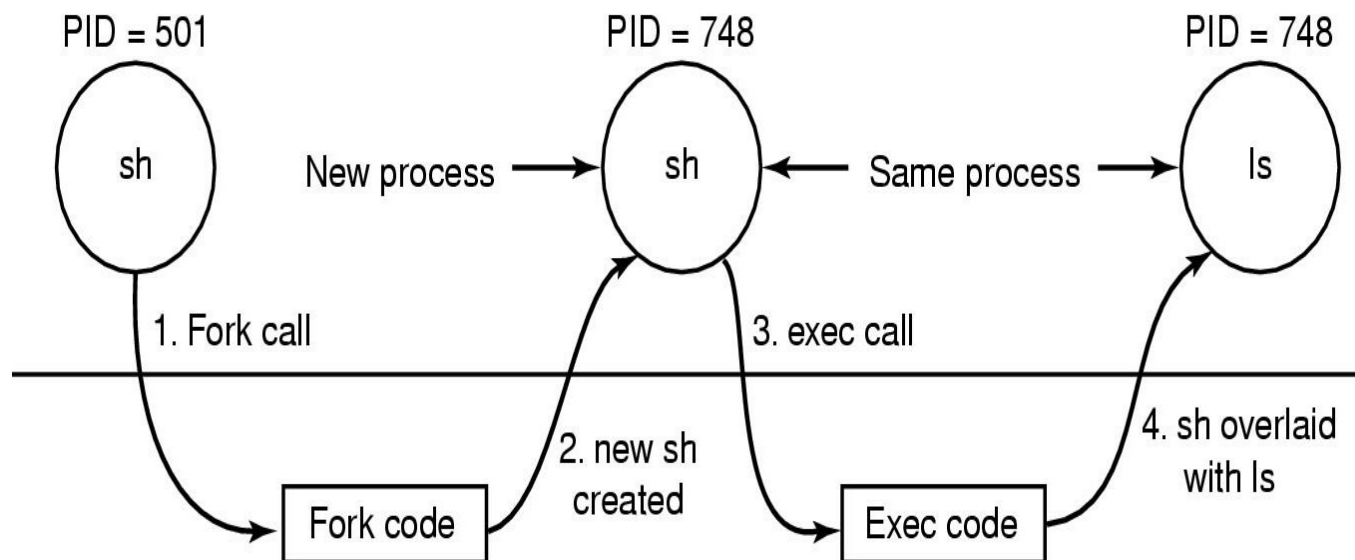
---

- An orphan process is a process that is still executing, but whose parent has terminated
  - If the parent terminates without calling `wait()`, the child is adopted by `init`
- Orphan processes do not become zombie processes
  - `init` periodically executes the `wait()` system call to avoid zombie processes



## exec () System Call

- Enable child process to run other program
- Replaces process's memory space with a new program
  - Loads binary file into memory and starts its execution
- Cannot create new process
  - Typically used after `fork()`



Steps in executing the command **ls** issued to the shell

2-Process

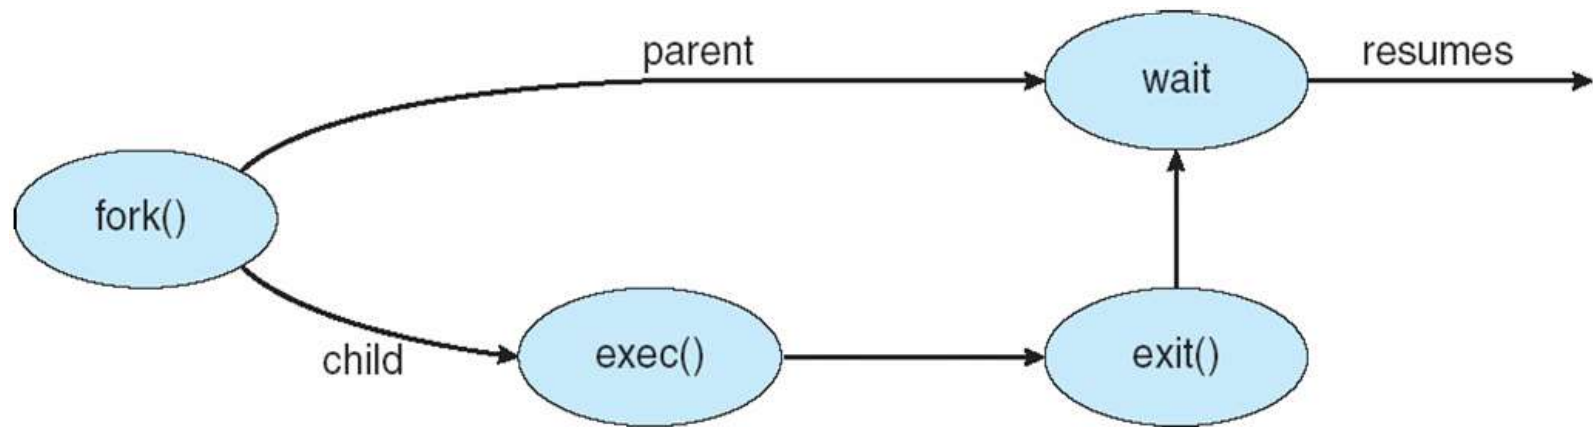
## exec () System Call Example

---

```
int main(int argc, char *argv[])
{
    pid_t cpid;
    cpid = fork();
    if (cpid == -1) {
        perror("fork");
        exit(EXIT_FAILURE);
    }
    if (cpid == 0) {
        /* Child code */
        execlp ("/bin/ls", "ls", NULL);
        /* Why no exit statement*/
    }
    else {
        /* parent code */
        wait(NULL);
        printf("child finished");
        exit(EXIT_SUCCESS);
    }
}
```

## exec () System Call Example

---



---

# Interprocess Communication

# Interprocess Communication

---

- A process has access to the memory which constitutes its own address space
- So far, we have discussed communication mechanisms only during process creation/termination
- When a child process is created, the only way to communicate between a parent and a child process is:
  - The parent receives the exit status of the child
- Processes may need to communicate during their life time

# Interprocess Communication

---

- Processes within a system may be **independent** or **cooperating**
  - Cooperating process can affect or be affected by other processes
- Reasons for processes corporation
  - Information sharing
  - Computation speedup
  - Modularity
  - Convenience
- Cooperating processes need interprocess communication (IPC)
- Two fundamental models of IPC
  - Shared memory
  - Message passing

# Communication Models

---

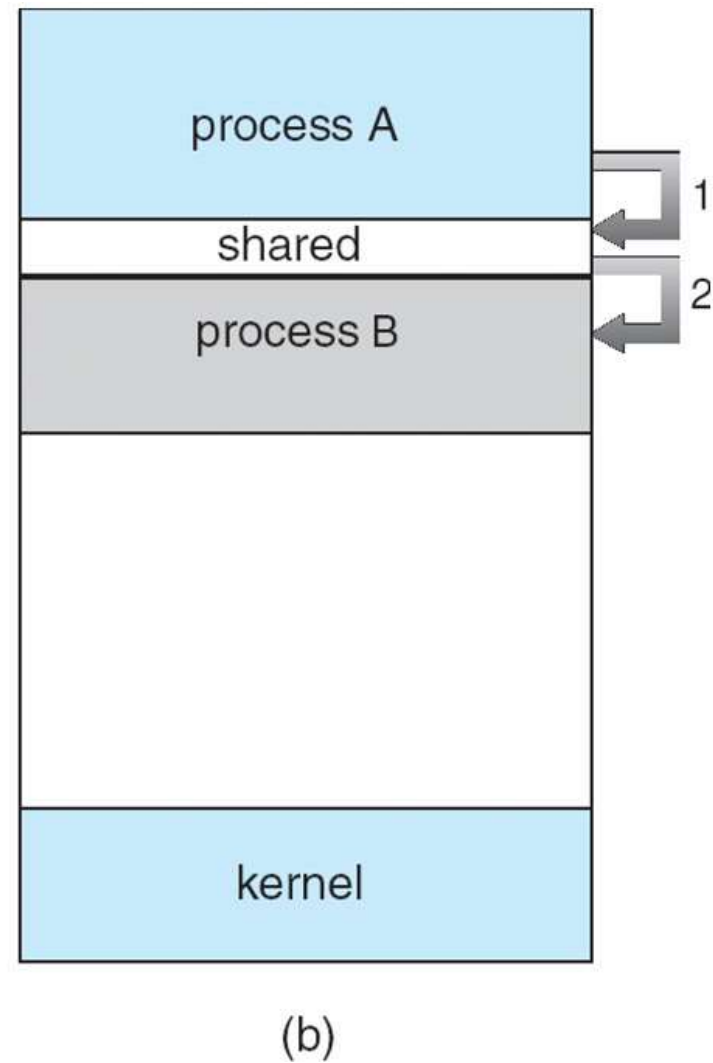
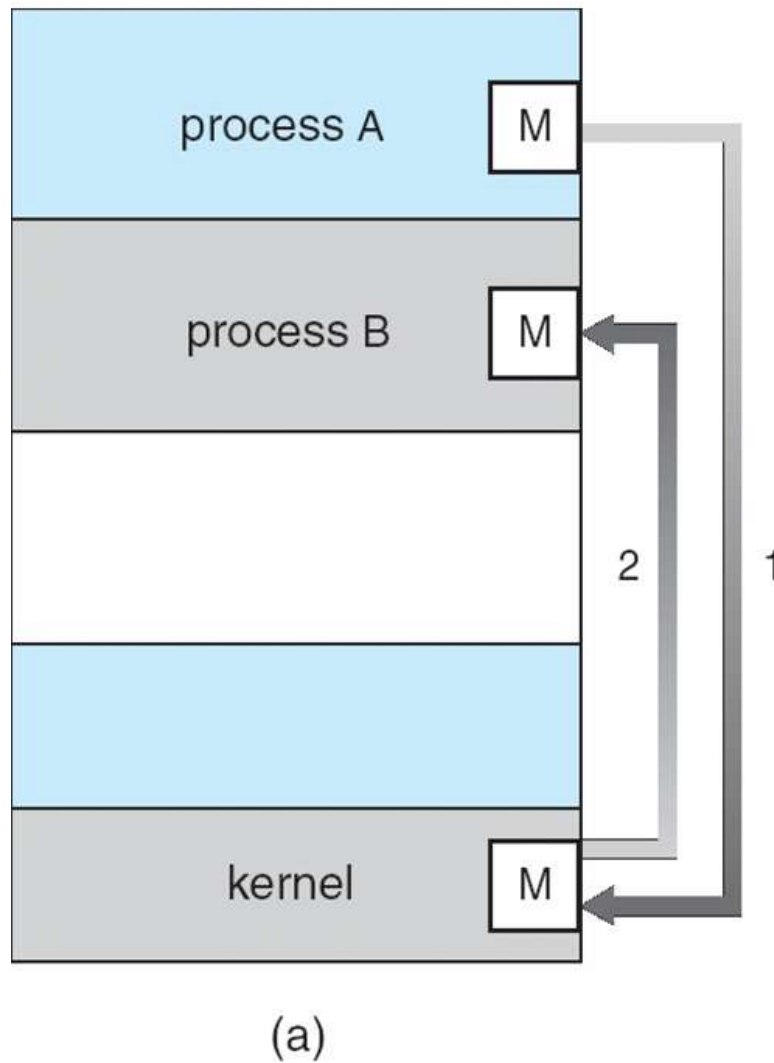
## Shared memory

- Need to establish a region of shared memory
- Usually resides in the address space of a process
- OS system calls are only needed to setup the memory
- Basic operations: `read/write`

## Message passing

- Send messages between processes
- Requires system calls to the OS for every message
- Often easier to realize for networked process communication
- Basic operations: `send/receive`

# Communication Models






# Shared Memory

---

- Usually resides in the address space of the process creating shared memory
- Require processes to coordinate their processing
  - Results of read/write not guaranteed to be deterministic
  - Concurrent writes to the same address
  - Result depends on the order of operations



# Shared Memory – POSIX API

---

- Process first creates shared memory segment
  - `id = shmget(IPC_PRIVATE, size, S_IRUSR | S_IWUSR);`
- Process wanting access to that shared memory must attach to it
  - `shared_memory = (char *) shmat(id, NULL, 0);`
- Now the process could write to the shared memory
  - `sprintf(shared_memory, "Writing to shared memory");`
- When done a process can detach the shared memory from its address space
  - `shmdt(shared_memory);`

# Message Passing

---

- No sharing of resources between processes needed
- Basic operations
  - `send({destination}, message)`: send a message
  - `receive({source}, message)`: receive a message
- If P and Q wish to communicate, they need to:
  - Establish a communication link between them
  - Exchange messages via send/receive
- Implementation of communication link
  - Physical (e.g., shared memory, hardware bus)
  - Logical (e.g., logical properties)

# Implementation Questions

---

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?

# Direct Communication

---

## Direct symmetric communication

- Each process that wants to communicate must name the recipient
  - `send(P, message)`
    - Send a message to process P
  - `receive(Q, message)`
    - Receive a message from process Q
- Communication link established automatically
- Link associated exactly between two processes

## Direct asymmetric communication

- Use instead
  - `send(P, message)`
    - Send a message to process P
  - `receive(message)`
    - Receive message from arbitrary process

# Indirect Communication

---

- Problem of Direct Addressing
  - Changing a process identifier may impact all other process definitions
  - Limited modularity
- Indirect addressing uses mailbox or port concept
  - Mailbox: Shared by multiple receiver
  - Port: Receiver Specific, e.g., WebServer uses port 80
  - `send(A, message)`
    - Send a message to mailbox/port named A
  - `receive(A, message)`
    - Receive a message from mailbox A
  - Link is only established if both members share a mailbox
  - A link may be associated with more than 2 processes
  - Each pair of processes may share several communication links

## One-to-One

---


- Private communication link between two processes



# Many-to-One

---

- Client/server interactions






# One-to-Many

---


- Multicast information from a source to a set of receivers



# Many-to-Many

---

- Allows for many-to-many communication



# Indirect Communication – Example

---

- Mailbox sharing
  - P1, P2, and P3 share mailbox A
  - P1, sends; P2 and P3 receive
  - Who gets the message?
- Possible solutions
  - Allow a link to be associated with at most two processes
  - Allow only one process at a time to execute a receive operation
  - Allow the system to select arbitrarily the receiver
    - Sender is notified of the receiver

# Indirect Communication – Ownership

---

- Could be either process or Operating System
- Case process is the owner of a mailbox
  - Owner performs as the receiver of messages
  - User sends messages to the mailbox
    - Mailbox destroyed with termination of the process
- Case OS is the owner
  - Need support for processes to
    - Create
    - Send and receive messages through the mailbox
    - Delete a mailbox
  - Manage ownership
  - Possibly multiple owners/ receivers

# Indirect Communication – Ownership

---

- Process owns (i.e. mailbox is implemented in user space)
  - Only the owner may receive messages through this mailbox
  - Other processes may only send
  - When process terminates any “owned” mailboxes are destroyed
- Kernel owns
  - Kernel provides mechanisms to create, delete, send and receive through mailboxes
  - Mailbox has existence of its own independent of any process
  - Process that creates mailbox owns it (and so may receive through it)
  - Process may transfer ownership to another process

# Synchronization

---

- Message passing may be either **blocking** or **non-blocking**
- **Blocking** is considered **synchronous**
  - Blocking send has the sender block until the message is received
  - Blocking receive has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
  - Non-blocking send has the sender send the message and continue
  - Non-blocking receive has the receiver receive a valid message or null
- Combinations of blocking and non blocking calls
  - Blocking send + blocking receive
    - Tight coupling between processes
  - Non blocking send + blocking receive

# Buffering

---

- All messaging system require framework to temporarily buffer messages (i.e., Queues)
- Zero capacity
  - No messages may be queued within the link
  - Requires sender to block until receiver retrieves message
- Bounded capacity
  - Link has finite buffer capacity to hold messages
  - If link is full then sender must block until one is freed up
- Unbounded capacity
  - Link has unlimited buffer space
  - Send never needs to block

---

## IPC Case Study: Unix Pipes



# Process Creation Recap

---

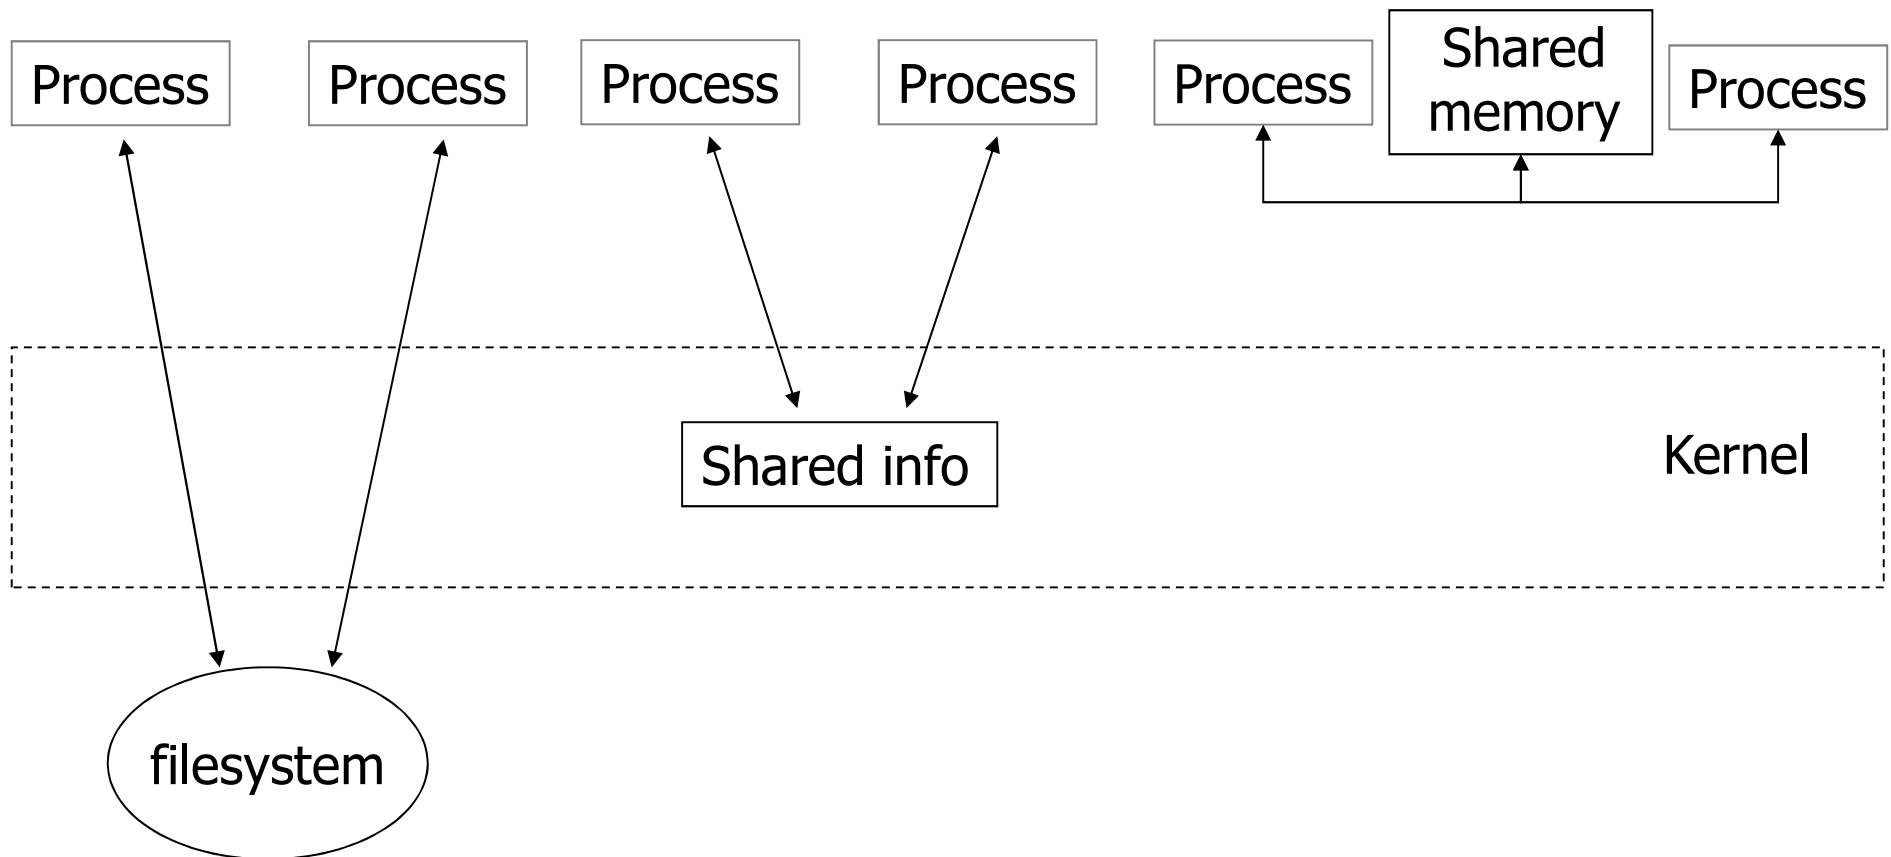
- Consider the following code

```
for(int i = 0; i < 4; i++) {  
    fork();  
}
```

- How many process would be created?
- Represent the processes in form of tree

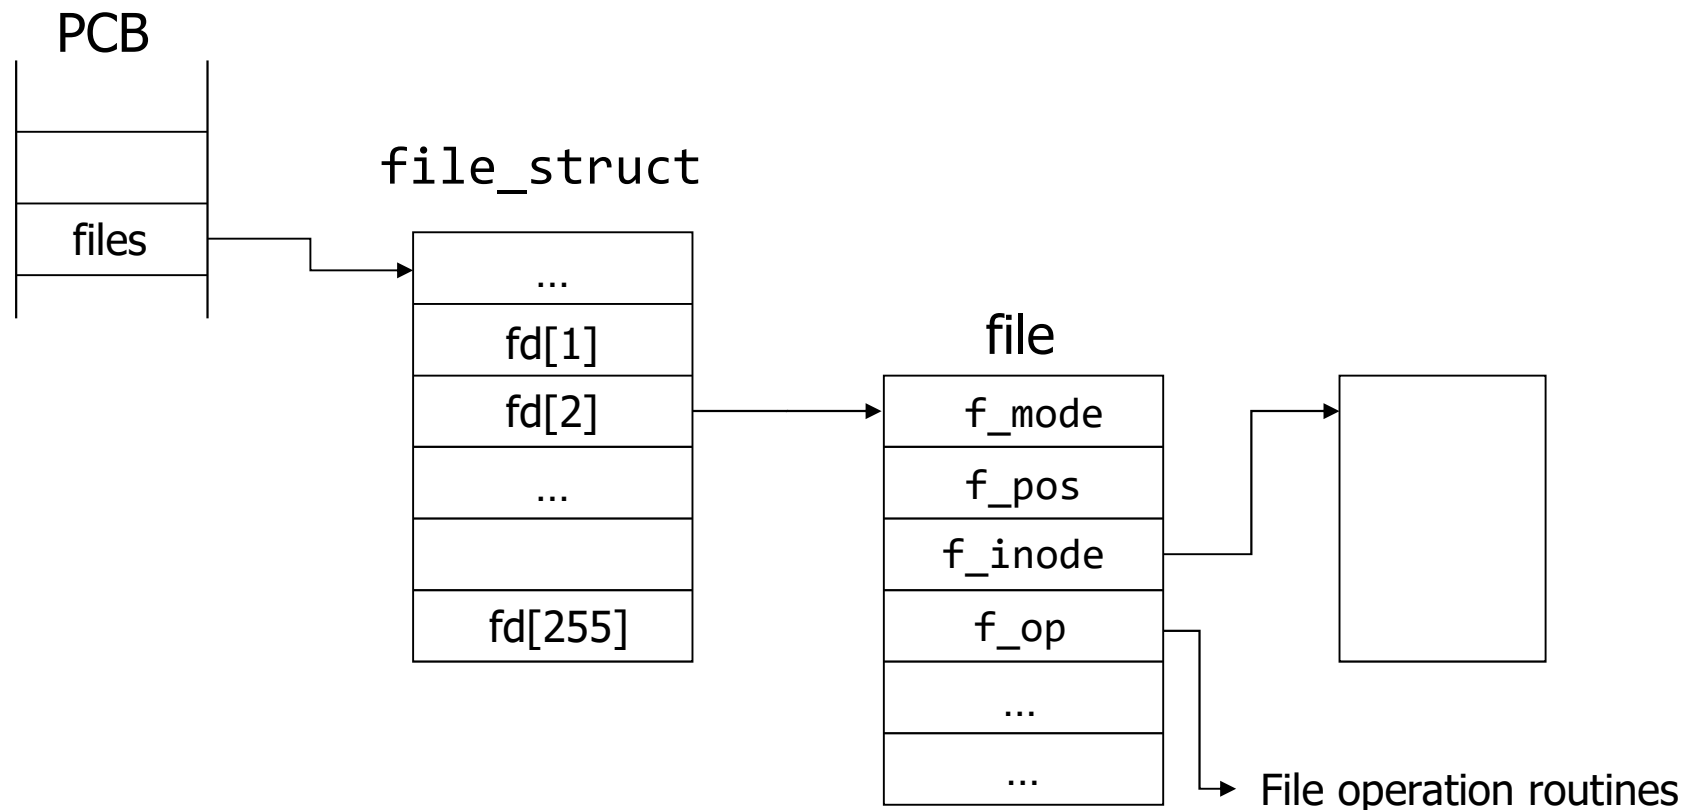
# Unix IPC

---



# File Descriptors


- PCB of each process keeps track of open files
- Pointer to `file_struct`, a kernel-resident array data structure containing the details of open files



# File Descriptors

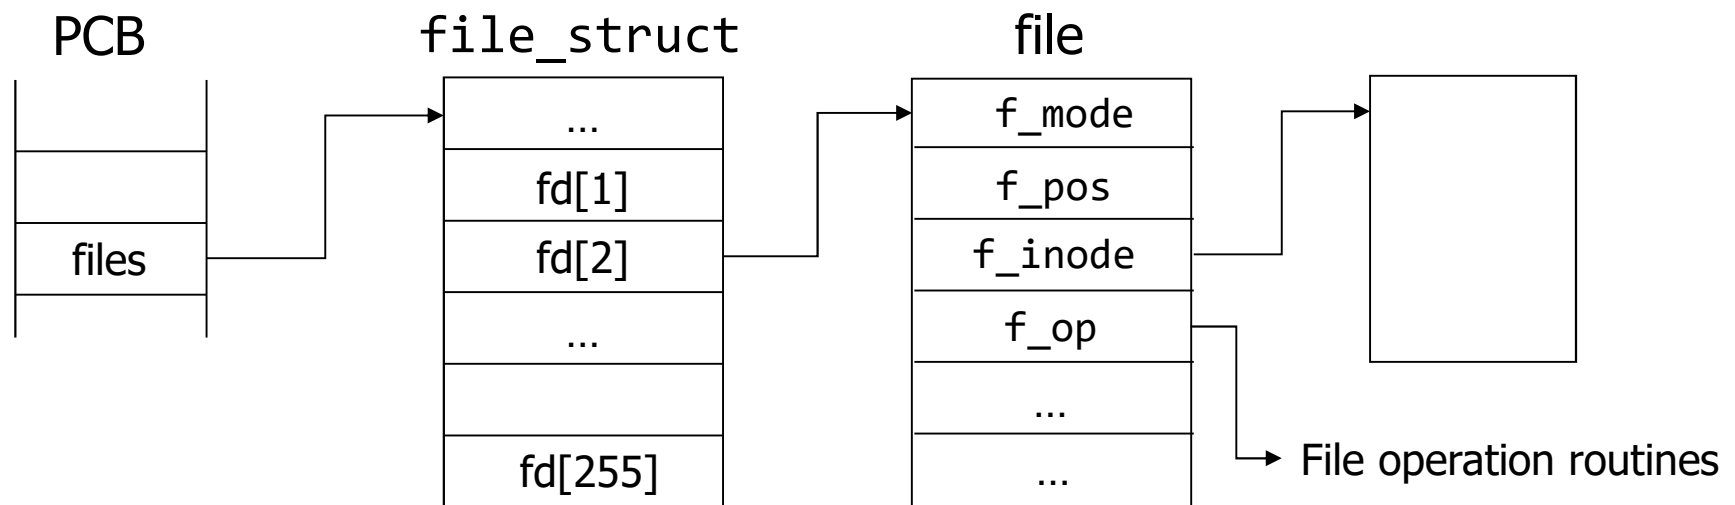
---

- `files_struct` contains pointers to file data structures
  - Each one describes a file being used by this process
- `f_mode`: Describes file mode, read only, read and write or write only
- `f_pos`: Holds the position in the file where the next read or write operation will occur
- `f_inode`: Points at the actual file



# File Descriptors

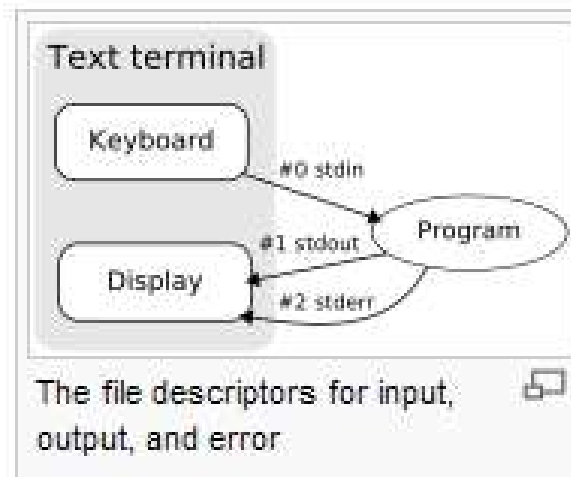
- When a process opens a file, one of the free file pointers in the `files_struct` is used to point to the new file structure
- When Unix processes do any sort of **I/O**, they do it by **reading or writing to a file descriptor**
  - A file descriptor is simply an integer associated with an open file
- All **accesses to files** are via **standard system calls** which pass or return file descriptors



# Standard Input, Output, Error

---

- Linux processes expect three file descriptors to be open when they start
  - Standard input: File descriptor 0 (`stdin`)
  - Standard output: File descriptor 1 (`stdout`)
  - Standard error: File descriptor 2 (`stderr`)
- These three are usually inherited from the creating parent process



# Standard Input, Output, Error: Example

---

- Read from standard input (by default it is keyboard)
  - `char buffer[10];`
  - `read(0,buffer,5);`
- Write to standard output (by default is is monitor))
  - `char buffer[10];`
  - `write(1,buffer,5);`
- By changing the file descriptors we can write to files
- `fread/fwrite` etc. are wrappers around the above `read/write` functions

# Unix Fact

---

- Everything in Unix is a file
- A file in Unix can be
  - A network connection
  - A FIFO queue
  - A pipe
  - A terminal
  - A real on-the-disk file
  - Or just about anything else



# Pipes

---


- Pipes represent a channel for Interprocess Communication
  - Provides a one-way flow of data
  - Can be thought as a special file that can store a limited amount of data in a first-in-first-out manner, exactly akin to a queue



# Pipes

---

- Shared info in kernel's memory

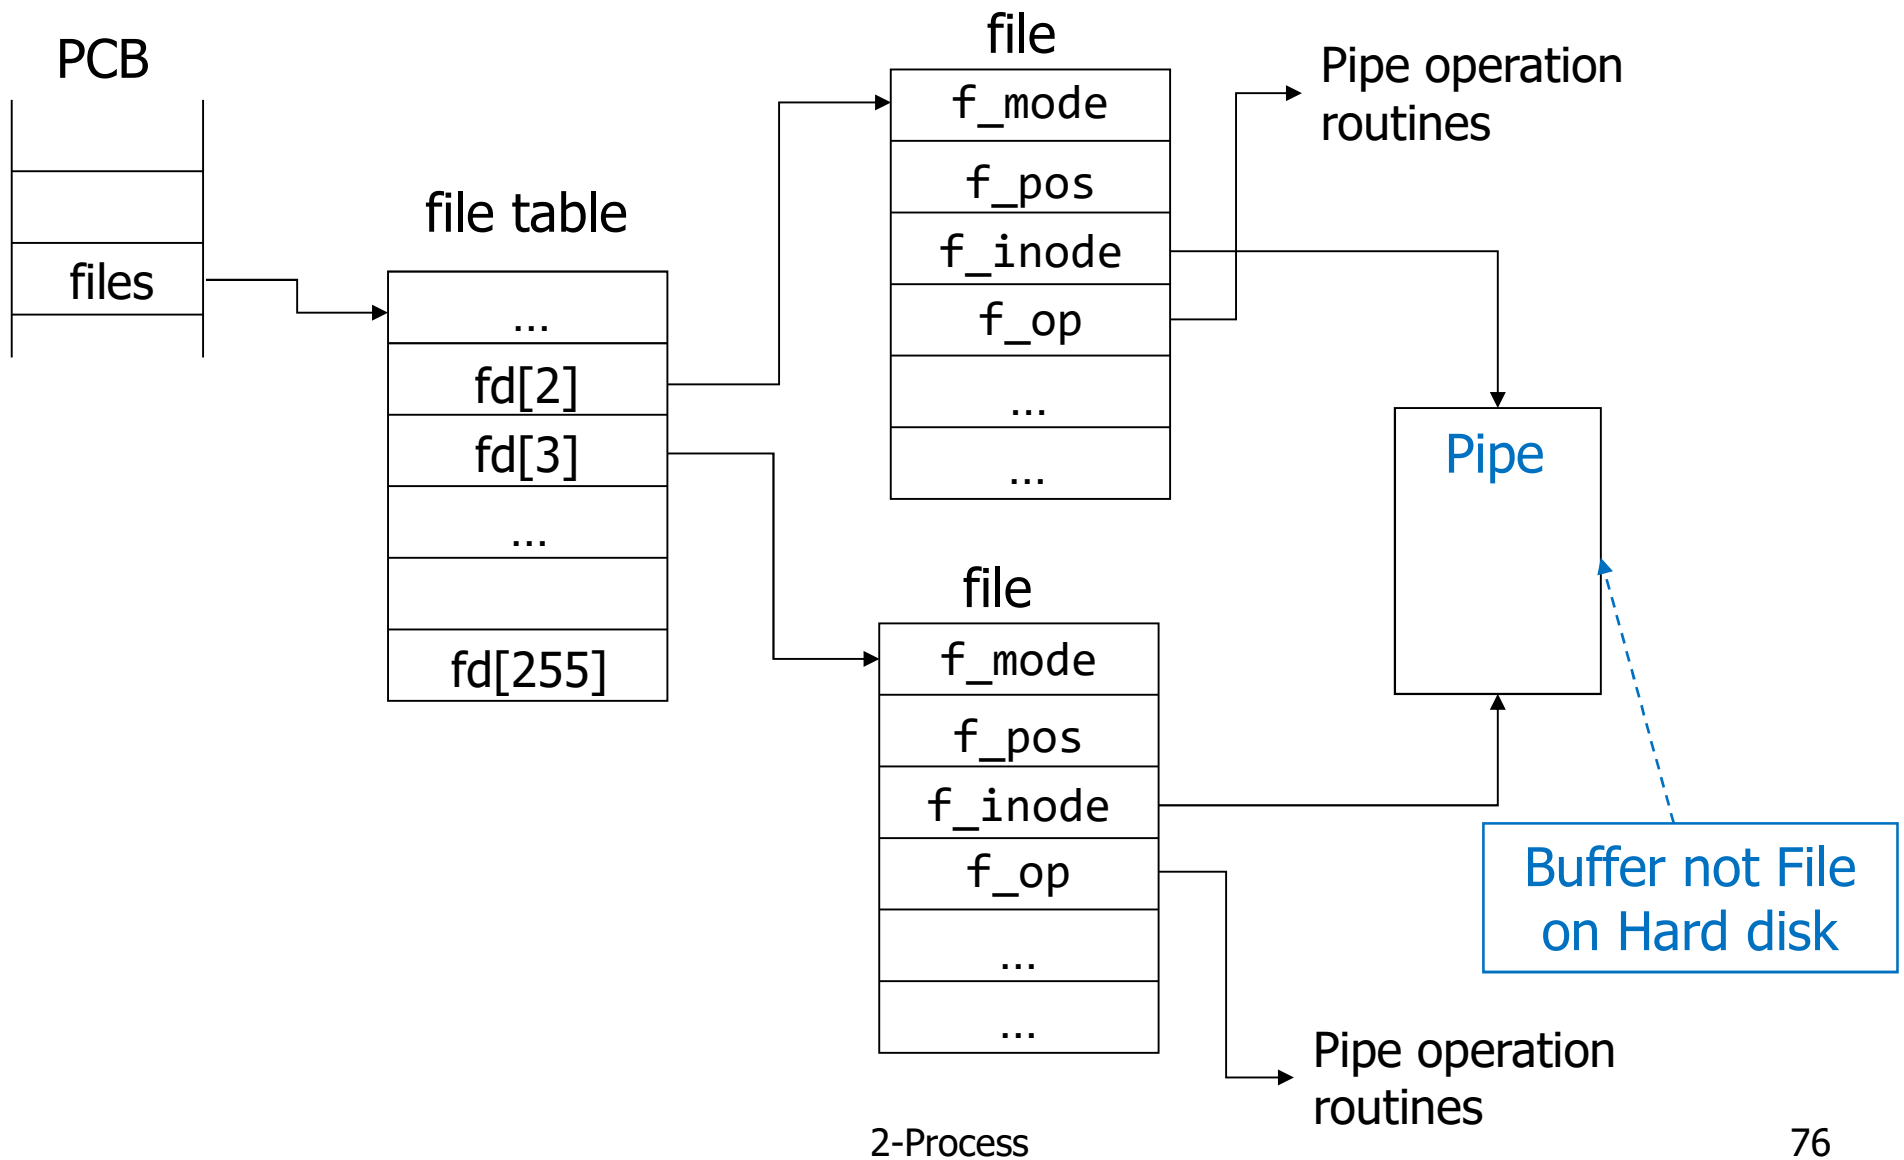


# Pipes

---

- A pipe is implemented using two file data structures which both point at the same temporary data node
- This hides the underlying differences from the generic system calls which read and write to ordinary files
- Thus, reading/writing to a pipe is similar to reading/writing to a file

# Pipes



# (Unnamed) Pipe Creation

---

Two methods for creating (unnamed) pipes

- **pipe system call**

```
#include <unistd.h>
int pipe(int filedес[2]);
```

- Creates a pair of file descriptors pointing to a pipe inode
- Places them in the array pointed to by `filedes`
  - `filedes[0]` is for reading
  - `filedes[1]` is for writing
- Return value: Success returns zero; error returns -1

- **popen system call**


```
FILE *popen(const char *command, const char *type);
FILE* file = popen("ntpdate", "r");
```

- Opens a process by creating a pipe, forking, and invoking the shell

# Pipe Creation


---

```
int main()
{
    int pfd[2];
    if (pipe(pfd) == -1) {
        perror("pipe");
        exit(1);
    }
}
```



# Pipe Example

```
int main()
{
    int pfd[2];
    if (pipe(pfd) == -1) {
        perror("pipe");
        exit(1);
    }
    printf("writing to file descriptor #%d\n", pfd[1]);
    write(pfd[1], "test", 5);
    printf("reading from file descriptor #%d\n", pfd[0]);
    read(pfd[0], buf, 5);
    printf("read %s\n", buf);
}
```



# A Channel Between Parent and Child

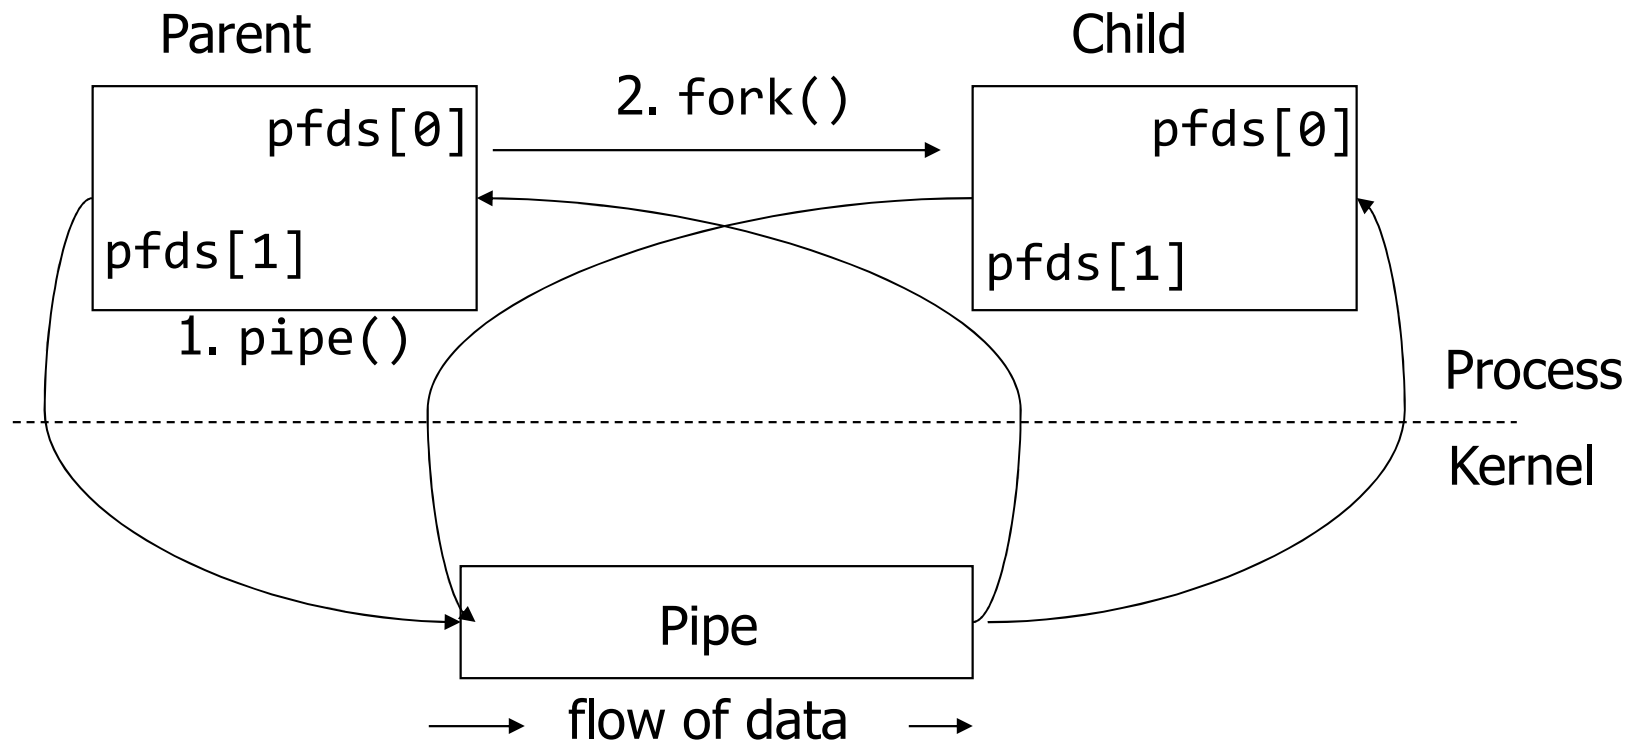
---

- Child is created by a `fork()` call executed by the parent
  - Child process is an image of the parent process
  - All the file descriptors that are opened by the parent are available in the child
- Pipe is inherited by the child
  - File descriptors refer to the same I/O entity
  - Pipe may be passed on to the grand-children by the child process or other children by the parent



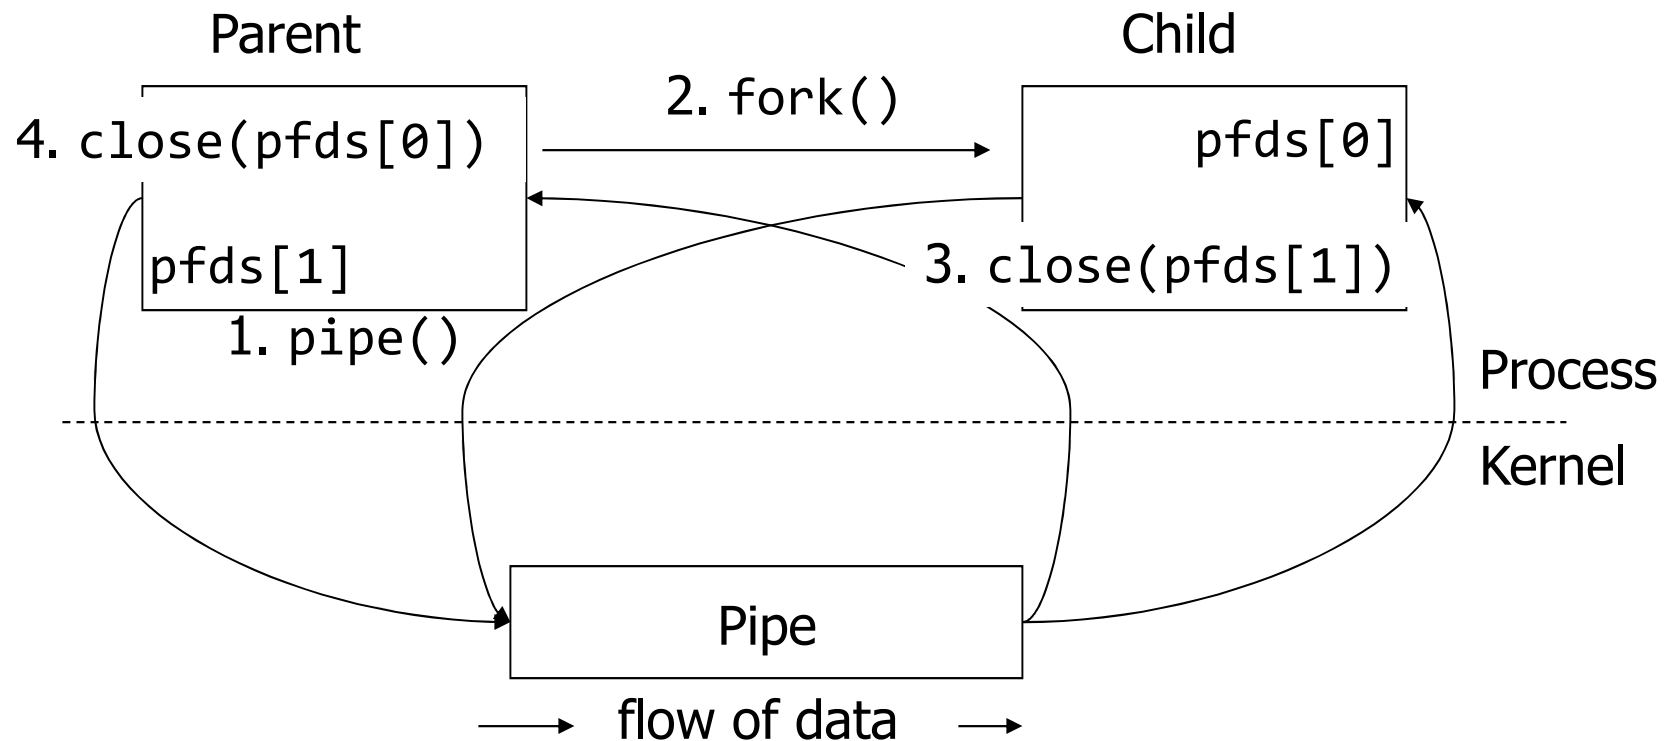
# Piping Between Parent and Child

---



# Piping Between Parent and Child

- To allow one way communication each process should close one end of the pipe



## Pipe Closing

---

- The file descriptors associated with a pipe can be closed with the `close(fd)` system call
- A pipe exists until both file descriptors are closed in all processes
- How would we achieve two way communication?

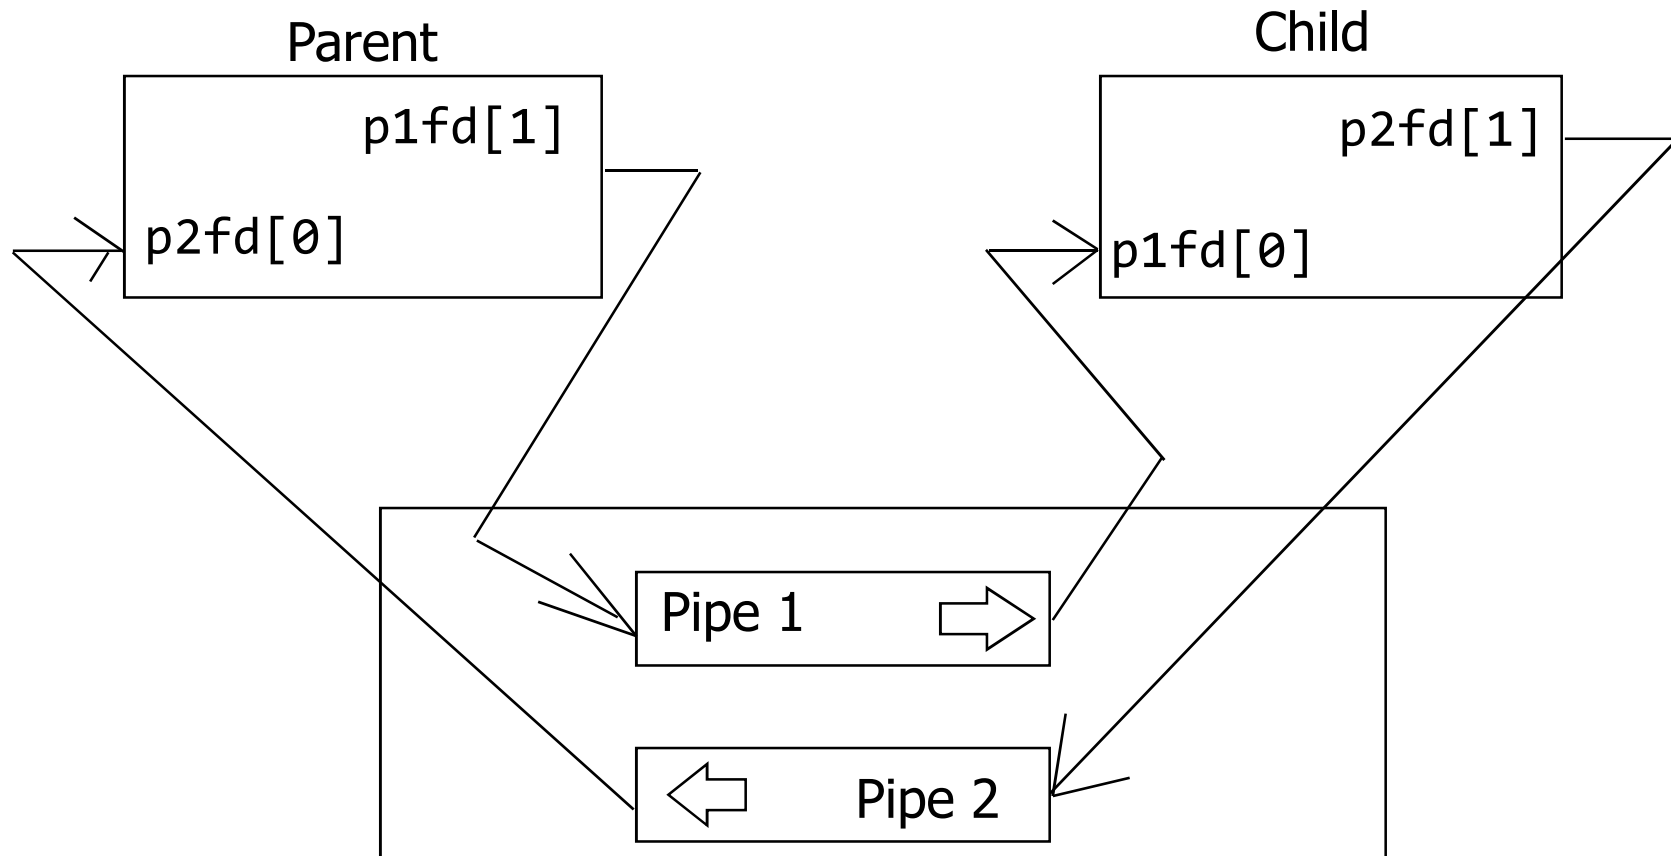
# Pipe Example

---

```
int main() {
    int pfd[2]; char buf[30];
    pipe(pfd);
    if (!fork()) {
        close(pfd[0]);
        printf(" CHILD: writing to the pipe\n");
        write(pfd[1], "test", 5);
        printf(" CHILD: exiting\n");
        exit(0);
    }
    else {
        close(pfd[1]);
        printf("PARENT: reading from pipe\n");
        read(pfd[0], buf, 5);
        printf("PARENT: read \"%s\"\n", buf);
        wait(NULL);
    }
}
```

# Full Duplex Communication via Two Pipes

---



# Named vs. Unnamed Pipes

---

## Unnamed pipe

- Unnamed pipes can only be used between related process, such as parent/child, or child/child process
- Unnamed pipes can exist only as long as the processes using them

## Named pipe

- When created, named pipes have a directory entry
  - Have file access permissions for unrelated processes to use the pipe
- Named pipes can be created by using `mkfifo` system call
  - `int mkfifo(const char *pathname, mode_t mode);`
  - Makes a FIFO special file with name `pathname` and FIFO's permissions
- Any process can open FIFO special file for reading or writing
  - Opening a FIFO for reading normally blocks until some other process opens the same FIFO for writing, and vice versa

# Named Pipe Example

## Process 1

```
int main() {
    int fd;
    char * myfifo = "/tmp/myfifo";
    /* create the FIFO(named pipe)*/
    mkfifo(myfifo, 0666);

    /* write "Hi" to the FIFO */
    fd = open(myfifo, O_WRONLY);
    write(fd, "Hi", sizeof("Hi"));
    close(fd);

    /* remove the FIFO */
    unlink(myfifo);

    return 0;
}
```

## Process 2

```
#define MAX_BUF 1024

int main() {
    int fd;
    char * myfifo = "/tmp/myfifo";
    char buf[MAX_BUF];

    /* open, read, and display the
       message from the FIFO */
    fd = open(myfifo, O_RDONLY);
    read(fd, buf, MAX_BUF);
    printf("Received: %s\n", buf);
    close(fd);

    return 0;
}
```

# Redirecting Standard I/O

---

- Redirection of stdin, stdout and stderr in Unix and Linux

<< or 0<<	redirect stdin within command
> or 1>	redirect stdout to file (overwrite if file not empty)
>> or 1>>	redirect stdout to file (append if file not empty)
2>	redirect stderr to file (overwrite if file not empty)
2>>	redirect stderr to file (append if file not empty)

- Redirecting data from one program to another
  - `ls | head -3`
  - `ls | head -3 | tail -1`
  - `ls | head -3 | tail -1 > myoutput`



# Redirecting Standard I/O

## Program

```
main( int ac, char *av[]){  
    int i;  
    printf ("No. of args: %d, Args:\n", ac);  
    for (i=0; i < ac; i++)  
        printf ("args[%d]: %s\n", i, av[i]);  
  
    fprintf (stderr, "Msg sent to stderr.\n");  
}
```

```
$ listargs > output 2> error
```

```
$ cat output
```

```
No. of args: 1, Args:  
args[0]: listargs
```

```
$ cat error
```

```
Msg sent to stderr.
```

```
$ listargs arg1 > output 2>error
```

```
$ cat output
```

```
No. of args: 2, Args:  
args[0]: listargs  
args[1]: arg1
```

# Redirection in C Programs

---

- Processes do not read from files, they read from file descriptors

- Close / Open Method:

```
close(0);           /* file descriptor 0 - i.e. stdin */
open(filename, O_RDONLY); /* redefines lowest file descriptor */
```

- Open / Close / Dup / Close / Method:

```
fd=open(filename, O_RDONLY); /* open file to read - fd */
close(0);                    /* close file descriptor 0 - i.e.
                             stdin */
newfd=dup(fd);               /* make "clone" of fd and
                             uses lowest next available fd -
                             I.e. 0 that was closed */
close(fd);                   /* close original fd associated
                             with file. Now, stdin is
                             associated with file... */
```

**dup2** is similar to dup, but dup2 will automatically close(0).  
For example, newfd=dup2(fd,0);

# Any Question So Far?

---



2-Process

91