

Operating Systems

4. Multithreaded Programming

Recall the Main Concepts Behind Processes

1. Resource Ownership


- Virtual address space to hold the process image
 - Program
 - Data
 - Stack
 - Attributes
- OS shields processes from interfering with each others resources
 - Protection

2. Scheduling/ Dispatching

- Execution may be interleaved with other processes
 - Maintain execution states, etc.
- Can we **decouple** this functionality?

Processes and Threads

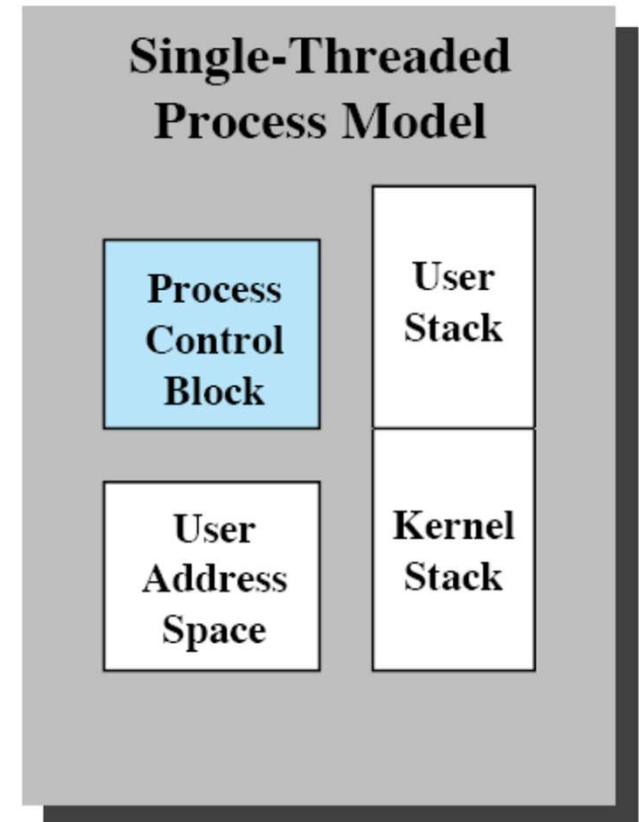
Multithreading: More than one entities can possibly execute in the same resource- (i.e., process-) environment (and collaborate better)



- Unit of **dispatching**
 - Referred to as a **thread**
- Unit of **resource ownership**
 - Referred to as a **process** or **task**

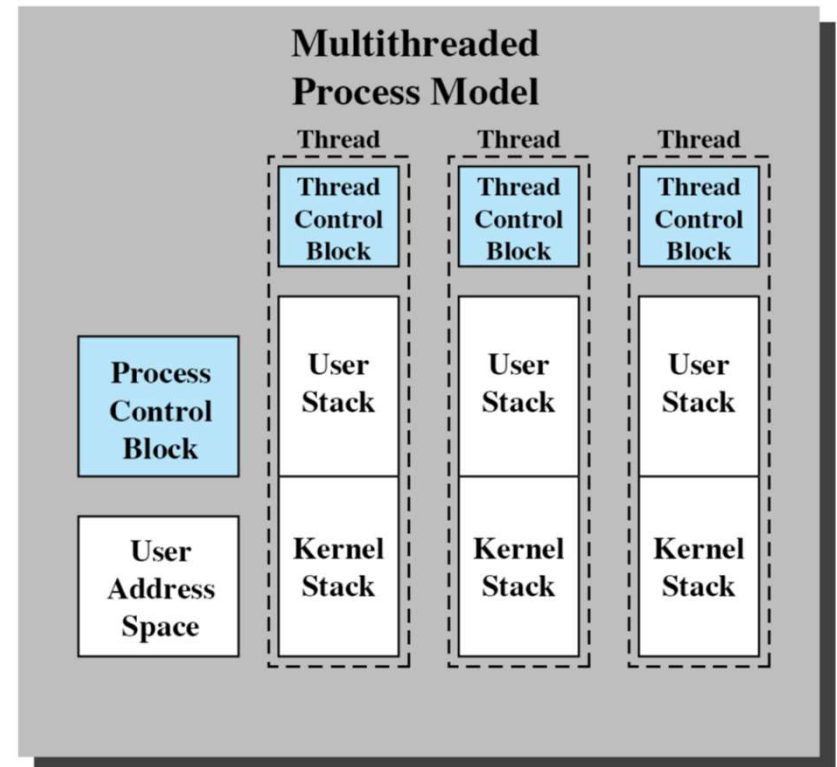
Processes vs. Threads

- Process characteristics
 - A **virtual address space**
 - Holds the process image
 - Global variables, files, child processes, signals and signal handlers

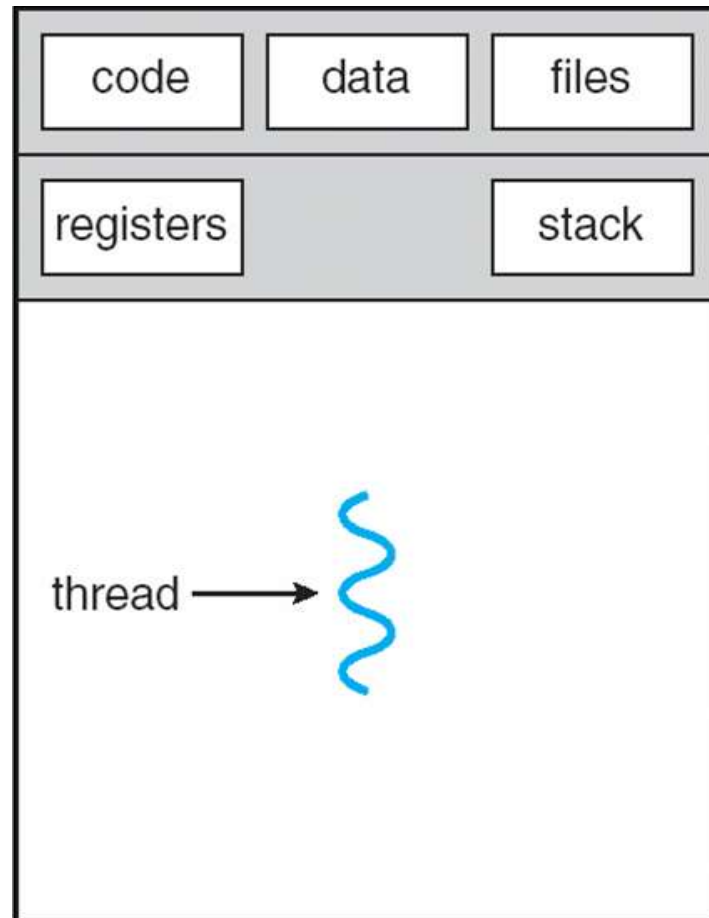


Processes vs. Threads

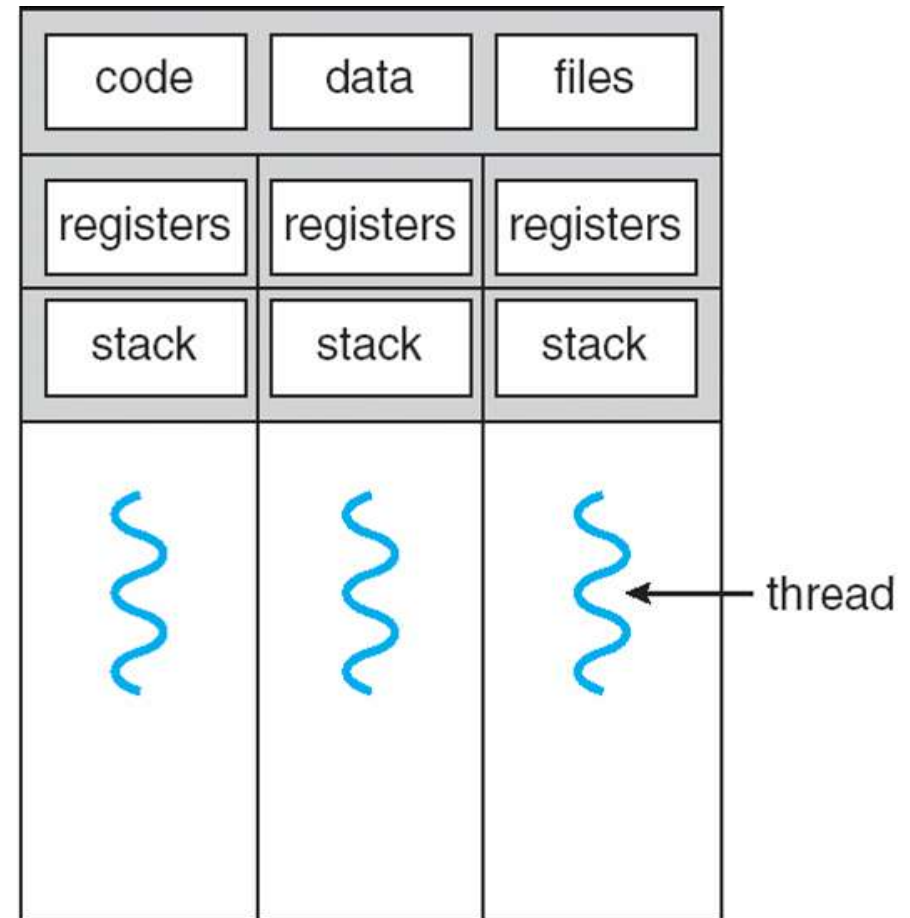
- Process characteristics
 - A **virtual address space**
 - Holds the process image
 - Global variables, files, child processes, signals and signal handlers
- Thread characteristics
 - An **execution state, stack and context** (saved when not running)
 - Access to the memory and resources of its process
 - All threads of a process share this
 - Some per-thread static storage for local variables



Processes vs. Threads

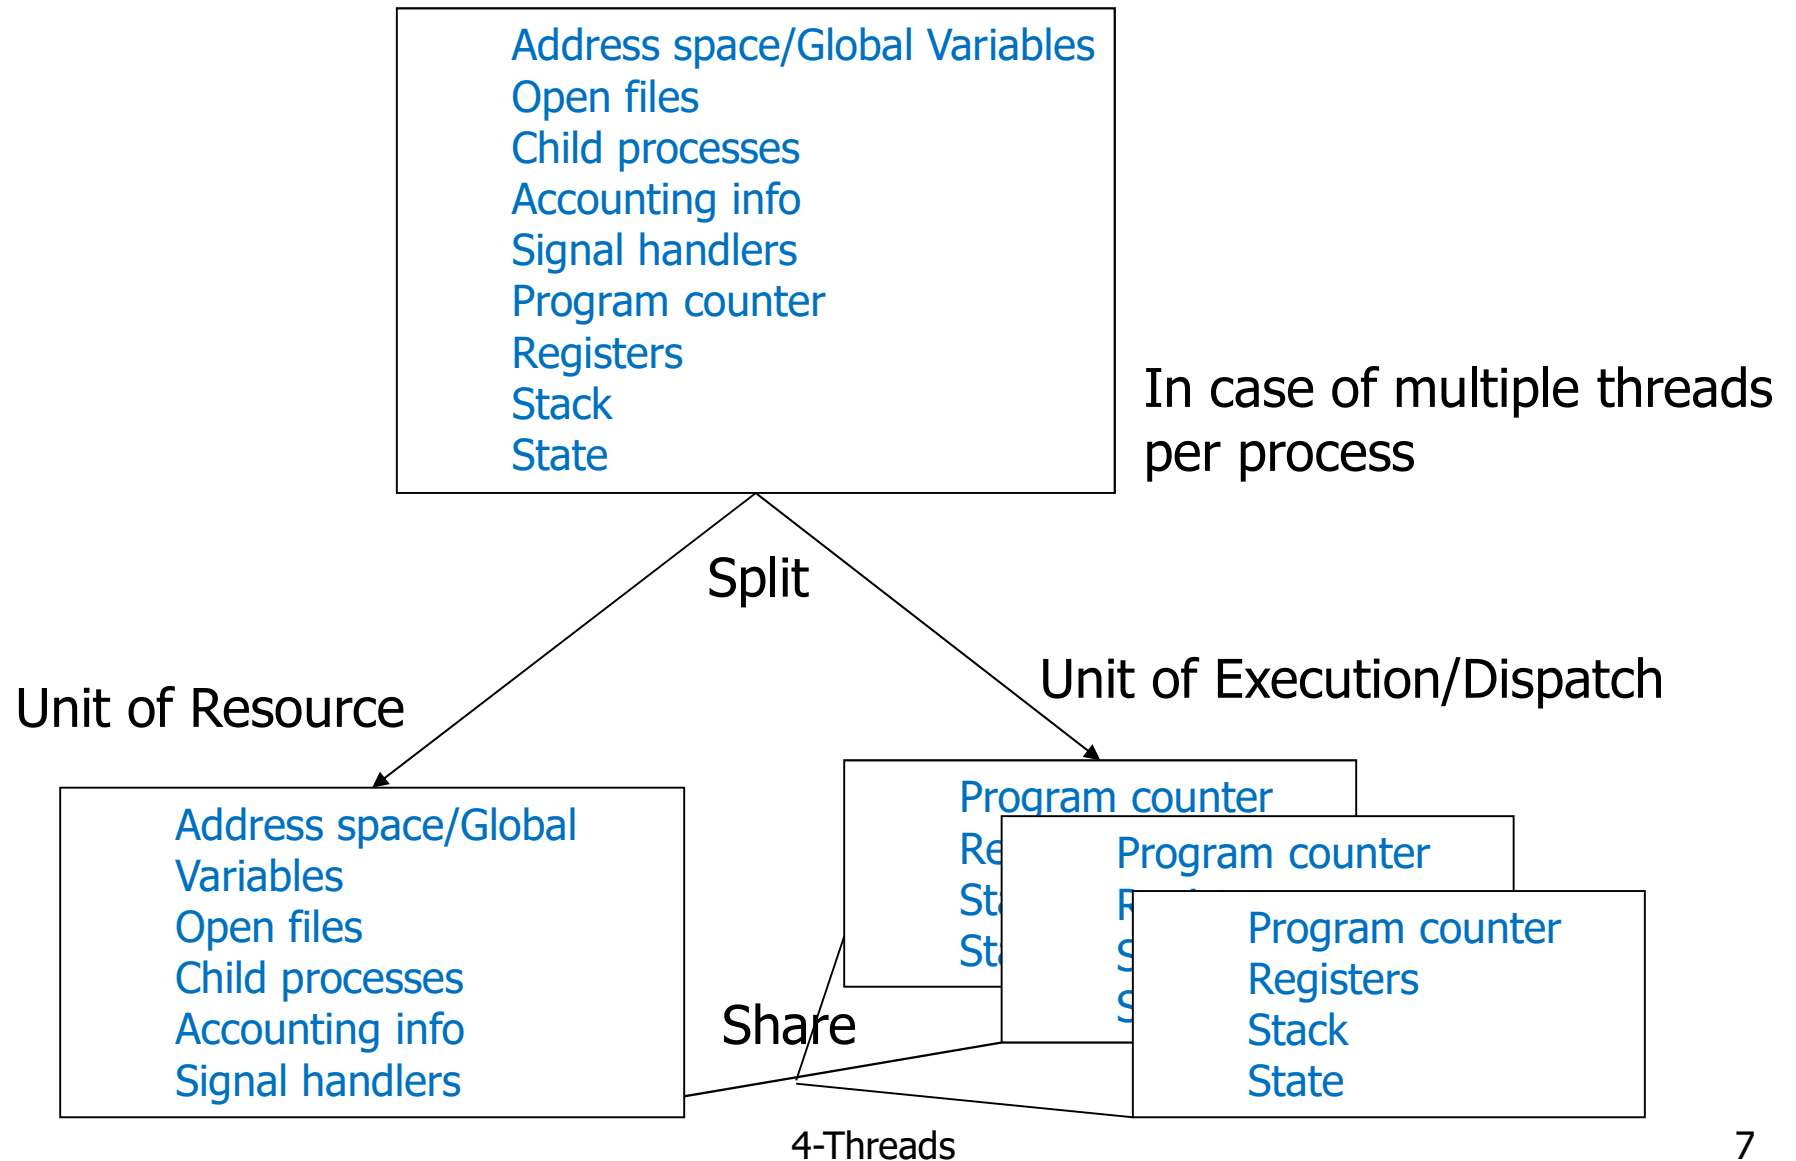


single-threaded process



multithreaded process

Processes vs. Threads



Benefits of Threads

- Easy/Lightweight Communication
 - Threads within the same process **share memory and files**
 - **Communication** does not invoke the kernel
- May allow parallelization within a process
 - **I/O and computation to overlap**
 - Recall historical evolution from uni-programming to multiprogramming
 - **Concurrent execution in multiprocessors**
- Takes **less time to**
 - **Create/terminate** a thread than a process
 - **Switch between two threads** within the same process

Uses of Threads

1. Overlap foreground with background work

- Decouple interactions from processing
- For instance processing message requests

Thread I
Wait for message
Store result in buffer

2. Asynchronous processing

- Backup while editing

Thread II
Wait for message
Store result in buffer

3. Speed up execution

- Parallelize independent actions

Thread III
Consume messages
in buffer


4. Modular program structure

- Must be careful here, not to introduce too much extra overhead


Thread IV
Consume messages
in buffer

⇒ Increase throughput
in message processing

Multithreaded Server Architecture – Example



Parallel execution on single core and multicore




Effect of Suspension/Termination

- Suspending a process involves suspending all threads of the process since all threads share the same address space
- Termination of a process, terminates all threads within the process

Implementing Threads in User Space


- The kernel is not aware of the existence of threads
- **Run-time system** (thread-library in execution)
 - Responsible for **bookkeeping**, **scheduling** of threads
- Allows for **customized scheduling**
- Support for **any OS**
- But: **problem** with **blocking system calls**
 - Blocking of a thread causes blocking of the process
 - Threads cannot execute if a thread of the same process is blocked



Implementing Threads in User Space

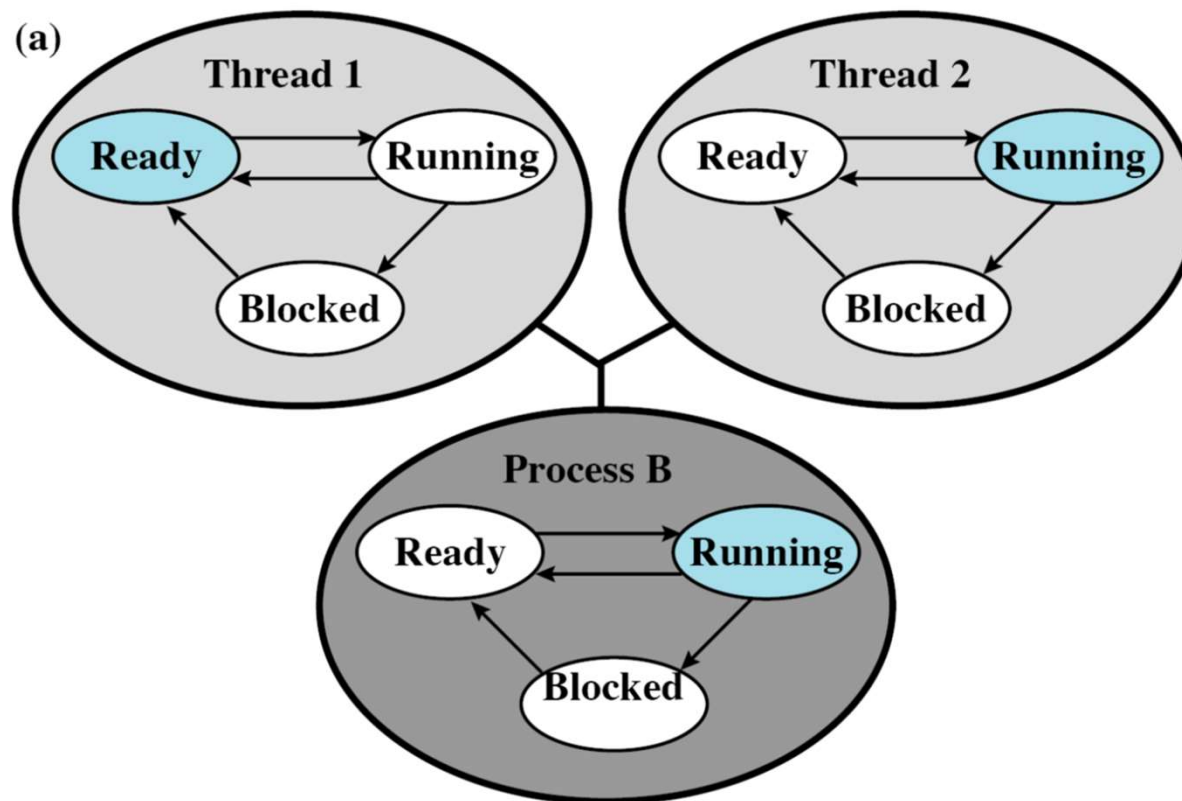
The thread library contains code for

- Creating and destroying threads
- Passing messages and data between threads
- Scheduling thread execution
 - Pass control from one thread to another
- Saving and restoring thread contexts
- Three primary thread libraries:
 - POSIX Pthreads
 - Win32 threads
 - Java threads




Example: User-Level Thread Execution

- Consider a thread makes a system call




Example: User-Level Thread Execution

- Process switches to blocked state
- State of Thread 2 remains running, but Thread 2 cannot execute instructions




Example: User-Level Thread Execution

- Process switches to ready state
 - Returned from the system call



Example: User-Level Thread Execution

- Thread 2 needs to wait for some action of Thread 1



User-Level Thread

Advantages

- Fast
 - Do not need context switching to kernel mode to perform thread management
- Scheduling can be application specific
 - Choose the best algorithm for the situation
- Can run on any OS
 - Only a thread library is required


Disadvantages

- Most system calls are blocking for processes
 - All threads within a process will be implicitly blocked
- Cannot benefit from the multiprocessor system
 - Kernel assign one process to only one CPU

Implementing Threads in Kernel/Lightweight Process

- Scheduling
 - Happens on a thread basis
- Does not suffer from the “blocking problem”
- Threads of a process can be scheduled on multiple processors
- Less efficient than user-level threads
 - Kernel is invoked for thread creation,
 - Termination
 - Switching

Kernel maintains context information for the process and the threads



Kernel-Level Thread Support

- Windows XP/2000
- Solaris
- Linux
- Tru64 UNIX
- Mac OS X

Kernel-Level Thread

Advantages

- In multiprocessor system, multiple threads from the same process can be scheduled on multiple processors
- Blocking at thread level, not process level
 - If a thread blocks, the CPU can be assigned to another thread in the same process
- Even the kernel routines can be multithreaded

Disadvantages

- Thread switching always involves the kernel
 - Two context switches per thread switch
- Slower compared to User Level Threads
 - Faster than a full process switch

Some Performance Data



Table 4.1 Thread and Process Operation Latencies (μ s) [ANDE92]

Operation	User-Level Threads	Kernel-Level Threads	Processes
Null Fork	34	948	11,300
Signal Wait	37	441	1,840

- Null Fork: Overhead of creating a process or thread
- Signal Wait: Overhead of synchronizing two processes or threads together

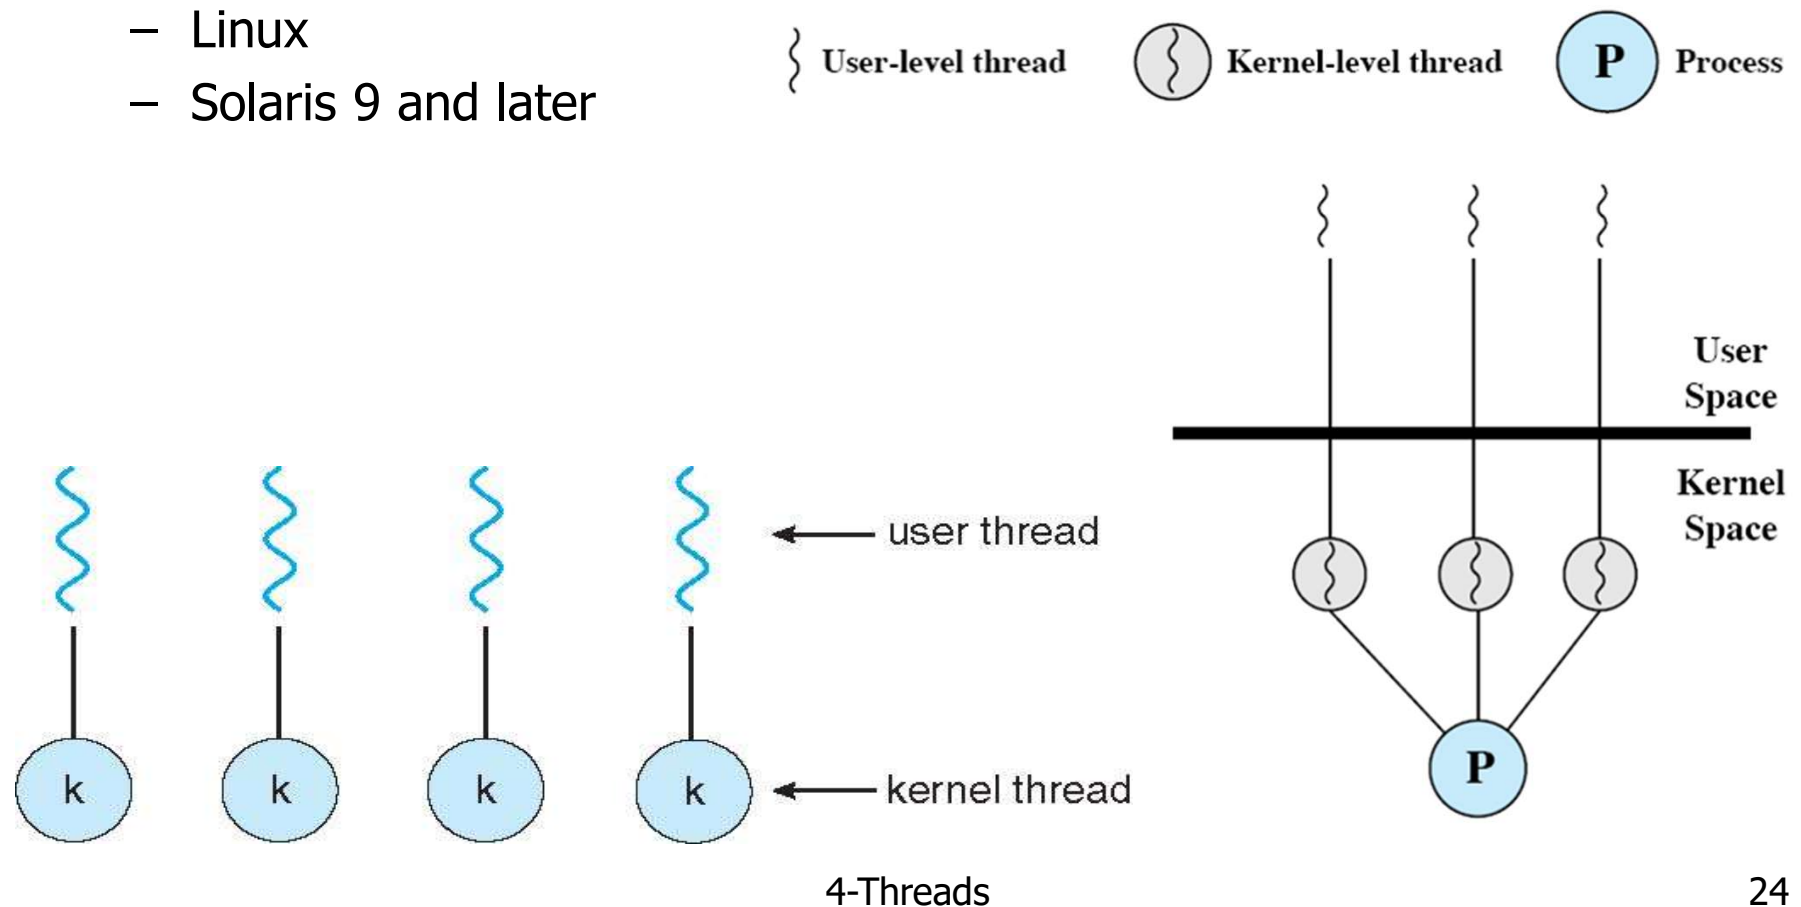
Multithreading Models: Many-to-one Model

- Many user-level threads mapped to single kernel thread
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads



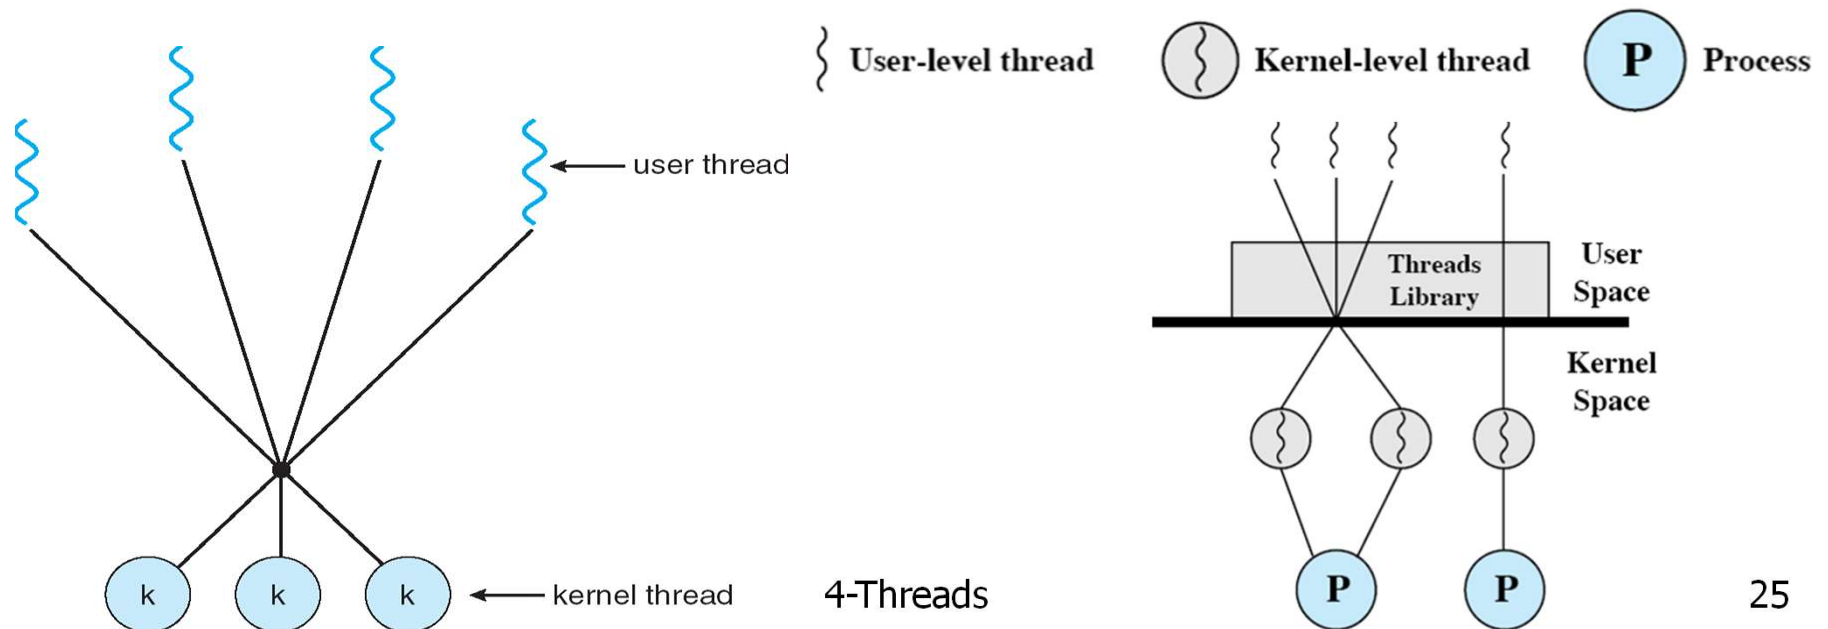
Multithreading Models: One-to-one Model

- Each user-level thread maps to kernel thread
- Examples
 - Windows NT/XP/2000
 - Linux
 - Solaris 9 and later

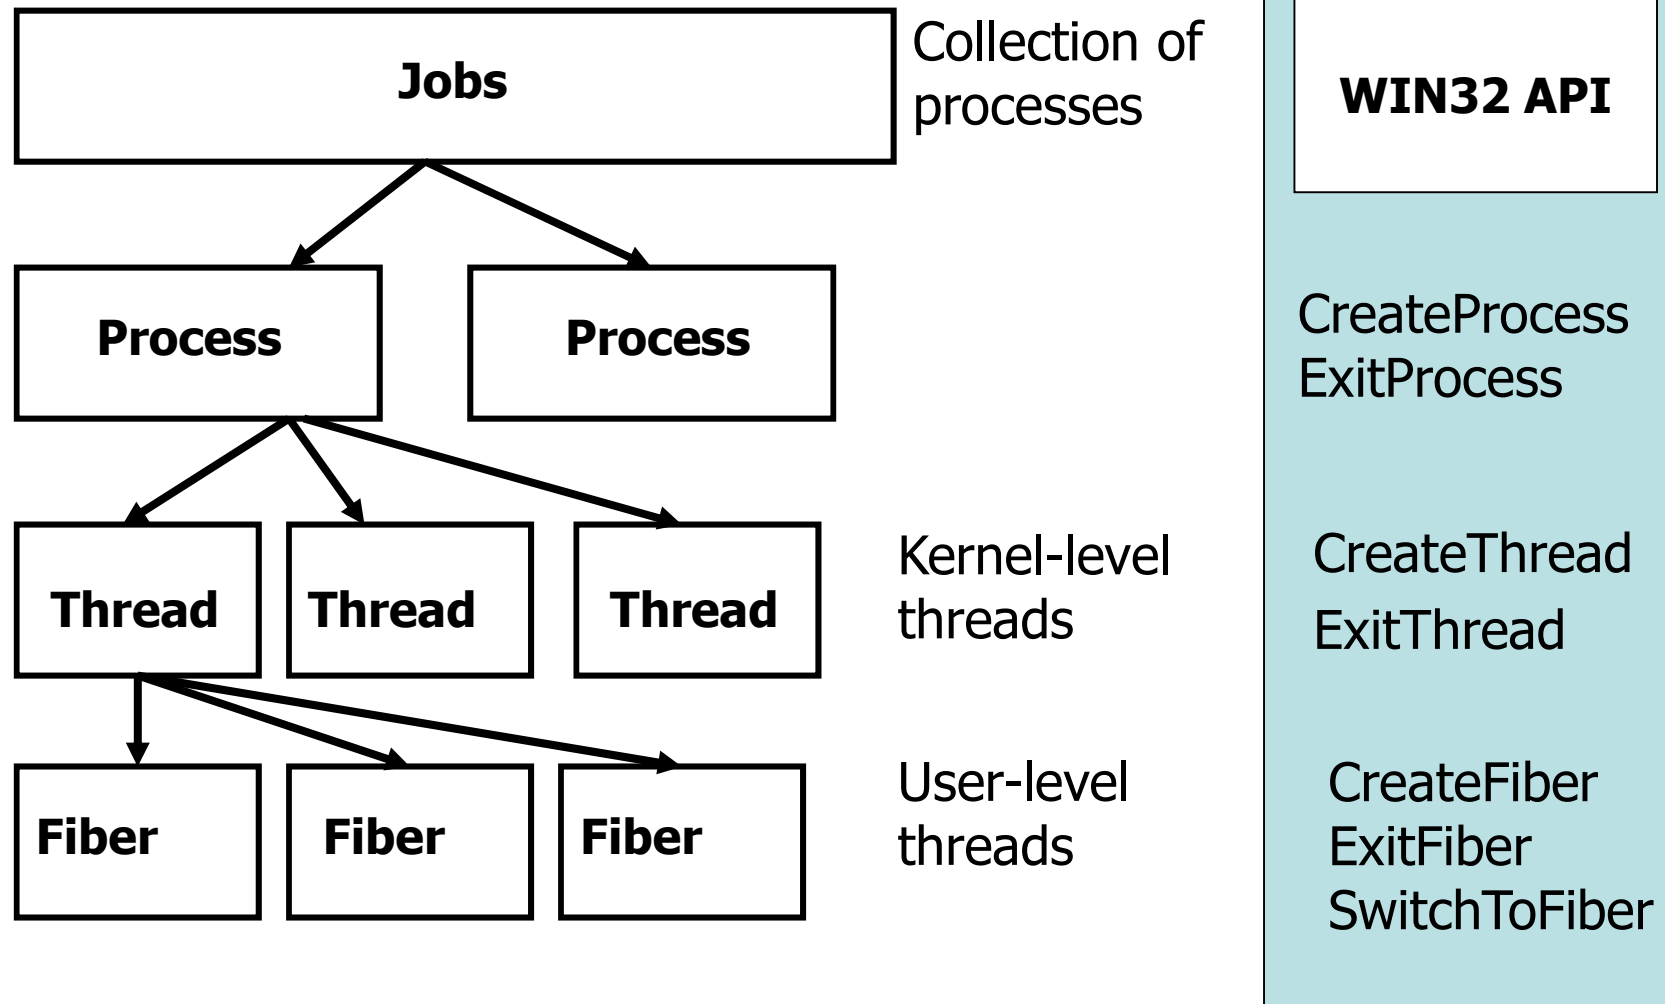


Multithreading Models: Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Example:
 - Solaris prior to version 9
 - Windows NT/2000 with the ThreadFiber package



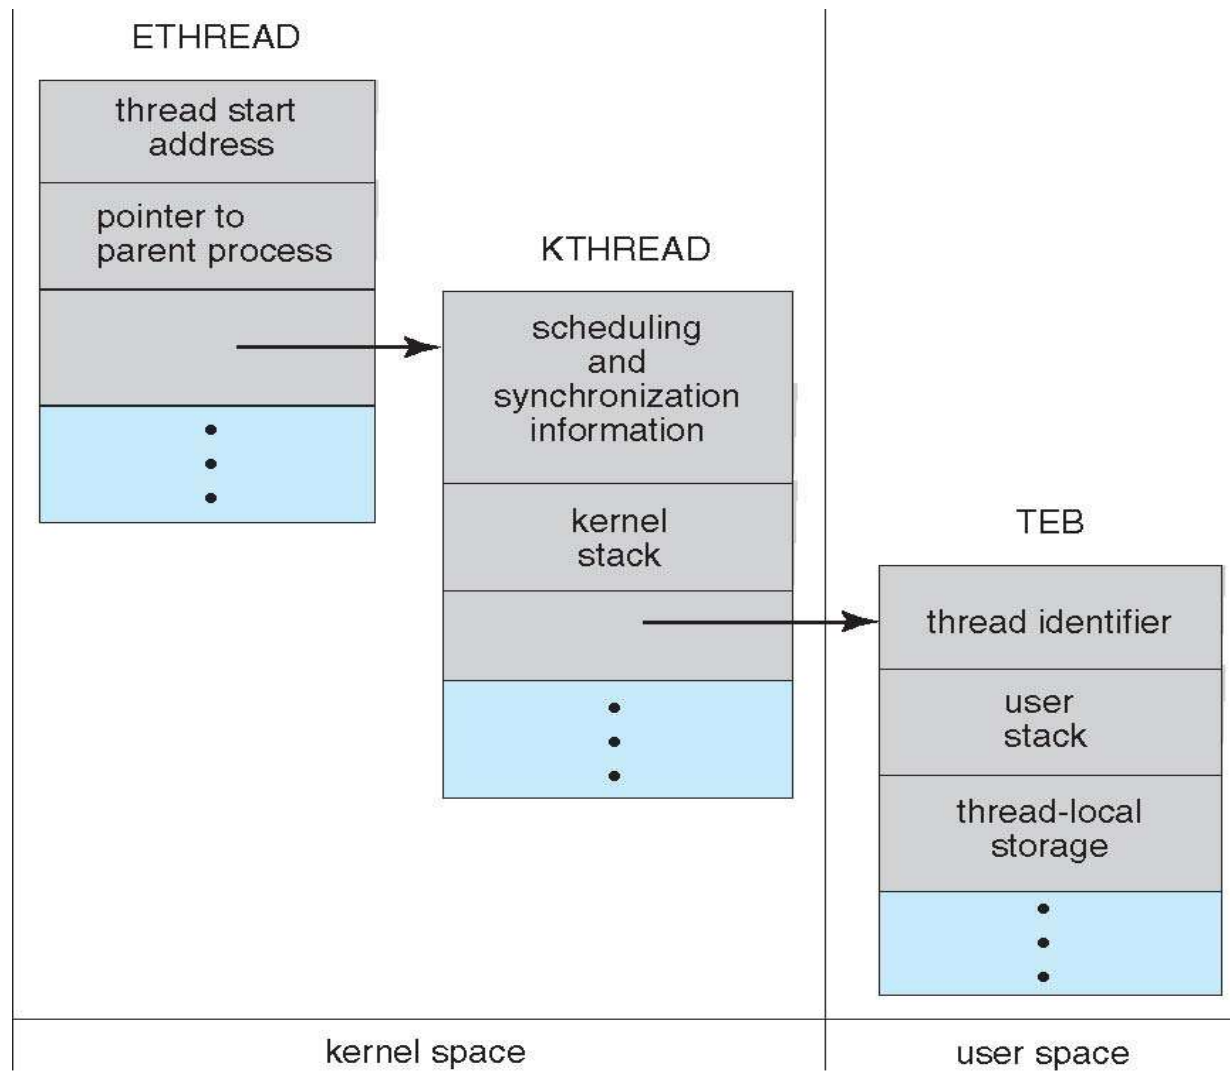
Windows: Threads and Processes



Windows XP: Threads and Processes

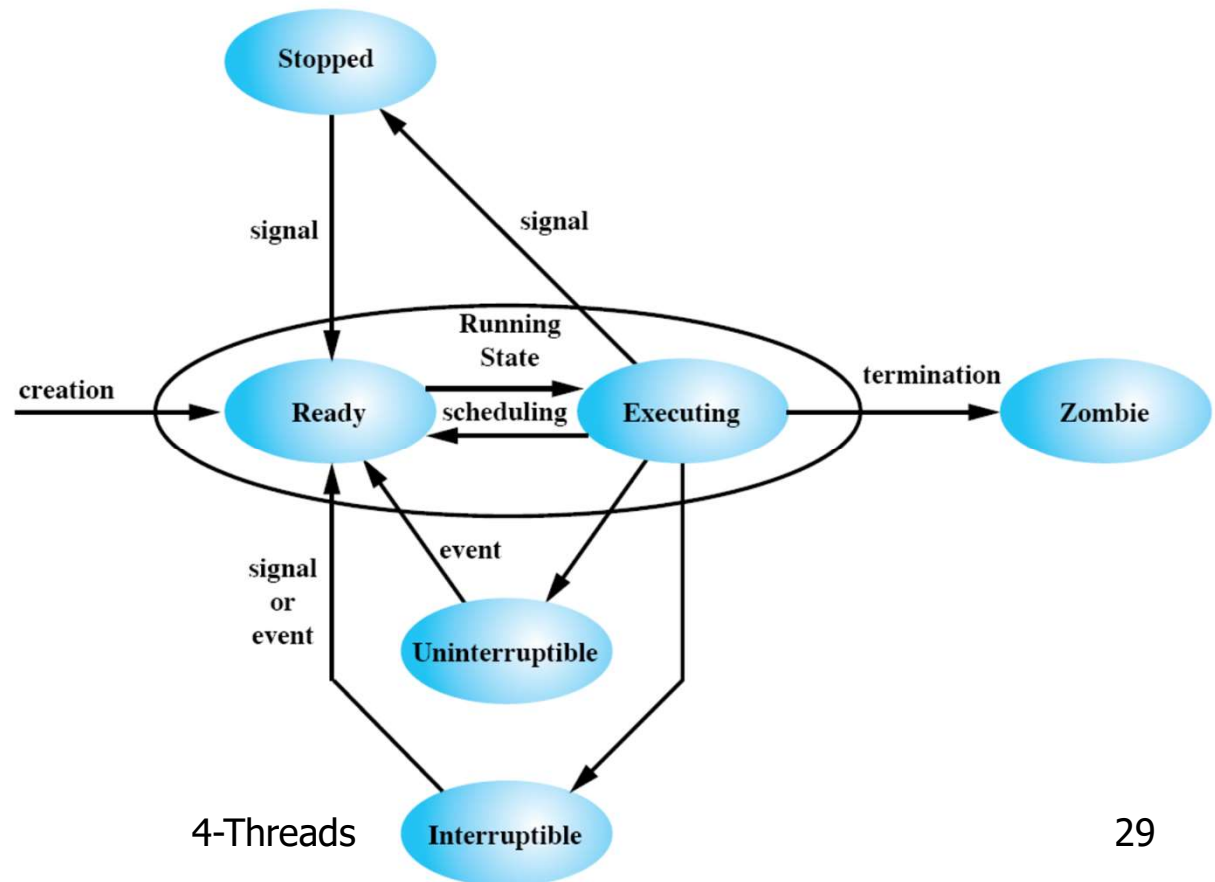
- Each thread contains
 - A thread id
 - Register set
 - Separate user and kernel stacks
 - Private data storage area
- The register set, stacks, and private storage area are known as the **context** of the threads
- The primary data structures of a thread include
 - ETHREAD (executive thread block)
 - KTHREAD (kernel thread block)
 - TEB (thread environment block)

Windows XP: Threads and Processes



Linux Task Management

- Linux uses **task** as a placeholder for processes and threads
- Task creation
 - Specify resources to be inherited from the parent
 - Share file system information
 - File descriptors
 - Memory space
 - ...



Linux Task Management

- `fork()` and `clone()` system calls
- `clone()` takes options to determine sharing between parent and child tasks

flag	meaning
<code>CLONE_FS</code>	File-system information is shared.
<code>CLONE_VM</code>	The same memory space is shared.
<code>CLONE_SIGHAND</code>	Signal handlers are shared.
<code>CLONE_FILES</code>	The set of open files is shared.

- No sharing takes place, if none of above flags is set when `clone()` is invoked
 - Functionality similar to that provided by `fork()` system call

Linux Task Management

- Unique kernel data structure (`struct task_struct`) for each task
 - Instead of storing data, contains pointers to other data structures
 - For example, data structure that represent
 - open files
 - Signal handlers
 - Virtual memory
- When `fork()` is invoked
 - New task is created with copy of all associated data structures of the parent process (task)
- When `clone()` is invoked
 - New task points to the data structure of the parent process (task)
 - Depending on the set of flags passed to `clone()`

Thread Libraries

- Thread library provides programmer with API for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS

Thread Libraries

- **POSIX threads** (Pthreads)
 - Execution model for threads (specified by IEEE)
 - Specifies application programming interface
 - Implementation (using user-/kernel-level threads) is up to developers
 - More common in UNIX systems
- **Win32 thread library**
 - Kernel-level library, windows systems
- **Java threads**
 - Supported by the JVM
 - Java thread API is implemented using a thread library available on the host operating system, e.g., Pthreads, Win32

Pthreads API Overview

Thread creation:	<pre>int pthread_create (pthread_t *thread_id, const pthread_attr_t *attributes, void *(*thread_function)(void *), void *arguments);</pre>
Thread termination:	<pre>int pthread_exit (void *status); int pthread_cancel (pthread_t thread);</pre>
Waiting for status of a thread:	<pre>int pthread_join (pthread_t thread, void **status_ptr);</pre>
Modify thread attributes:	<pre>int pthread_attr_init(pthread_attr_t *attr); int pthread_attr_destroy(pthread_attr_t *attr);</pre>

Thread Creation

```
int pthread_create (pthread_t *thread_id,  
                    const pthread_attr_t *attributes,  
                    void *(*thread_function)(void *),  
                    void *arguments);
```

- This routine creates a new thread and makes it executable
- Thread stack is allocated and thread control block is created
- Once created, a thread may create other threads
 - When process starts an "initial thread" exists by default and is the thread which runs main
- Function pthread_create returns
 - zero, if success
 - Non-zero if error

Thread Creation

```
int pthread_create (pthread_t *thread_id,  
                    const pthread_attr_t *attributes,  
                    void *(*thread_function)(void *),  
                    void *arguments);
```

pthread_t *thread_id

- Hold the identifier of the newly created thread
- Caller can use this thread ID to perform various operations

Thread Creation

```
int pthread_create (pthread_t *thread_id,  
                   const pthread_attr_t *attributes,  
                   void *(*thread_function)(void *),  
                   void *arguments);
```

`pthread_attr_t *attributes`

- Used to set thread attributes
- To create thread with default attributes
 - Attribute object is not specified, i.e., NULL is used
- Default thread is created with following attributes
 - It is non-detached (i.e., joinable)
 - Detached: On termination all thread resources are released by OS
 - It has a default stack and stack size
 - It inherits the parent's priority

Thread Creation

```
int pthread_create (pthread_t *thread_id,  
                    const pthread_attr_t *attributes,  
                    void *(*thread_function)(void *),  
                    void *arguments);
```

`void *(*thread_function)(void *)`

- The C routine that the thread will execute once it is created

Thread Creation

```
int pthread_create (pthread_t *thread_id,  
                    const pthread_attr_t *attributes,  
                    void *(*thread_function)(void *),  
                    void *arguments);
```

void *arguments

- Arguments to be passed to **thread_function**
- Arguments must be passed by reference and cast to **void***
- These pointers must be cast as pointers of type void

Thread Creation: Example

```
#include <pthread.h>
#define NUM_THREADS 5

void * PrintHello(void *threadid) {
    printf("\n%d: Hello World!\n", threadid);
}

int main() {
    pthread_t threads[NUM_THREADS];
    int rc, t;
    for(t=0; t < NUM_THREADS; t++) {
        printf("Creating thread %d\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)&t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
}
```


Possible Output Sequence


- No deterministic execution!
 - Thread execution can interleave in multiple ways

Possible Output

```
In main: creating thread 0
In main: creating thread 1
Hello World! It's me, thread #0!
In main: creating thread 2
Hello World! It's me, thread #1!
Hello World! It's me, thread #2!
In main: creating thread 3
In main: creating thread 4
Hello World! It's me, thread #3!
Hello World! It's me, thread #4!
```

Thread Creation

- Once created, threads are peers, and may create other threads
 - No implied hierarchy or dependency between threads



Thread Termination

Several ways of termination

- The thread makes a call to the `pthread_exit()` routine
- The thread is canceled by another thread via `pthread_cancel()` routine
- The entire process is terminated due to a call to the `exit()`
- The thread returns from its starting routine/function
 - Same as if there was an implicit call to `pthread_exit()` using the return value of thread function as the exit status
 - Behavior of thread executing `main()` differs
 - Effect: Implicit call to `exit()` using the return value of `main()` as the exit status

Thread Termination

```
int pthread_exit (void *status);
```

- The `pthread_exit()` routine does not close files
 - Any files opened inside the thread will remain open after the thread is terminated
- If the "initial thread" exits with `pthread_exit()` instead of `exit()`, other threads will continue to execute
- The programmer may specify a termination status
 - Stored as a void pointer for any thread that may join the calling thread i.e., wait for this thread

Thread Termination: Example

```
#include <pthread.h>
#define NUM_THREADS 5

void * PrintHello(void *threadid) {
    printf("\n%d: Hello World!\n", threadid);
    pthread_exit(NULL);
}

int main() {
    pthread_t threads[NUM_THREADS];
    int rc, t;
    for(t=0; t < NUM_THREADS; t++) {
        printf("Creating thread %d\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)&t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```

Passing Arguments to Threads

- The `pthread_create` permits the programmer to pass one argument to the thread start function
 - Argument must be passed by reference and cast to `(void *)`
- What if we want to pass multiple arguments?
 - Create a structure which contains all of the arguments
 - Pass a pointer to the structure in the `pthread_create`
- Once created threads have non-deterministic start-up & scheduling
 - How to safely pass data/argument to newly created threads?
 - Make sure that all passed data is thread safe
 - i.e., it cannot be changed by other threads
 - The calling function must ensure that argument remains valid for the new thread throughout its lifetime
 - At least until thread has finished accessing it

Passing Arguments to Threads

- **Incorrectly** passed arguments

```
for(t=0;t < NUM_THREADS;t++) {  
    printf("Creating thread %d\n", t);  
    rc = pthread_create(&threads[t], NULL, printHello,(void *) &t);  
    ...  
}
```

- **Correctly** passed arguments

```
int *tids[NUM_THREADS];  
for(t=0;t < NUM_THREADS;t++) {  
    tids[t] = new int;  
    *tids[t] = t;  
    printf("Creating thread %d\n", t);  
    rc = pthread_create(&threads[t],NULL,PrintHello,(void*)tid[t]);  
    ...  
}
```

Passing Structure as Argument

```
struct thread_data {  
    int thread_id;  
    int sum;  
};  
thread_data thread_data_array[NUM_THREADS];  
int main() { ...  
    thread_data_array[t].thread_id = t;  
    thread_data_array[t].sum = sum;  
    rc = pthread_create(&threads[t], NULL, PrintHello,(void *)&thread_data_array[t]) ;  
    ...  
}  
void *PrintHello(void *threadarg) {  
    thread_data *my_data;  
    my_data = (struct thread_data *) threadarg;  
    taskid = my_data->thread_id;  
    sum = my_data->sum;  
    ...  
}
```


Thread Identifier

```
pthread_t pthread_self(void);
```

- Returns the unique thread ID of the calling thread

```
int pthread_equal(pthread_t t1, pthread_t t2)
```

- Compares two thread IDs
 - If the two IDs are different 0 is returned
 - Otherwise a non-zero value is returned

Changing Thread Attributes

```
int pthread_attr_init(pthread_attr_t *attr);  
int pthread_attr_destroy(pthread_attr_t *attr);
```

To change the (default) attributes of the thread, the following steps can be performed

- Initialize attributes object with `pthread_attr_init()` function
- Afterwards, individual attributes of the object can be set using various related functions
 - `pthread_attr_setdetachstate()` to set the detached state of thread
 - `pthread_attr_setstacksize()` function sets the stack size
- When a thread attributes object is no longer required, free library resources using the `pthread_attr_destroy()` function

Changing Thread Attributes: Example

```
#define NUM_THREADS 5
#define MEGEXTRA 1000000
#define N 1000

pthread_attr_t attr;
int main() {
    pthread_t threads[NUM_THREADS];
    int rc, t;
    size_t stacksize
    pthread_attr_init( &attr);
    pthread_attr_getstacksize(&attr, &stacksize);
    printf("Default stack size - %li\n", stacksize);
    stacksize = sizeof(double) * N * N + MEGEXTRA;
    pthread_attr_setstacksize (&attr, stacksize);
    for(t=0;t < NUM_THREADS; t++) {
        printf("Creating thread with stacksize - %li\n", stacksize);
        rc = pthread_create(&threads[t], &attr, PrintHello, (void *)&t);
        . . .
    }
    pthread_attr_destroy(&attr):
    . . .
```

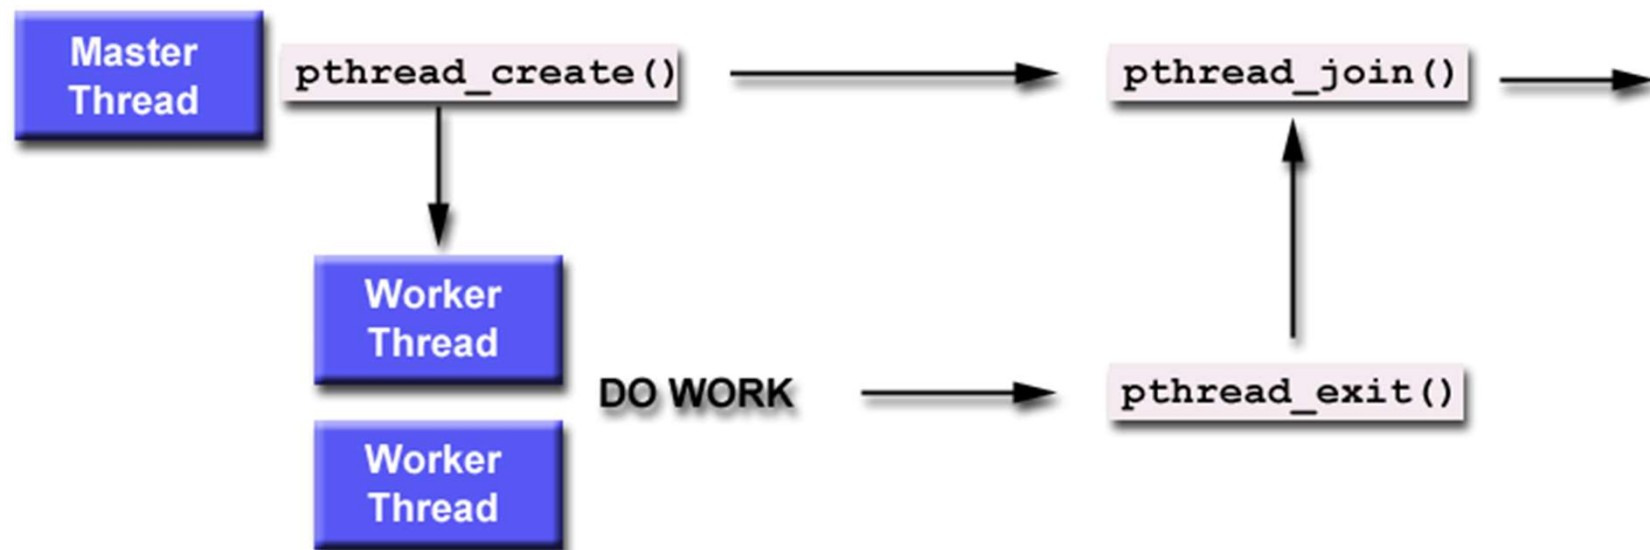
Thread Suspension and Termination

- Similar to UNIX processes, threads have the equivalent of the `wait()` and `exit()` system calls

```
int pthread_join (pthread_t thread,  
                 void **status_ptr);
```

- Use to block threads
- To instruct a thread to block and wait for a thread to complete
- Any thread can call join on (and hence wait for) any other thread
- Serves as a mechanism to get a return value from a thread

Thread Suspension and Termination



Joinable Threads

- When a thread is created, one of its attributes defines whether it is **joinable** or **detached**
 - Only threads that are created as joinable can be joined
 - If a thread is created as detached, it can never be joined
- Detached Thread
 - On termination all thread resources are released by the OS
 - A detached thread cannot be joined
 - No way to get at the return value of the thread
- Joinable Thread
 - On thread termination thread ID and exit status are saved by the OS
 - Joining a thread means waiting for a thread with a certain ID
 - One way to accomplish synchronization between threads

Joinable Threads

- Multiple threads cannot wait for the same thread to terminate
- If multiple thread simultaneously try to join, the results are undefined
 - For example, only one thread returns successfully
 - The other threads fail with an error of ESRCH or EINVAL
- After `pthread_join()` returns, any stack storage associated with the thread can be reclaimed by the process
- Threads which have exited but have not been joined are equivalent to zombie processes
 - Their resources cannot be fully recovered

Joinable Threads: Example

```
#include <stdio.h>
#include <pthread.h>
```

```
int x = 5;
```

```
int child_code()
{
    while (x>0){
        x=x-1;
        printf("x has changed to %d\"
            ,x);
    }
    return 0;
}
```

```
int main(int argc, char *argv[])
{
    int rc;
    void** status=0;
    pthread_t child_thread;
    pthread_attr_t attr;
```

```
/* Set threads properties */
pthread_attr_init(&attr);
pthread_attr_setdetachstate(&attr,
    PTHREAD_CREATE_JOINABLE);
```

```
/* Create new thread */
rc = pthread_create(&child_thread,
    &attr, (void*) &child_code, NULL);
/* Free attributes object */
pthread_attr_destroy(&attr);
```

```
/* concurrent computation */
x=x+5;
printf("x has change to %d\n",x);
```

```
/*wait for thread to terminate */
pthread_join(child_thread, status);
pthread_exit(NULL);
}
```


Possible Outcome

No deterministic execution!

- Two threads can interleave in multiple ways
- Requires further synchronization

Child: x=5

```
while (x>0){  
    x=x-1;  
    printf("x has changed to %d \n",x);  
}
```

- x has changed to 4
- x has changed to 3

Parent: x=3

```
x=x+5;  
printf("x has changed to %d\n",x);
```

- x has changed to 8
- x has changed to 7

Child: x=3

```
while (x>0){  
    x=x-1;  
    printf("x has changed to %d \n",x);  
}
```

- ...
- x has changed to 1
- x has changed to 0

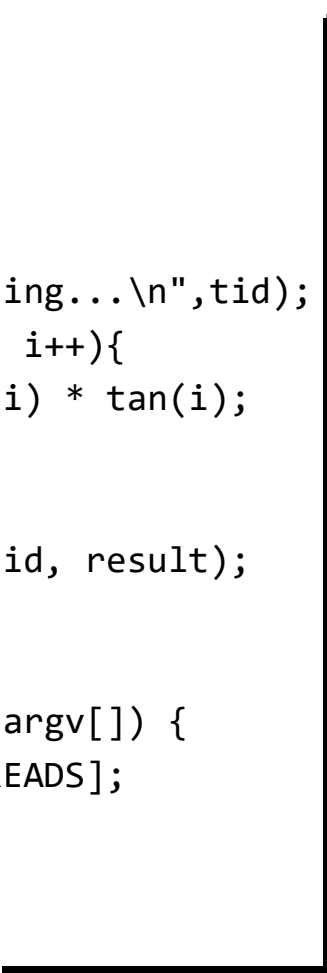
Joinable Threads: Example

```
#include <pthread.h>
#define NUM_THREADS 4

void *BusyWork(void *t){
    long tid = (long)t;
    double result=0.0;
    printf("Thread %ld starting...\n",tid);
    for (int i=0; i<1000000; i++){
        result = result + sin(i) * tan(i);
    }
    printf("Thread %ld done.
           Result = %e\n",tid, result);
    pthread_exit((void*) t);
}

int main (int argc, char *argv[]) {
    pthread_t thread[NUM_THREADS];
    pthread_attr_t attr;
    int rc;
    long t;
    void *status;

    /* Initialize & set thread attribute */
    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr,
        PTHREAD_CREATE_JOINABLE);
    for(t=0; t<NUM_THREADS; t++) {
        printf("Main:creating thread %ld\n", t);
        pthread_create(&thread[t], &attr,
            BusyWork, (void *)t);
    }
    /* Free attribute & wait for threads */
    pthread_attr_destroy(&attr);
    for(t=0; t<NUM_THREADS; t++) {
        pthread_join(thread[t], &status);
        printf("Main: completed join with thread
               %ld with status %ld\n", t,
               (long)status);
    }
    printf("Main: program completed. \n");
    pthread_exit(NULL);
}
```



Possible Output

```
Main: creating thread 0
Main: creating thread 1
Thread 0 starting...
Main: creating thread 2
Thread 1 starting...
Main: creating thread 3
Thread 2 starting...
Thread 3 starting...
Thread 1 done. Result = -3.153838e+06
Thread 0 done. Result = -3.153838e+06
Main: completed join with thread 0 with status of 0
Main: completed join with thread 1 with status of 1
Thread 3 done. Result = -3.153838e+06
Thread 2 done. Result = -3.153838e+06
Main: completed join with thread 2 with status of 2
Main: completed join with thread 3 with status of 3
Main: program completed. Exiting.
```

Making a Thread Detached

```
int pthread_detach (pthread_t thread_id);
```

- Function marks the thread identified by `thread_id` as detached
 - Return 0 on success
- To avoid memory leaks a thread should
 - Either be joined
 - Or detached by a call to `pthread_detach()`
- Once a thread is detached using this function, it cannot be made joinable again
- Threads can detach themselves by calling `pthread_detach()` with an argument of `pthread_self()`
 - i.e., `pthread_detach(pthread_self())`

Any Question So Far?

