

Lecture 02

Process Management

Dr. Ehtesham Zahoor

Some review

- How distributed systems are related to the operating systems?
- What is Cloud computing and the virtualization?
- How blockchain is related to our course?

Operating Systems Research

- This year, the course project would involve writing term paper and presentation
- Some words about the last year project

Process


- A computer program in execution on a machine is a process
- More formally:
 - A **Sequential stream of Execution** in its own **address space**

Why use processes?

- Why use processes?
 - Express concurrency
 - Systems have many concurrent jobs going on
 - E.g. Multiple users running multiple shells, I/O, ...
- General principle of divide and conquer
 - Decompose a large problem into smaller ones
 - easier to think of well contained smaller problems

Process =? Program

- Program: series of commands (e.g. C statements, assembly commands, shell commands)



Program != process

- Program: static code + static data
- Process: dynamic instantiation of code + data + more
- Program process: no 1:1 mapping
 - Process has more than code and data
 - one program runs many processes
 - many processes of same program

Process

- A computer program in execution on a machine is a process
- More formally:
 - A **Sequential stream of Execution** in its own **address space**

Process Address Space

- A list of memory locations from some min to some max that a process can read and write.
- Contains
 - the executable program
 - program's data
 - Stack
 - Associated with a process is a set of registers e.g. PC, SP and other information to run the program.

CPU State

- CPU registers contain the current state

Instruction Register (IR):

Program Counter (PC):

Stack Pointer (SP):

General purpose registers:

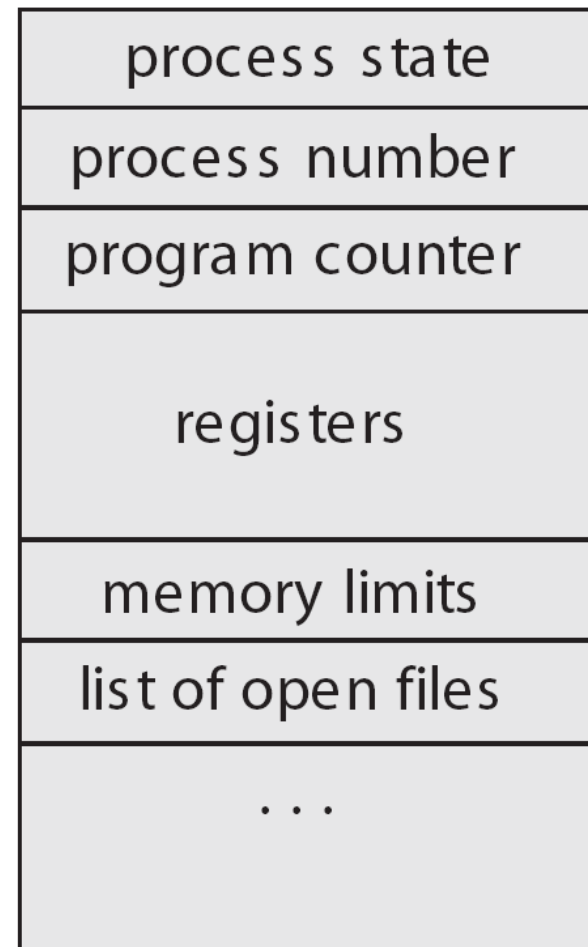
...

Environment

- Contains the relationships with other entities
- A process does not exist in a vacuum
- It typically has connections with other entities, such as
 - A terminal where the user is sitting.
 - Open files
 - Communication channels to other processes, possibly on other machines.

Process Control Block

- The OS keeps all the data it needs about a process in the process control block (PCB)



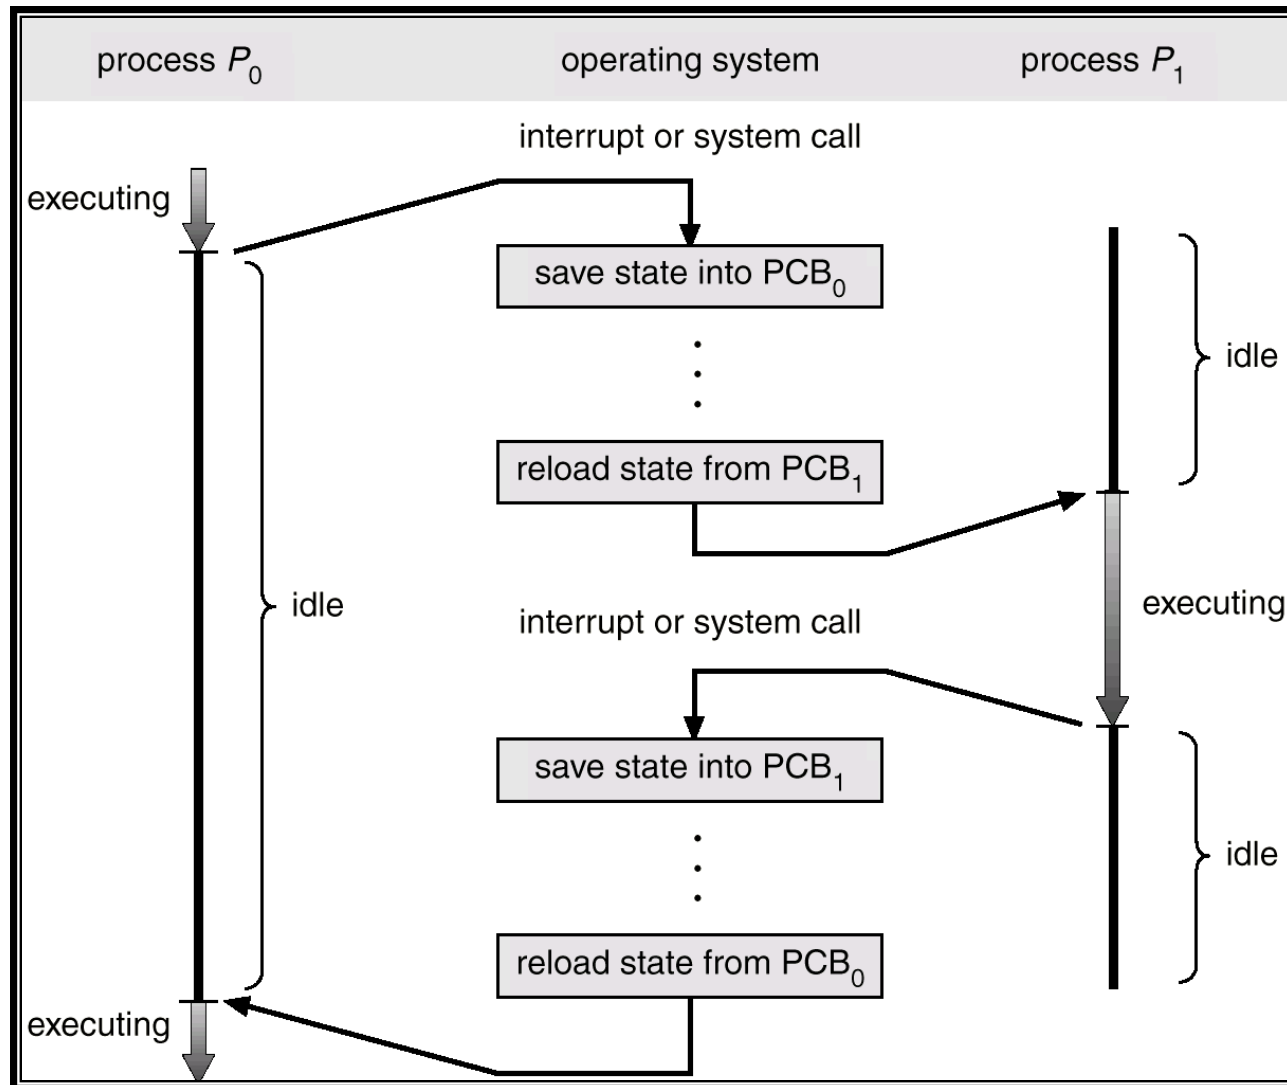
PCB

- PCB is "the manifestation of a process in an operating system".
- **Data Structure** defined in the operating system kernel containing the information needed to manage a particular process.
- It must be kept in an area of memory protected from normal user access.

PCB

- In general a PCB includes:
 - The identifier of the process (a process identifier, or PID)
 - Register values for the process including, notably, the program counter and stack pointer values for the process.
 - The address space for the process
 - Priority (in which higher priority process gets first preference. e.g., nice value on Unix operating systems)
 - Process accounting information, such as when the process was last run, how much CPU time it has accumulated, etc.
 - I/O Information (i.e. I/O devices allocated to this process, list of opened files, etc.)

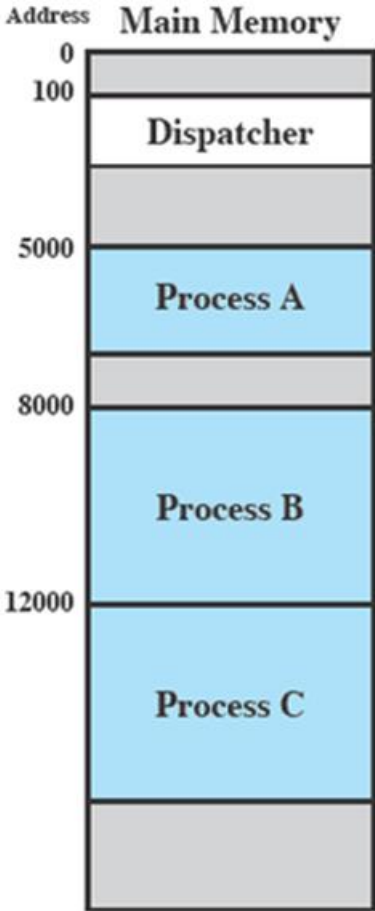
CPU Switch From Process to Process



CPU Switch From Process to Process

- Switching a process requires
 - Saving the state of old process
 - Loading the saved state of the new process
- This is called **Context Switch**
- Part of OS responsible for switching the processor among the processes is called **Dispatcher**

Lecture 02: Process Management



5000	8000	12000
5001	8001	12001
5002	8002	12002
5003	8003	12003
5004		12004
5005		12005
5006		12006
5007		12007
5008		12008
5009		12009
5010		12010
5011		12011

(a) Trace of Process A


(b) Trace of Process B

(c) Trace of Process C

5000 = Starting address of program of Process A
8000 = Starting address of program of Process B
12000 = Starting address of program of Process C

Figure 3.3 Traces of Processes of Figure 3.2

Trace from Processors point of view



Process States

- At any given time a process is either running or not running
- Number of states
 - Running
 - Not Running
- When the OS creates a process, the process is entered into which state?
 - Not Running

Five-state Process Model

- Running: currently being run
- Ready: ready to run
- Blocked: waiting for an event (I/O)
- New: just created, not yet admitted to set of run-able processes
- Exit: completed/error exit

Five-state Process Model

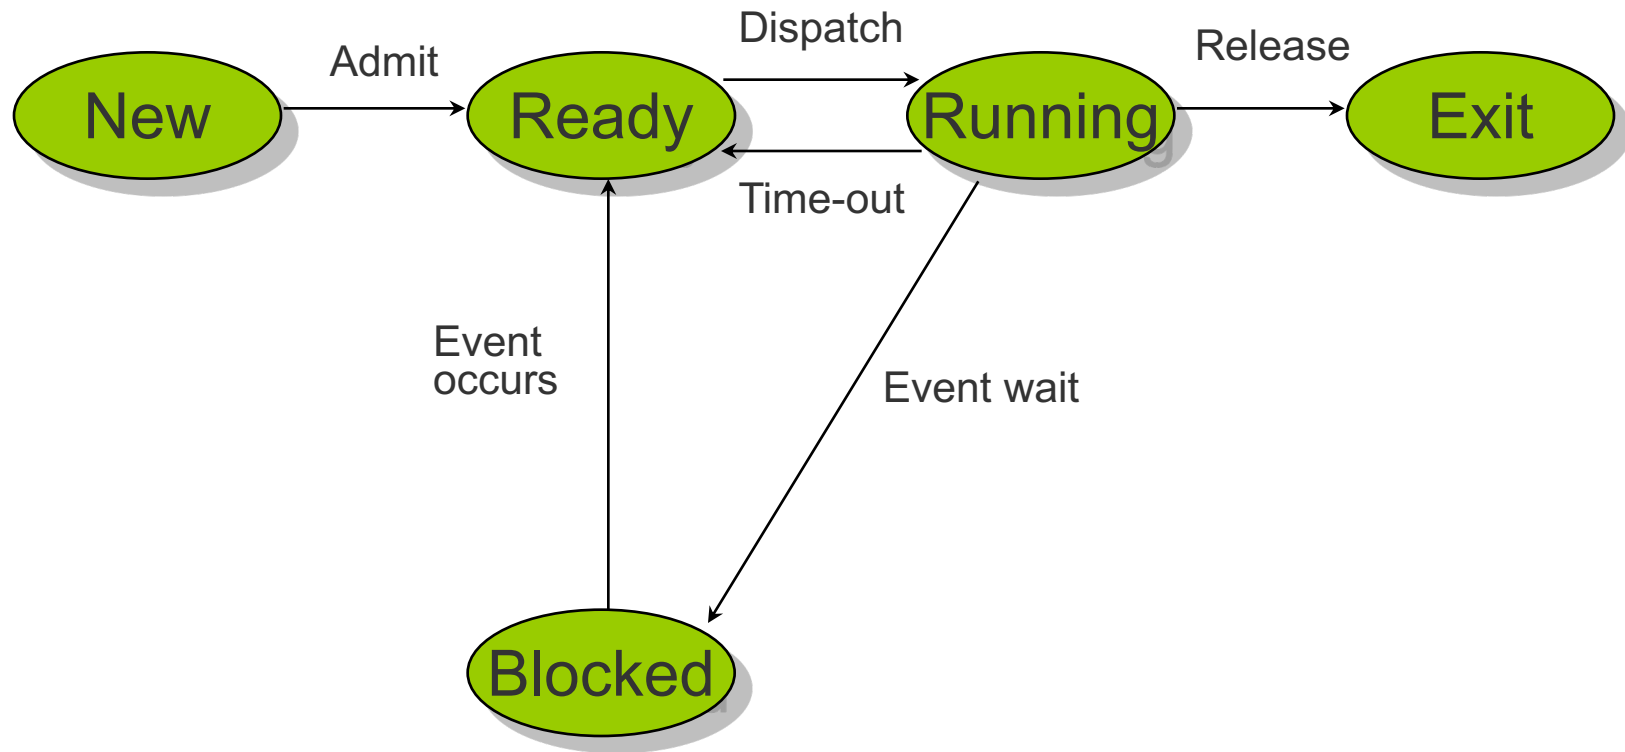

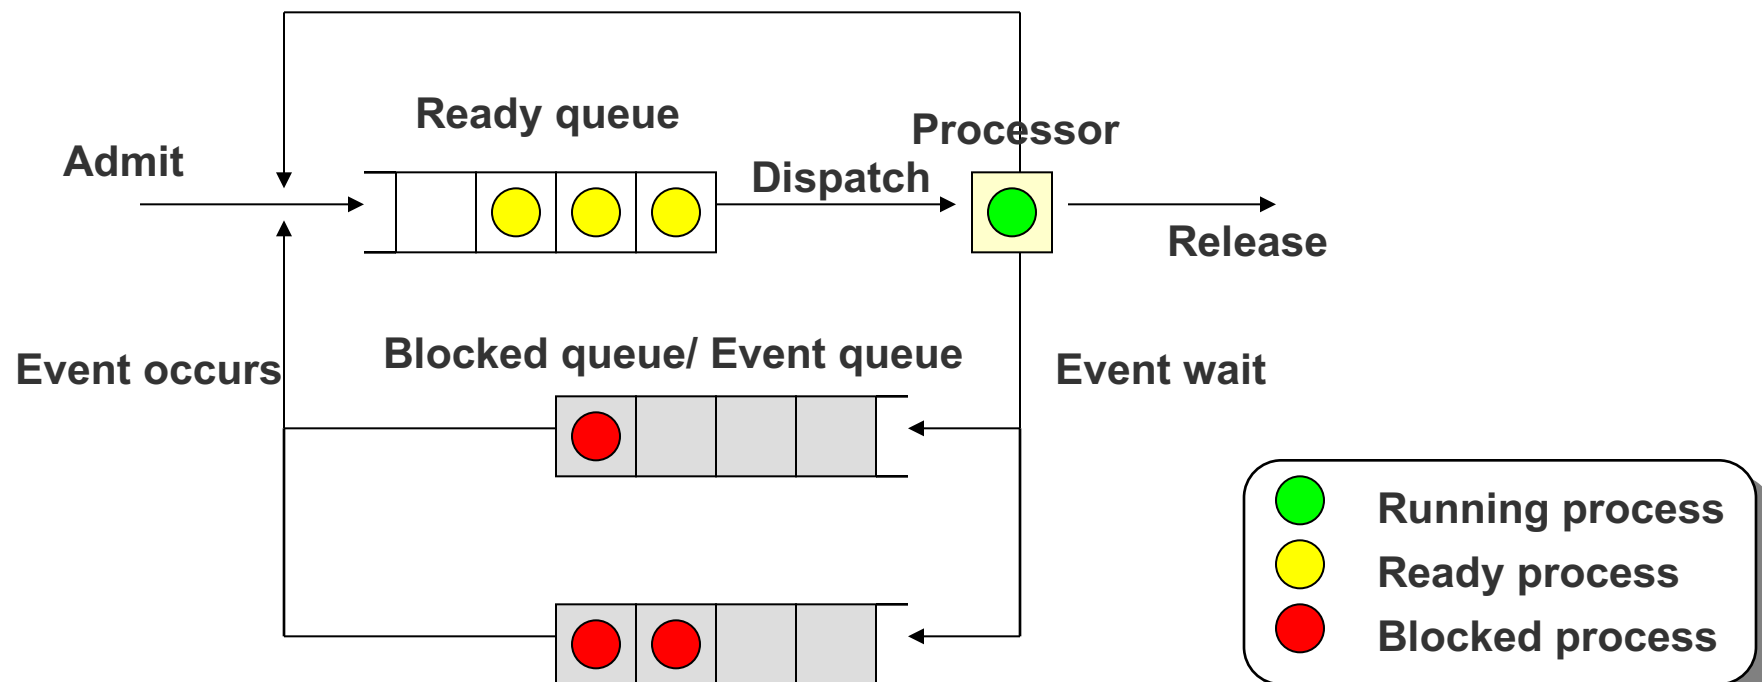


Diagram of Process State

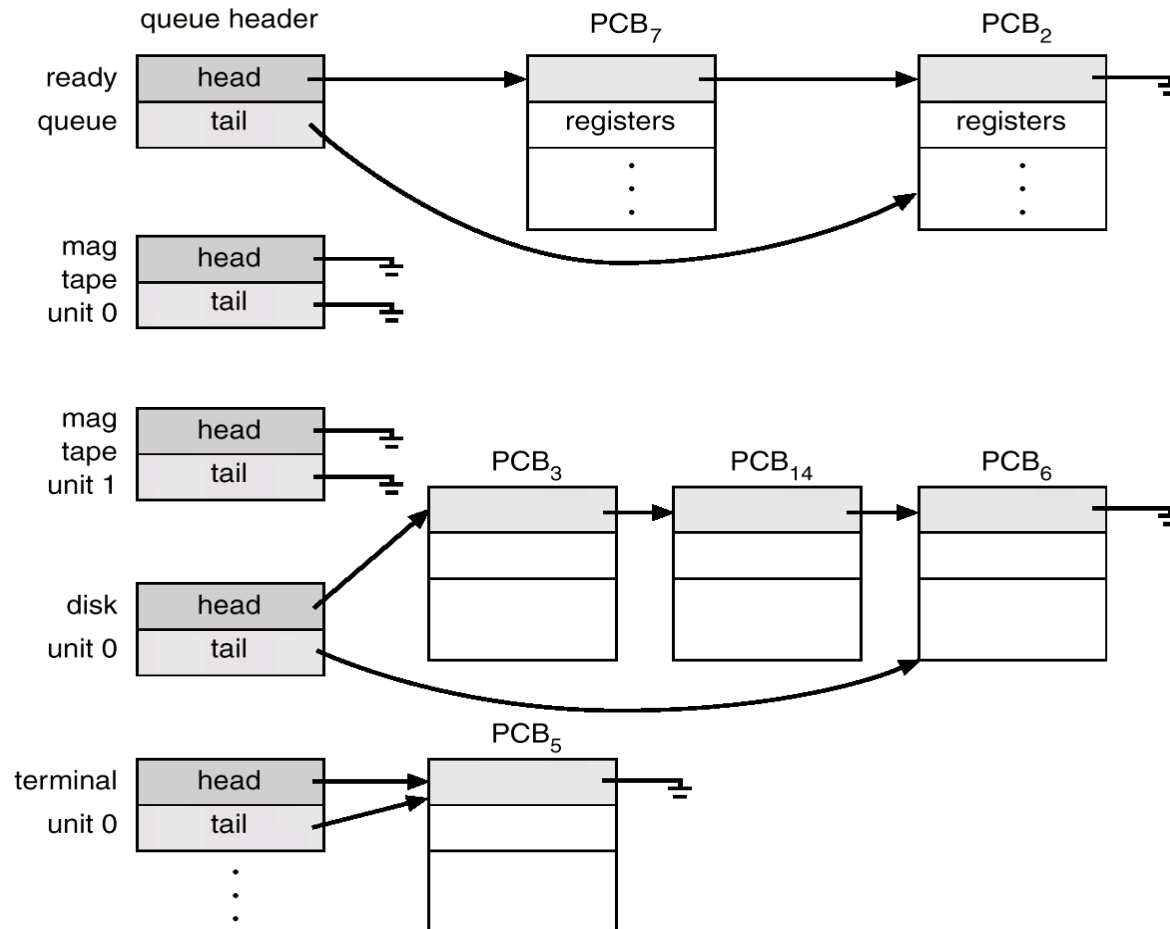


Blocked Queues



Further Enhancement:
A separate queue holds the processes waiting for different event.

Scheduling Queues



Schedulers

- Short term Scheduler or CPU Scheduling
 - Which program is to be run next
- Long term Scheduler or Job Scheduler
 - Which ready jobs should be brought to memory
 - May need to invoke only when a process leaves the system
 - Must make a careful selection

What's Happening at Process Creation?

- The OS:
 - Assigns a unique process identifier
 - Allocates space for the process
 - Initializes process control block
 - Creates or expand other data structures
- Traditionally, the OS created all processes
 - But it can be useful to let a running process create another
 - This action is called ***process spawning***
 - ***Parent Process*** is the original, creating, process
 - ***Child Process*** is the new process

OS Control Tables

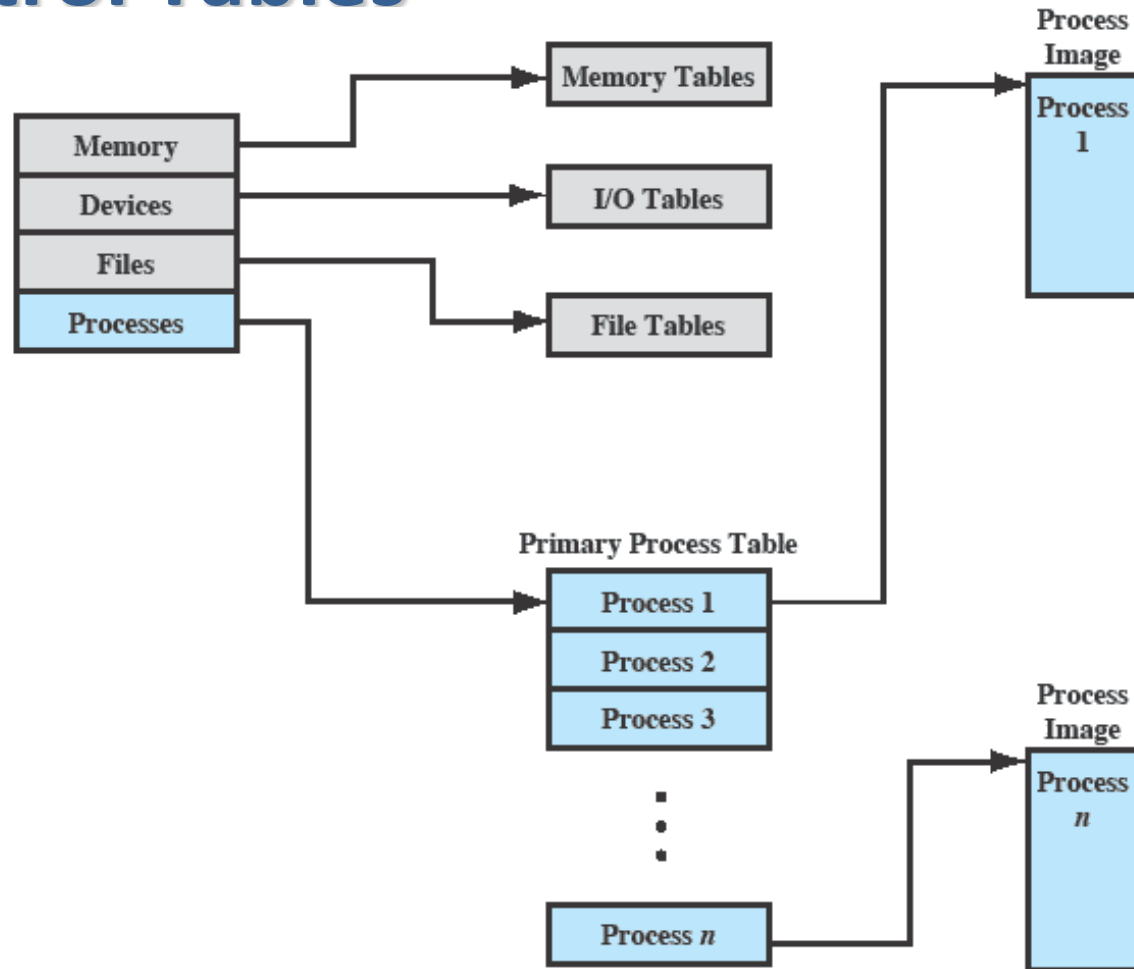


Figure 3.11 General Structure of Operating System Control Tables

Process Termination

- User logs off
- Process executes a service request to terminate
- Parent terminates so child processes terminate
- Operating system intervention
 - such as when deadlock occurs
- Error and fault conditions
 - E.g. memory unavailable, protection error, I/O failure, invalid instruction

Unix Process Creation

- When the system starts up it is running in kernel mode
- There is only one process, the kernel process.
- At the end of system initialization, the initial process starts up another kernel process.
 - The **init** kernel process has a process identifier of 1.

Process Creation

- These new processes may themselves go on to create new processes.
- All of the processes in the system are descended from the init kernel thread.
- You can see the family relationship between the running processes in a Linux system using the **ps tree** command
- A new process is created by a **fork()** system call

Compiling C++ code

- `g++ test.cpp -o Output`
- Running the code:
- `./Output`

The fork() system call

At the end of the system call there is a new process waiting to run once the scheduler chooses it

- A new data structure is allocated
- The new process is called the child process.
- The existing process is called the parent process.
- The parent gets the child's pid returned to it.
- The child gets 0 returned to it.
- Both parent and child execute at the same point after fork() returns

Unix Process Control

```
int main()
{
    int pid;
    int x = 0;

    x = x + 1;
    fork() ;
    x = 3;
    printf("%d",x) ;
}
```

But we want the child process to do something

```
int pid;  
int status = 0;  
  
pid = fork();  
if (pid > 0) {  
    /* parent */  
    .....  
    pid = wait(&status);  
} else {  
    /* child */  
    .....  
    exit(status);  
}
```

*The **fork** syscall returns a zero to the child and the child process ID to the parent*

*Parent uses **wait** to sleep until the child returns and*

***Fork** creates an exact copy of the parent process*

***Wait** variants allow wait on a specific child or*

*Child process passes status back to parent on **exit**, to report success/failure*

Child Process Inherits

- Stack
- Memory
- Environment
- Open file descriptors.
- Current working directory
- Resource limits
- Root directory

Child process DOESNOT Inherit

- Process ID
- Different parent process ID
- Process times
- Own copy of files
- Resource utilization (initialized to zero)

Lecture 02: Process Management

```
#include <iostream>
#include <unistd.h>

using namespace std;

int main()
{
    int pid;
    string x;

    fork();
    x = "hello";
    cout << "\n" << x;
}
```

Lecture 02: Process Management

```
#include <iostream>
#include <unistd.h>

using namespace std;

int main()
{
    int status;
    int pid = fork();
    if (pid == 0) {
        cout << "Hello - from child having id " <<
getpid() << " having parent " << getppid() << endl;
    }
    else {
        cout << "World - from Parent having id " <<
getpid() << endl;
    }
}
```

How can a parent and child process communicate?

- Through any of the normal IPC mechanism schemes.
- But have special ways to communicate
 - For example
 - The variables are replicas
 - The parent receives the exit status of the child

The wait() System Call

- A child program returns a value to the parent, so the parent must arrange to receive that value
- The wait() system call serves this purpose
 - `pid_t wait(int *status)`
 - it puts the parent to sleep waiting for a child's result
 - when a child calls `exit()`, the OS unblocks the parent and returns the value passed by `exit()` as a result of the wait call (along with the pid of the child)
 - if there are no children alive, `wait()` returns immediately
 - also, if there are zombies, `wait()` returns one of the values immediately (and deallocates the zombie)

What is a zombie?

- In the interval between the child terminating and the parent calling `wait()`, the child is said to be a 'zombie'.
- Even though its not running its taking up an entry in the process table.
- The process table has a limited number of entries.

What is a zombie?

- If the parent terminates without calling `wait()`, the child is adopted by `init`.

The solution is:

- Ensure that your parent process calls `wait()` or `waitpid` or etc, for every child process that terminates.

exit()

void exit(int *status*);

- After the program finishes execution, it calls *exit()*
- This system call:
 - takes the “result” of the program as an argument
 - closes all open files, connections, etc.
 - deallocates memory
 - deallocates most of the OS structures supporting the process
 - checks if parent is alive:
 - If so, it holds the result value until parent requests it, process does not really die, but it enters the zombie/defunct state
 - If not, it deallocates all data structures, the process is dead

Lecture 02: Process Management

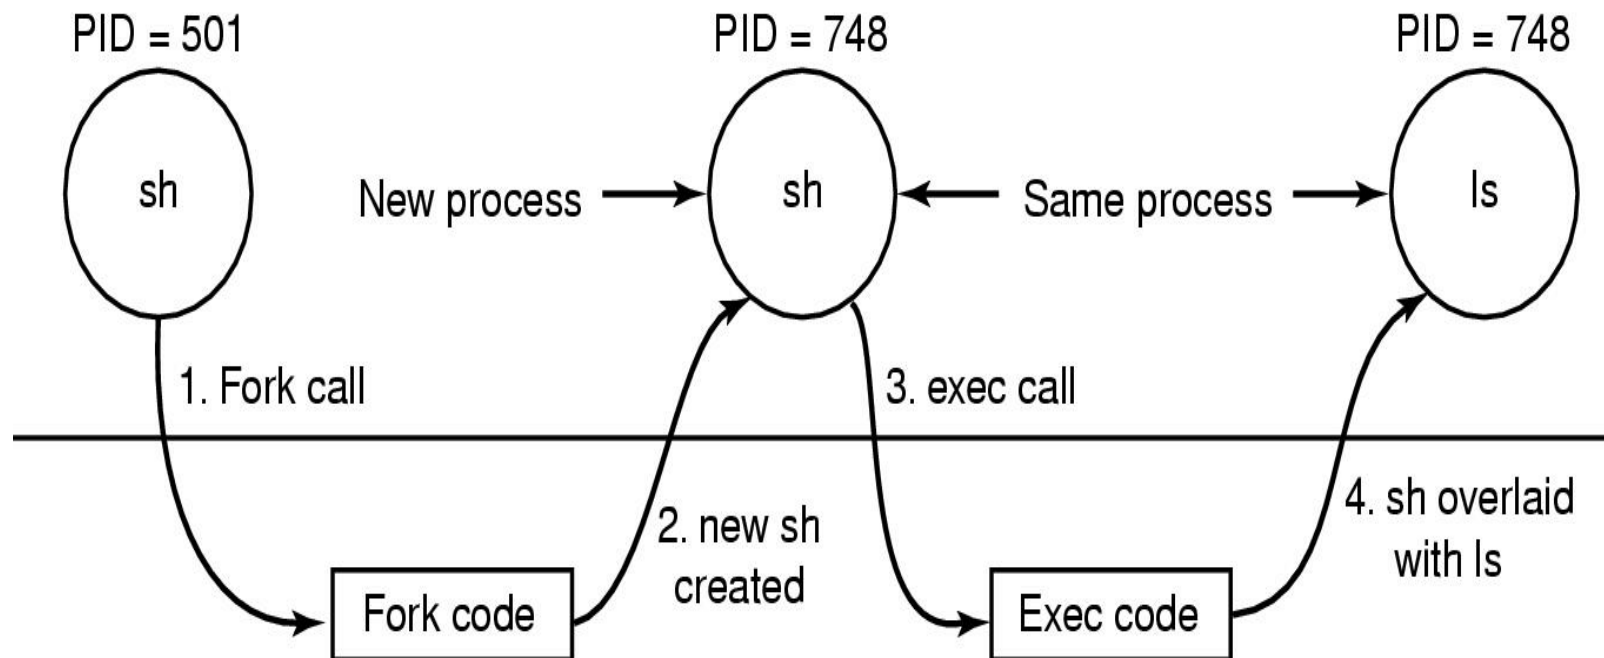
```
#include <iostream>
#include <unistd.h>
using namespace std;

int main()
{
    int status;
    int pid = fork();
    if (pid == 0) {
        cout << "Hello - from child having id " << getpid()
        << " having parent " << getppid() << endl;
        exit(2);
    }
    else {
        cout << "World - from Parent having id " <<
        getpid() << endl;
        wait(&status);
        cout << WEXITSTATUS(status) << endl;
    }
}
```

exec()

- We usually want the child process to run some other executable
- For Example, *ls - list directory contents*
- Overlays, performed by the family of exec system calls, can change the executing program, but can not create new processes.

The *ls* Command



Steps in executing the command *ls* type to the shell

execv

- `int execv(const char *path, char *const argv[]);`
- the current process image with a new process image.
- **path** is the filename to be executed by the child process
- When a C-language program is executed as a result of this call, it is entered as a C-language function call as follows:
 - `int main (int argc, char *argv[]);`
- The **argv** array is terminated by a null pointer.
- The null pointer terminating the **argv** array is not counted in **argc**.

Lecture 02: Process Management

//helloAOS.cpp

```
#include <iostream>  
#include <unistd.h>
```

```
#include <iostream>  
using namespace std;
```

```
using namespace std;
```

```
int main()  
{
```

```
    int status;  
    int pid = fork();  
    if (pid > 0) {
```

```
        cout << "World - from Parent having id " << getpid()  
<< endl;
```

```
    }
```

```
    else {
```

```
        execv("./helloAOS", NULL);
```

```
    }
```

```
int main()  
{
```

```
    cout << "Hello AOS  
\n";
```

```
    return 0;
```

```
}
```


Assignment-01

- Will be announced in the next class.