

**Lecture 04**  
**Inter-Thread and**  
**Inter-Process Communication**

**Dr. Ehtesham Zahoor**

# Operating systems research

- How to choose a good research paper?

### Some review

- Why use threads instead of processes?
- Is there any relation between threads and Docker containers?
- How to create threads?

# The Pthreads API

- The most important of thread APIs, in the Unix world, is the one developed by the group known as POSIX.
- POSIX is a standard API supported
- Portable across most UNIX platforms.
- PTHREAD library contains implementation of POSIX standard
- To link this library to your program use *-lpthread*
  - `gcc MyThreads.c -o MyThreadExecutable - lpthread`

# Thread Creation

```
pthread_create( pthread_t *threadid  
    , const pthread_attr_t *attr, void  
    * (*start_routine) (void *) , void *arg) ;
```

- This routine creates a new thread and makes it executable.
- Thread stack is allocated and thread control block is created
- Once created, a thread may create other threads.
- Note that an "initial thread" exists by default and is the thread which runs main.
- Returns zero, if ok
- Returns Non-zero if error

# Thread Creation

```
pthread_create( pthread_t *threadid  
    , const pthread_attr_t *attr, void  
    *(*start_routine) (void *) , void *arg) ;
```

- **threadid**
  - The routine returns the new thread ID via the threadid
  - The caller can use this thread ID to perform various operations
  - This ID should be checked to ensure that the thread was successfully created.
- **attr**
  - used to set thread attributes.
  - NULL for the default values.
- **start\_routine**
  - The C routine that the thread will execute once it is created.
- **arg**
  - Arguments are passed to *start\_routine* via *arg*.
  - Arguments must be passed by reference as pointers
  - These pointers must be cast as pointers of type void.


# Thread Creation

```
#include <pthread.h>  
#include <iostream>  
using namespace std;
```

```
void* PrintHello(void* arg)  
{  
    cout << "Hello World! " << endl;  
    pthread_exit(NULL);  
}  
int main()  
{  
    pthread_t threadID;  
    pthread_create(&threadID, NULL, PrintHello, NULL);  
    cout << "Hello World! " << endl;  
  
    pthread_exit(NULL);  
}
```

# Thread Creation

- Once created, threads are peers, and may create other threads. There is no implied hierarchy or dependency between threads.





# Passing Arguments

```
#include <pthread.h>
#include <iostream>
using namespace std;

void* PrintHello(void* arg) {
    cout << *(string*) arg << endl;
    pthread_exit(NULL);
}

int main()
{
    pthread_t threadID;
    string threadArg = "Hello";
    pthread_create(&threadID, NULL, PrintHello, (void*)&threadArg);
    pthread_exit(NULL);
}
```

# Passing Arguments To Threads

- The `pthread_create()` routine permits the programmer to pass one argument to the thread start routine.
- What if we want to pass multiple arguments.
- Create a structure which contains all of the arguments
- Pass a pointer to the structure in the `pthread_create()` routine.
- Argument must be passed by reference and cast to `(void *)`.

# Passing Arguments – The wrong way

```
#include <pthread.h>
#include <iostream>
using namespace std;

void* PrintHello(void* arg)
{
    for (int counter=0; counter<5;counter++) {
        cout << *(string*) arg;
    }
    pthread_exit(NULL);
}

int main()
{
    pthread_t threadID;
    string threadData = "Hello";
    pthread_create(&threadID, NULL, PrintHello, (void*)&threadData);
    threadData = "World";
    pthread_create(&threadID, NULL, PrintHello, (void*)&threadData);

    pthread_exit(NULL);
}
```

# The problem?

- Threads initially access their data structures in the parent thread's memory space.
- That data structure must not be corrupted/modified until the thread has finished accessing it.

# Passing Arguments – better approach

```

#include <pthread.h>
#include <iostream>
using namespace std;
#define NUM_THREADS 3

void* PrintHello(void* arg)
{
    for (int counter=0; counter<2;counter++) {
        cout << *(string*) arg <<endl; }
    pthread_exit(NULL);
}


int main()
{
    pthread_t threadID[NUM_THREADS];
    string threadData[NUM_THREADS] = {"Hello", "ez", "World"};

    for (int counter=0; counter<NUM_THREADS;counter++) {
        pthread_create(&threadID[counter], NULL, PrintHello,
        (void*)&threadData[counter]);
    }
    pthread_exit(NULL);
}

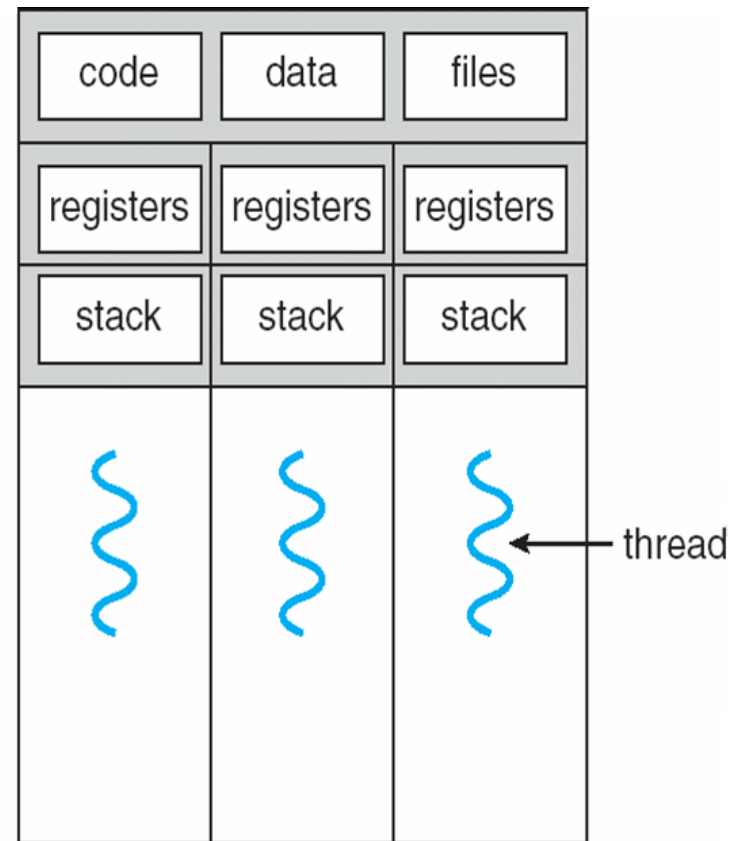
```

# **Inter-Thread Communication**

## Inter-Thread Communication



single-threaded process



multithreaded process

# Threads share Global variables!

```
#include <pthread.h>
#include <iostream>
#include <unistd.h>
using namespace std;
#define NUM_THREADS 10

int sharedData = 0;
void* incrementData(void* arg) {
    sharedData++;
    pthread_exit(NULL);
}
int main()
{
    pthread_t threadID;
    for (int counter=0; counter<NUM_THREADS;counter++) {
        pthread_create(&threadID, NULL, incrementData, NULL);
    }
    cout << "ThreadCount:" << sharedData <<endl;
    pthread_exit(NULL);
}
```



# What should be the output?

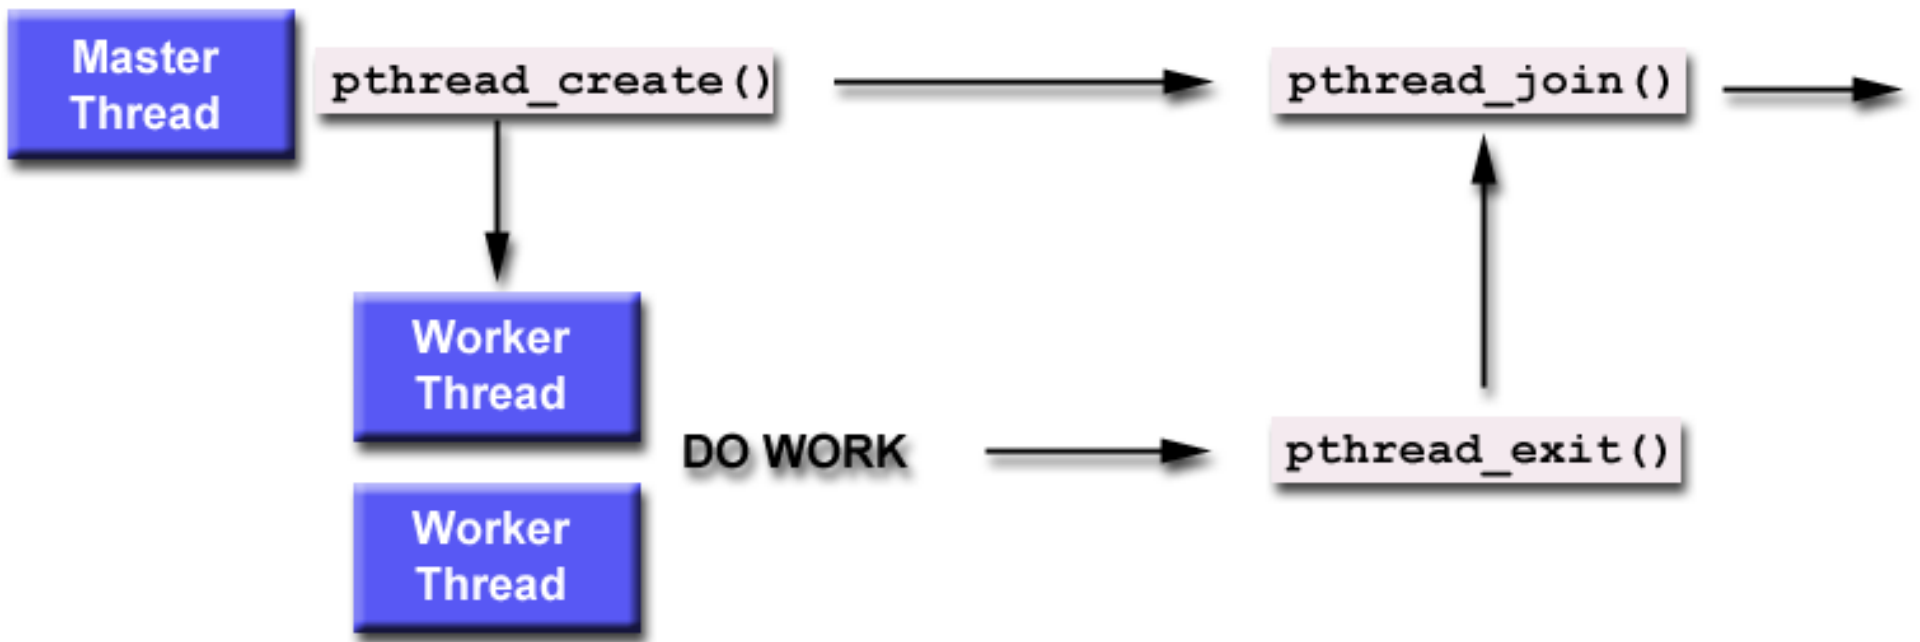
```
>./5globalData  
ThreadCount:10
```

```
>./5globalData  
ThreadCount:8
```

# Thread Suspension and Termination

- Similar to UNIX processes, threads have the equivalent of the `wait()` and `exit()` system calls
  - `pthread_join()` Used to block threads
  - `pthread_exit()` Used to terminate threads
- To instruct a thread to block and wait for a thread to complete, use the **`pthread_join()`** function.
- Any thread can call join on (and hence wait for) any other thread.

# Thread Suspension and Termination



# Joining thread

- \* **Joinable**: on thread termination the thread ID and exit status are saved by the OS.
- \* Joining a thread means waiting for a thread
- \* **pthread\_join(threadid, status)**
  - \* "Joining" is one way to accomplish synchronization between threads.
  - \* subroutine blocks the calling thread until the specified *threadid* thread terminates.
    - \* The programmer is able to obtain the target thread's termination return status (if specified) in the *status* parameter.
- \* It is impossible to join a detached thread

# ThreadCount: A better implementation

```

#include <pthread.h>
#include <iostream>
#include <unistd.h>
using namespace std;
#define NUM_THREADS 100
int sharedData = 0;
void* incrementData(void* arg)
{
    sharedData++;
    pthread_exit(NULL); }
int main()
{
    pthread_t threadID[NUM_THREADS];
    for (int counter=0; counter<NUM_THREADS;counter++) {
        pthread_create(&threadID[counter], NULL, incrementData, NULL);
    }
    //waiting for all threads
    int statusReturned;
    for (int counter=0; counter<NUM_THREADS;counter++) {
        pthread_join(threadID[counter], NULL);
    }
    cout << "ThreadCount:" << sharedData <<endl;
    pthread_exit(NULL);
}

```

# Is the problem solved?

- Unfortunately, not yet :(
- The output from running it with 1000 threads is as below:

```
>./6join  
ThreadCount:990  
>./6join  
ThreadCount:978  
>./6join  
ThreadCount:1000  
>
```

- Reasons? What can be done?
- Lets postpone this discussion till synchronization class.

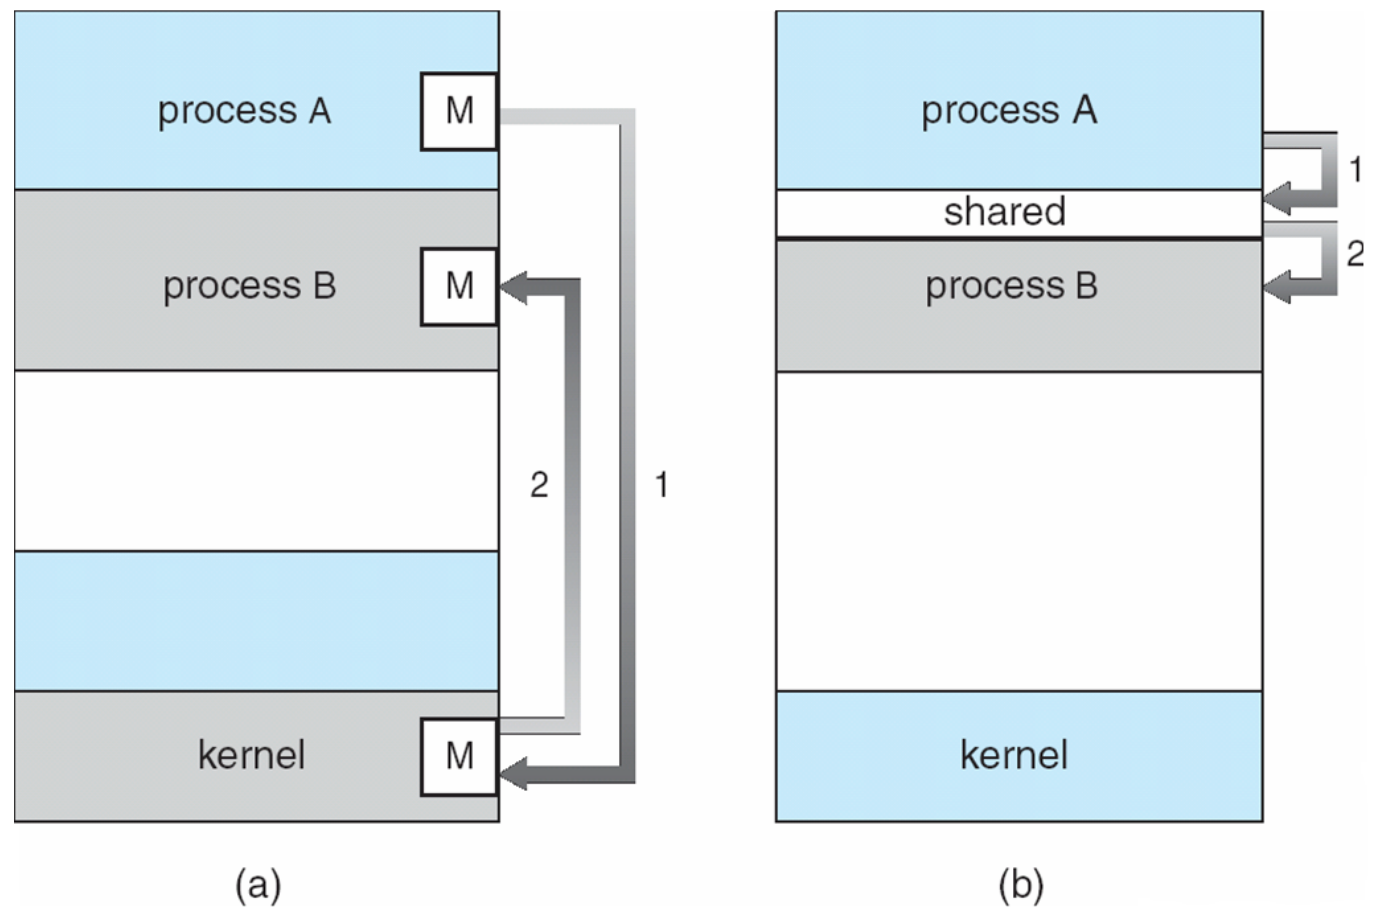
# **Inter-Process Communication**

# Interprocess Communication

- A process has access to the memory which constitutes its own address space.
- So far, we have discussed communication mechanisms only during process creation/termination
  - The parent receives the exit status of the child
- Processes may need to communicate during their life time.



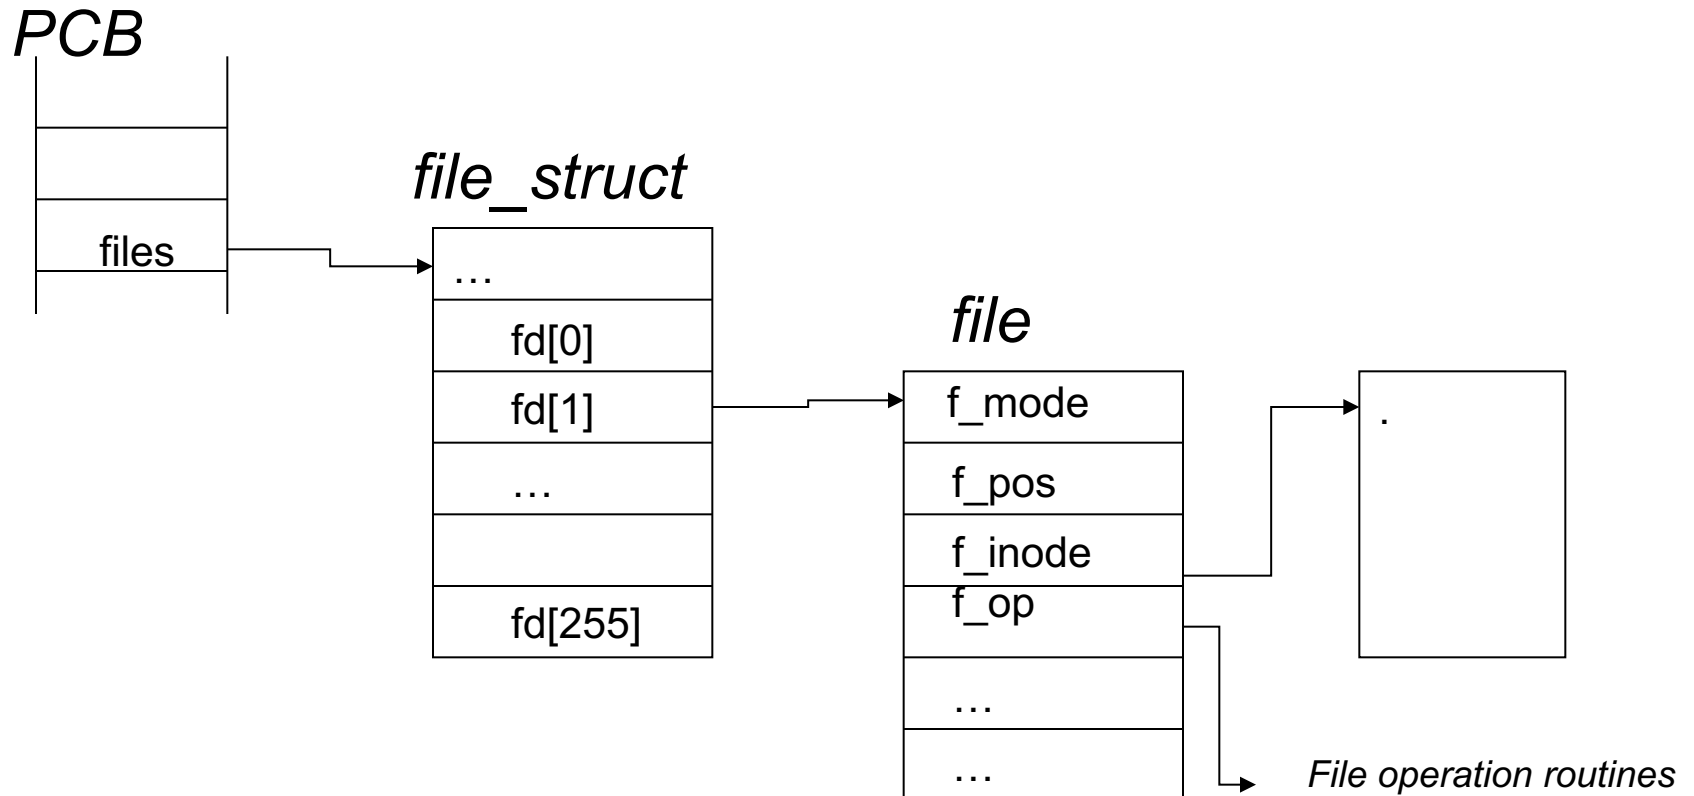
# Communications Models



# **UNIX IPC - Pipes**

# File Descriptors

- The PCB of each process contains a pointer to a *file\_struct*, a kernel-resident array data structure containing the details of open files



# File Descriptors

- The *files\_struct* contains pointers to file data structures
- Each one describes a file being used by this process.
- *f\_mode*: describes file mode, read only, read and write or write only.
- *f\_pos*: holds the position in the file where the next read or write operation will occur.
- *f\_inode*: points at the actual file

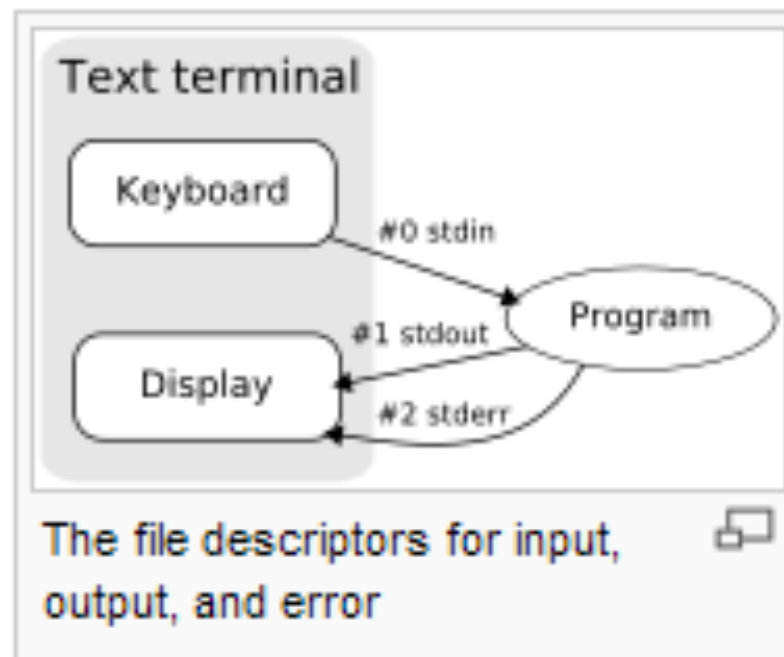
# File Descriptors

- Every time a file is opened, one of the free file pointers in the *files\_struct* is used to point to the new file structure.
- Linux processes expect three file descriptors to be open when they start.
- These are known as *standard input*, *standard output* and *standard error*

# File Descriptors

- The program treat them all as files.
- These three are usually inherited from the creating parent process.
- All accesses to files are via standard system calls which pass or return file descriptors.
- *standard input, standard output and standard error* have file descriptors 0, 1 and 2.

## Lecture 04: Inter-Thread and Inter-Process Communication



# File Descriptors

- `char buffer[10];`
- Read from standard input (by default it is keyboard)
  - `read(0,buffer,5);`
- Write to standard output (by default is is monitor))
  - `write(1,buffer,5);`
- By changing the file descriptors we can write to files



## Lecture 04: Inter-Thread and Inter-Process Communication

```
#include <iostream>  
#include <unistd.h>  
  
using namespace std;  
  
int main()  
{  
  char myBuffer[10];  
  cout << "Please enter your name: \n";  
  int count = read(0,myBuffer,10);  
  write(1,myBuffer,count);  
  
}
```

## The Unix fact

- When Unix programs do any sort of I/O, they do it by reading or writing to a file descriptor
- A file descriptor is simply an integer associated with an open file

# The Unix fact

- A file in Unix can be
  - A network connection
  - A pipe
  - A terminal
  - A real on-the-disk file
  - Or just about anything else


***EVERYTHING* IN UNIX IS A FILE!**

# Pipes

- Provides an interprocess communication channel



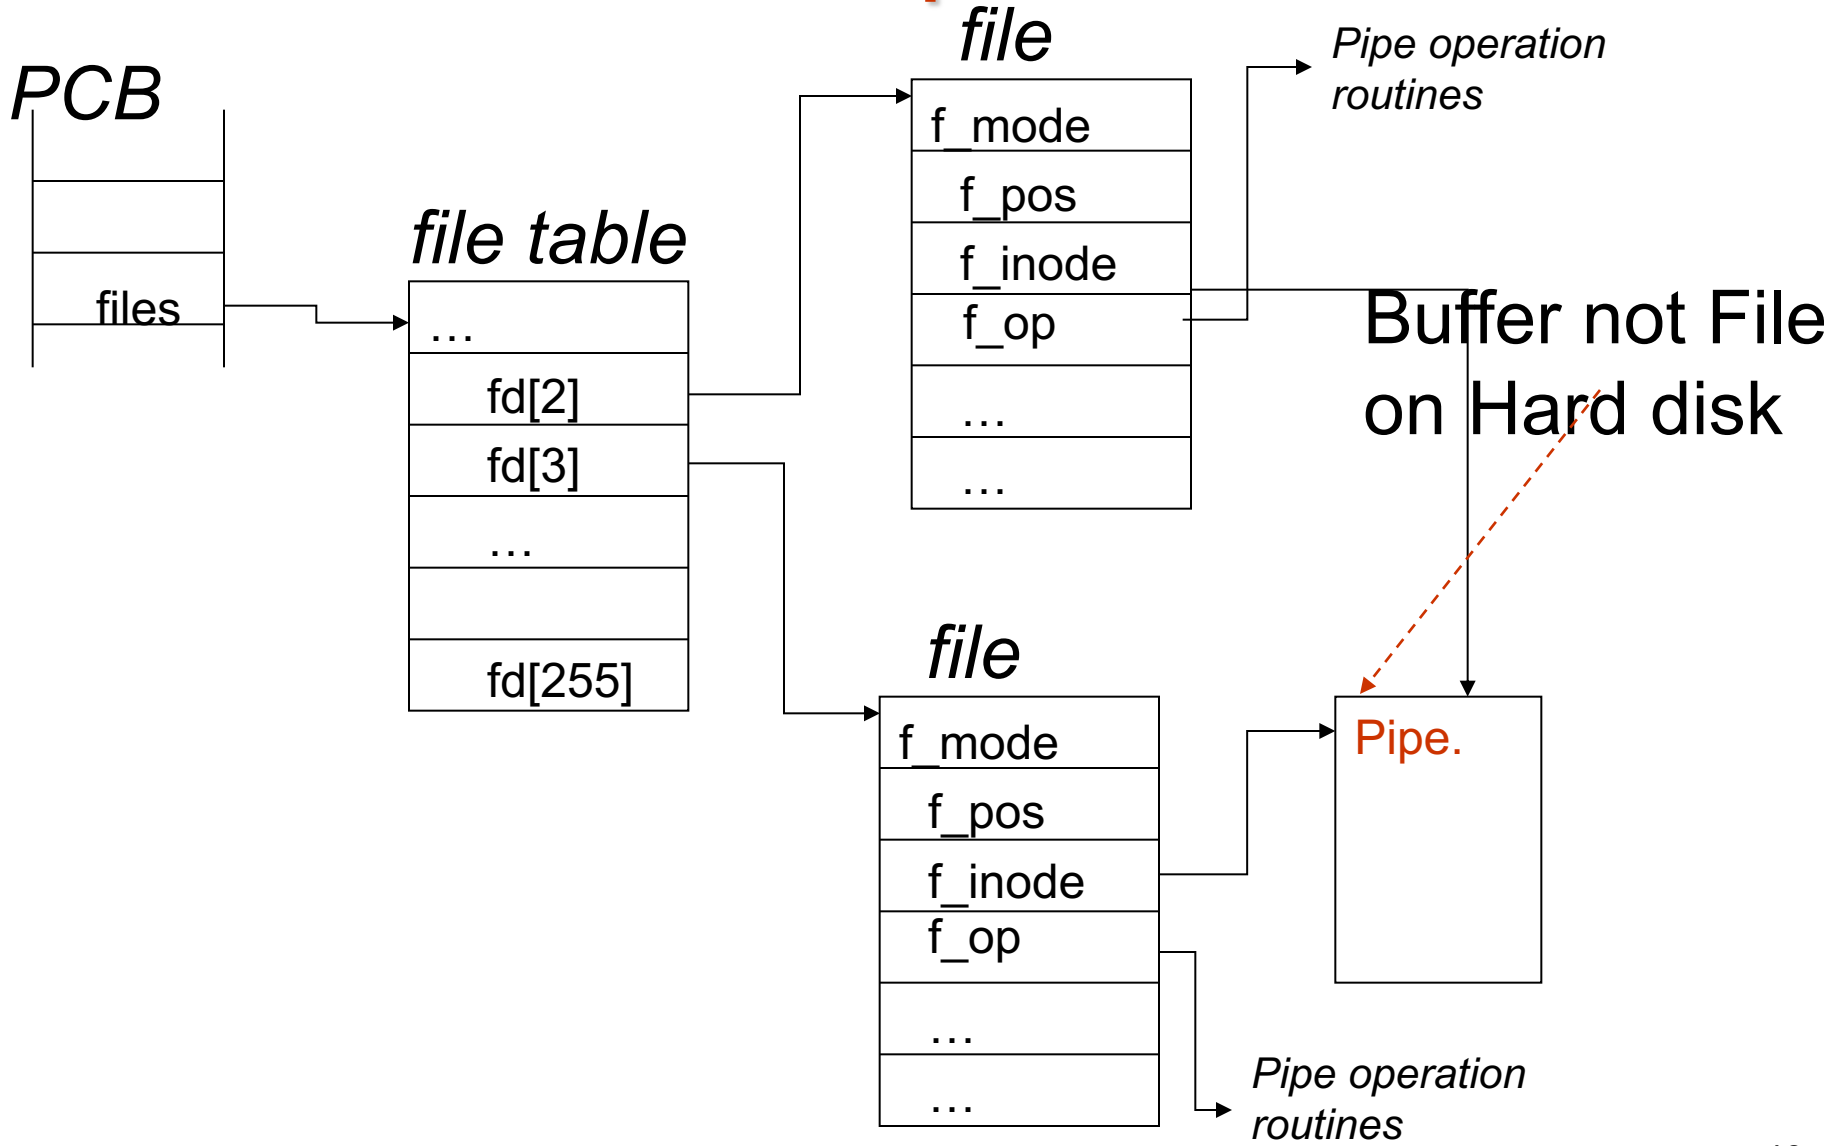
# Pipes: Shared info in kernel's memory



# Pipes

- A pipe is implemented using two file data structures which both point at the same temporary data node.
- This hides the underlying differences from the generic system calls which read and write to ordinary files
- Thus, reading/writing to a pipe is similar to reading/writing to a file

# Pipes





# Pipe Creation

```
#include <unistd.h>
```

```
int pipe(int filedes[2]);
```

- Creates a pair of file descriptors pointing to a pipe inode
- Places them in the array pointed to by *filedes*
  - *filedes[0]* is for reading
  - *filedes[1]* is for writing.
- On success, zero is returned.
- On error, -1 is returned

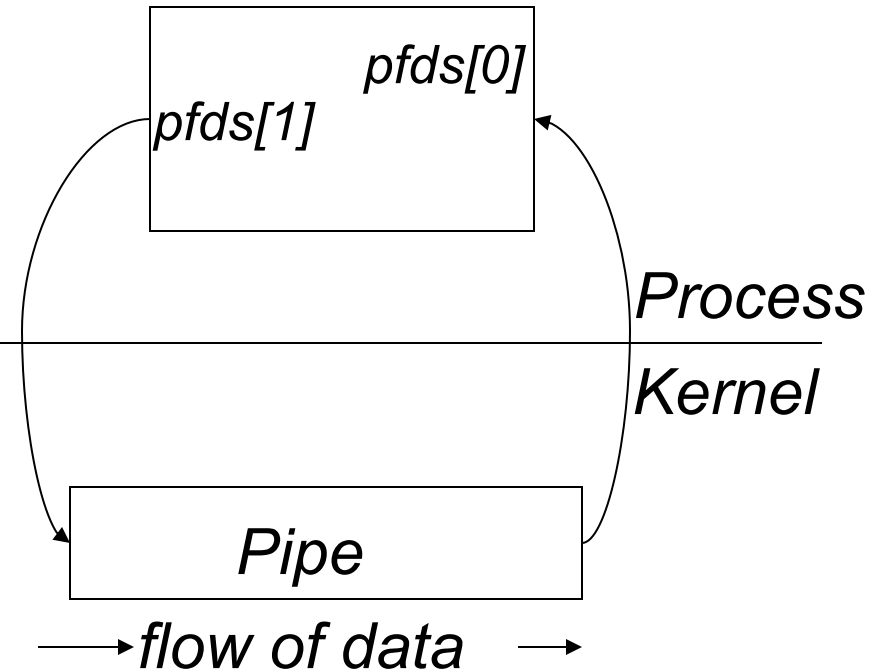
# Lecture 04: Inter-Thread and Inter-Process Communication

```
#include <iostream>
#include <unistd.h>
```

```
using namespace std;
```

```
int main()
{
    int pfd[2];
    if (pipe(pfd) != -1){
        cout << pfd[0] << pfd[1];
    } else {
        cout << "Error creating pipe!";
        exit(1);
    }
}
```

```
return 0;
}
```



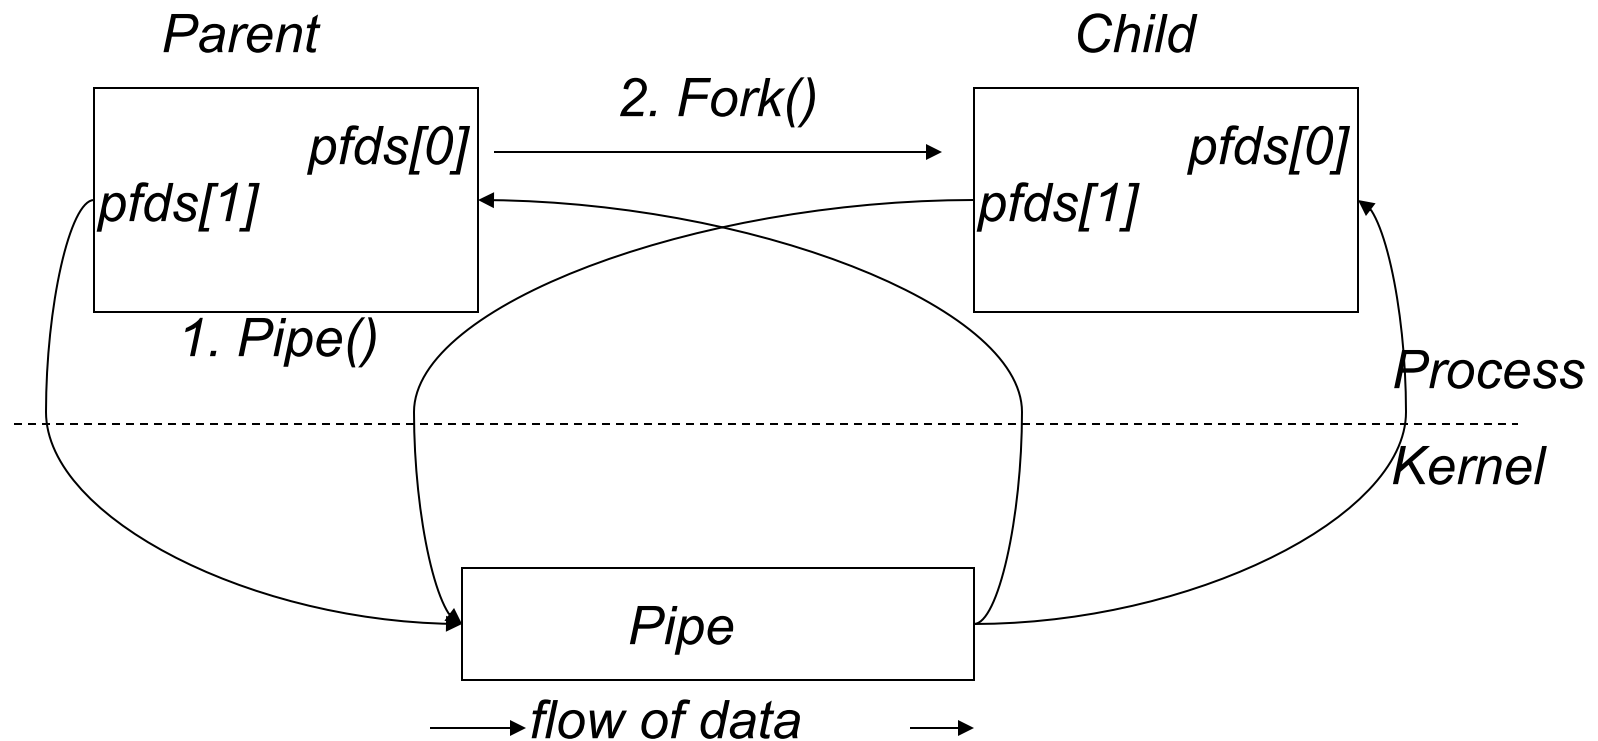
# A Channel between two processes

- Remember: the two processes have a parent / child relationship
- The child was created by a `fork()` call that was executed by the parent.
- The child process is an image of the parent process
- Thus, all the **file descriptors** that are opened by the parent are now available in the child.

## A Channel between two processes

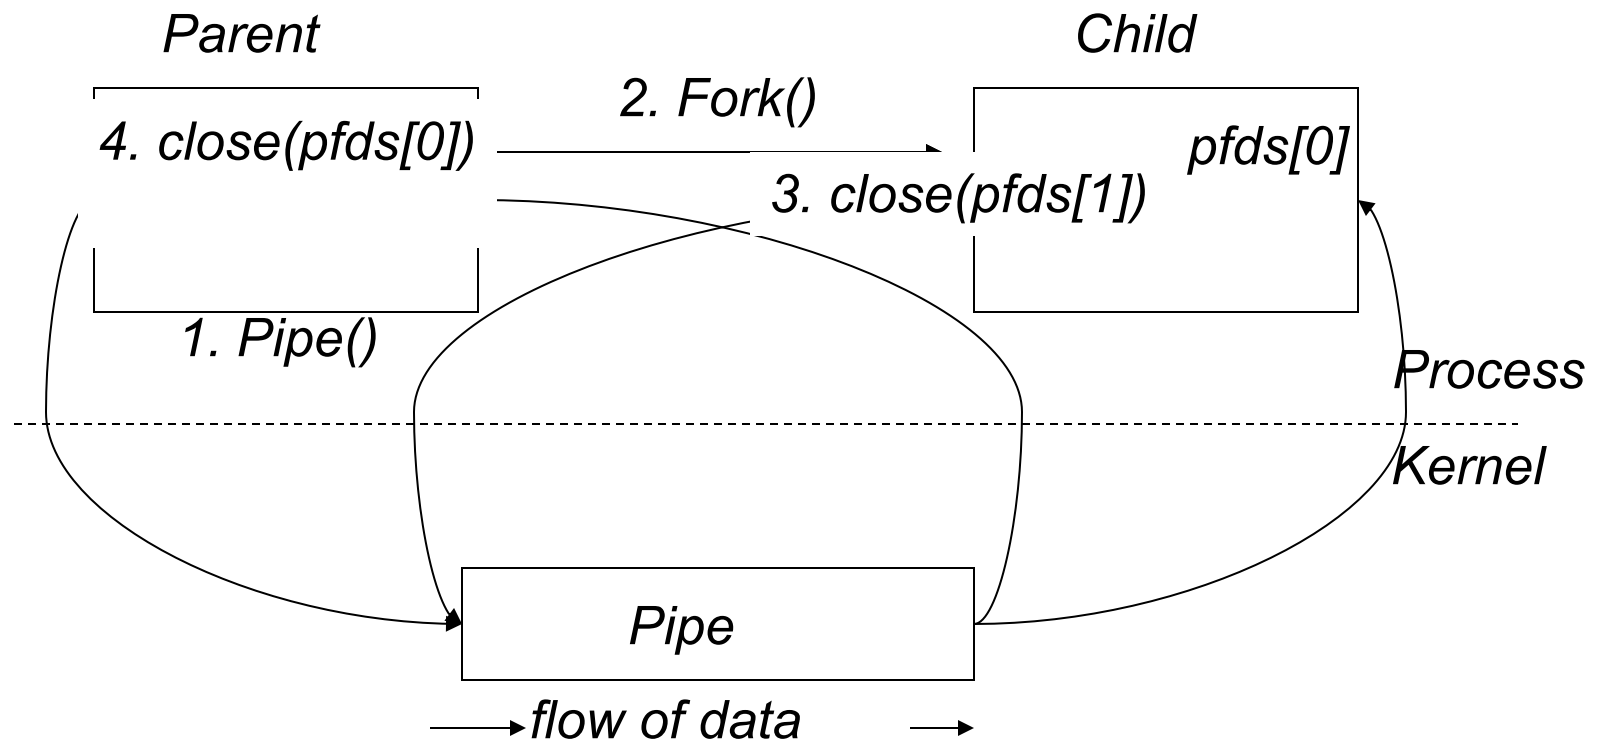
- The **file descriptors** refer to the same I/O entity, in this case a pipe.
- The pipe is inherited by the child
- And may be passed on to the grand-children by the child process or other children by the parent.

## A Channel between two processes



# A Channel between two processes

- To allow one way communication each process should close one end of the pipe.



## Closing the pipe

- The **file descriptors** associated with a pipe can be closed with the `close(fd)` system call
- How would we achieve two way communication

## Lecture 04: Inter-Thread and Inter-Process Communication

```
int main()
{
    int pfd[2];
    char buf[30];
    if (pipe(pfd) != -1){
        if (!fork()) {
            close(pfd[0]);
            cout << "CHILD: writing to the pipe\n";
            write(pfd[1], "Sample Text", 11);
            cout << "CHILD: exiting\n";
            exit(0);
        } else {
            close(pfd[1]);
            cout << "PARENT: reading from pipe\n";
            read(pfd[0], buf, 11);
            cout << "PARENT has read: " << buf;
            wait(NULL);
        }
    } else {
        cout << "Error creating pipe!";
        exit(1);
    }
    return 0; }
```